

Nonblocking commit protocols

Dale Skeen, SIGMOD'81

Jingchao Fang, Zhuoer Tong

Abstract

- Protocols that allow operational sites to continue transaction processing even though site failures have occurred are called **nonblocking**.
- This paper investigates the properties of nonblocking protocols.
- This paper also introduced central site nonblocking protocol and decentralized site nonblocking protocol.

Background

- Researchers have started to focus on distributed database systems since *late 70s*.
- Distributed crash recovery is vital.
- Some operations on the data are logically indivisible, and these operations are called **transactions**.

Background

- Atomic operation:** either executes to completion or it appears never to have executed. By definition, transaction on distributed database system is atomic operation.
- In reality, a transaction is *rarely* a physically atomic operations.
- The gap between logical atomicity and physical atomicity causes significant problems in distributed database implementation.

Background

- Preserving transaction atomicity is well understood in *single site case*.
- Site decides to **commit** or **abort** when reaching **commit point**.
- Unconditionally guaranteed, irreversible.
- Problem occurs when more than one site is involved.

Two Phase Commit Protocol

-The simplest commit protocol that allows unilateral abort.

-An example of a blocking protocol.

SITE 1

(1) Transaction is received.
"Start Xact" is sent.

(2) The vote is received.
If vote="yes" and site 1 agrees,
then "commit" is sent;
else, "abort" is sent.

SITE 2

"Start Xact" is received.
Site 2 votes:
 "yes" to commit,
 "no" to abort.
The vote is sent to site 1.

Either "commit" or "abort" is
received and processed.

Two Phase Commit Protocol: Process

-First Phase: coordinator

distribute the transaction to all sites, and each site individually votes.

-Second Phase: coordinator

collects all the votes and informs outcome.

SITE 1

- (1) Transaction is received.
"Start Xact" is sent.

- (2) The vote is received.
If vote="yes" and site 1 agrees,
then "commit" is sent;
else, "abort" is sent.

SITE 2

"Start Xact" is received.
Site 2 votes:
 "yes" to commit,
 "no" to abort.
The vote is sent to site 1.

Either "commit" or "abort" is
received and processed.

Two Phase Commit Protocol: Summary

- Simplest and cheapest in number of messages, but can be block on site failures.
- An example of a **blocking protocol**: operational sites sometimes wait on the recovery of a failed sites. *Locks* must be held on the database while the transaction is blocked.
- A protocol that never requires operational sites to block until a failed site has recovered is called a **nonblocking protocol**.

Termination and Recovery Protocols

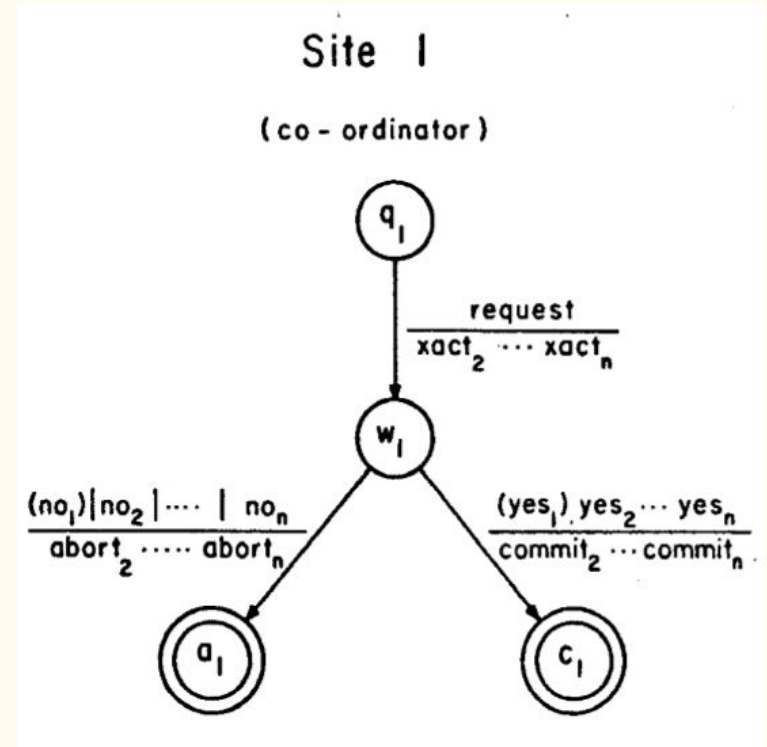
- Termination Protocol:** to terminate transaction execution as quickly as possible at the operational sites.
- Recovery Protocol:** invoked by failed sites to resume transaction processing. (not discussed in this paper)

Formal Model Describing Commit Protocols

- Transaction execution at each site models as **finite state automaton** (FSA).
- Network models as input/output tapes to all sites.
- The states of the FSA for site i are called **local states** of site i .
- State of transition: reading, writing, and moving to the next local state.


FSA for Two Phase Commit Protocol

- Each FSA has four (local) states: initial (q_i), wait (w_i), abort (a_i), and commit (c_i).
- Abort and commit are final states.



FSA for Two Phase Commit Protocol

- Local state for site i are subscripted with i .
- Messages sent or received by a slave are subscripted with that slave's site number.



Properties for Commit Protocol

- The FSA's are nondeterministic.
- The final states of the FSA's are partitioned into two sets: the abort states, and the commit states.
- The act of committing or aborting is irreversible.
- The state diagram describing a FSA is acyclic.

Global Transaction States

- Global state** defines the complete processing state of a transaction.
- A global state is said to be **inconsistent** if it contains both a local commit state and a local abort state.

Global Transaction States

- A global state is said to be a **final state** if all local states contained in the state vector are final states.
- A global state is said to be a **terminal state** if from it there is no immediately reachable successors.
- A terminal state that is not a final state is a **deadlock state**.
- The **concurrency set** of s_i state is the set of all local states s_j , where $i \neq j$, such that s_i and s_j are contained in the same reachable global state.

Committable States

- A local state is called **committable** if occupancy of that state by any site implies that all sites have voted yes on committing the transaction.
- For **nonblocking protocol**: have *more than one* committable state (assert in paper and without proof)

The Central Site Class Commit Protocol

-The central site class: uses one site (**coordinator**) to direct transaction processing at all participating sites (**slaves**).

-Properties:

- 1) Single coordinator;
- 2) All other participants execute the slave protocol;
- 3) Slaves can communicate only with coordinator;
- 4) Coordinator sends the same message and wait.

The Central Site Class Commit Protocol


- Advantages:** Relatively cheap, conceptually simple.
- Disadvantage:** Vulnerability to a coordinator failure.
- Synchronous within one state transition:** one site never leads another site by more than one state transition during the execution of the protocol.

The Decentralized Class Commit Protocol

- The (fully) decentralized class: each site participates as an equal in the protocol and executes the same protocol.
- Every site communicates with every other site.
- Process:** each site will send the identical message to every other site, then waits until it has received messages from all its cohorts.

Decentralized Two Phase Commit Protocol

- Simplest decentralized commit protocol.
- First phase:** each site receives the “start xact”, make the decision and sends to other cohorts.
- Second phase:** each site accumulates all the abort decisions and move to a final state.
- Also synchronous within one state transition.



The Fundamental Nonblocking Theorem

- Blocking situation arises.

- The fundamental nonblocking theorem.

- A protocol is **nonblocking** *if and only if* it satisfies:


- 1) there exists no local state such that its concurrency set contains both an abort and a commit state;

- 2) there exist no noncommitable state whose concurrency set contains a commit state.

The Canonical Two Phase Commit Protocol


-The concurrency set of state **q** contains **q**, **w** and **a**. The concurrency set of state **w** contains all of the local states of the protocol.

-Lemma: A protocol that is synchronous within one state transition is nonblocking if and only if: 1) it contains no local state adjacent to both a commit and an abort state, and 2) it contains no noncommittable state that is adjacent to a commit state.



Buffer States and Canonical Nonblocking Protocol

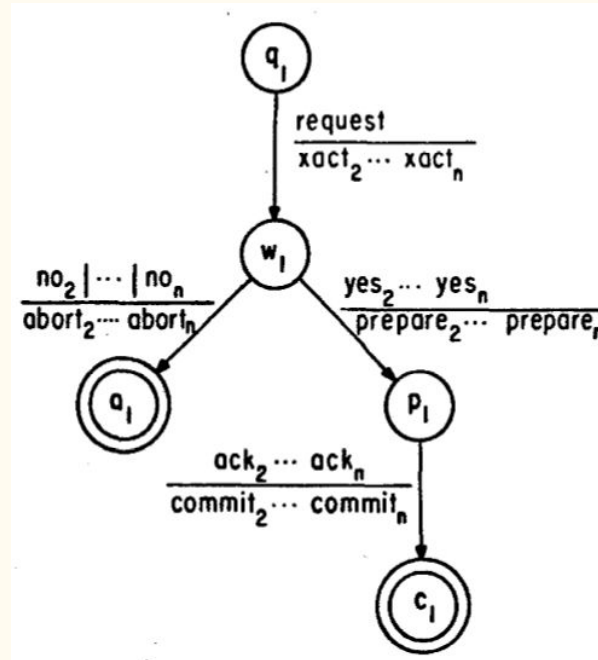
- The buffer state can be thought of a “prepare to commit” state (labelled as **p** in the graph).
- This protocol can be referred as *canonical nonblocking protocol*, which is a three phase protocol.



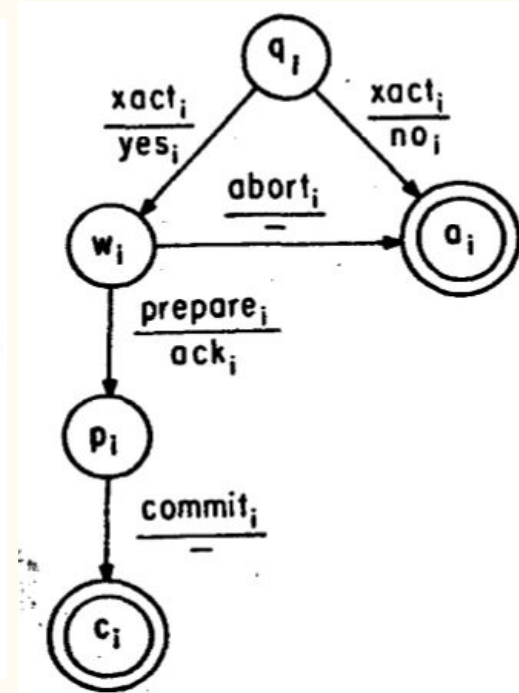
Nonblocking Central Site Protocol

-The slave protocol is the three phase protocol.

-The coordination protocol is also a three phase protocol that is an extension of the two phase coordinator protocol.



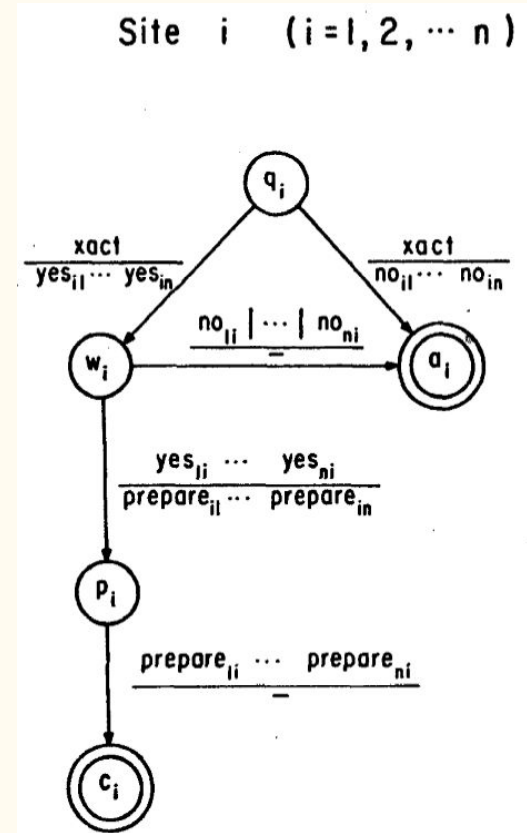
coordination



slave

Nonblocking Decentralized Protocol

-All protocols are the canonical nonblocking protocol.



Termination Protocols

- Site failures may occur and would make the continued execution of the commit protocol impossible. *Solution*: termination protocol.
- A termination protocol can accomplish its task only if the current state of at least one operational site obeys the conditions given in the fundamental nonblocking theorem.
- But in the worst case, it will be able to terminate correctly only if all of the operational sites obey the fundamental nonblocking theorem.

Central Site Termination Protocol and Backup

- Choose a **backup coordinator** from the set of operational sites.
- The backup coordinator will complete the transaction by directing all the remaining sites toward a commit or an abort. The protocol must be reentrant.
- Decision Rule For Backup Coordinators:** If the concurrency set for the current state of the backup contains a commit state, then the transaction is committed. Otherwise, it is aborted.
- Phase 1:** Issue a message and wait.
- Phase 2:** Issue a commit or abort.

Summary

- Two phase commit protocol
- Two most popular commit classes: central site and decentralized
- Fundamental nonblocking theorem
- Three phase commit protocol
- Termination protocol