

Brewer's conjecture and the feasibility of consistent, available, partition-tolerant web services

Presented by
Vikram Rao & Steven Guan
UC DAVIS





Figure: Prof. Eric Brewer, University of California, Berkeley



PODC 2000 Invited Talk

July 19, 2000

Towards Robust Distributed Systems

Eric A. Brewer

University of California, Berkeley



Motivation

Prof. Brewer's talk:

Current distributed systems, even the ones that work, tend to be very fragile: they are hard to keep up, hard to manage, hard to grow, hard to evolve, and hard to program

In this talk, I look at several issues in an attempt to clean up the way we think about these systems

...



Motivation

Prof. Brewer's talk:

These are not (yet) provable principles, but merely ways to think about the issues that simplify design in practice

They draw on experience at Berkeley and with giant-scale systems built at Inktomi, including the system that handles 50% of all web searches

...



Motivation



Figure: New Game APP



Motivation



Figure: Multi Player User Game APP



Motivation



Figure: Your MMORPG Game APP is awesome and more Users play!



Motivation



Connection lost.
Please wait - attempting to reestablish.

Figure: You scale up then Server crashes frequently, users complaint server issues and you don't know how to solve this issue perfectly!

- ?? Provide users 24/7 game run despite high load ??
- ?? Provide users latest game updates ??
- ?? Equally accessible to users across globe ??



Contents

1 Formal Model

2 Cases

3 Conclusion

4 References



Atomic Data Objects

There must exist a total order on all operations such that each operation looks as if it were completed at a single instant. for ex: return the right response to each request.



Available Data Objects

Every request received by a non-failing node in the system must result in a response. for ex: any service must eventually terminate



Partition Tolerance

When the network is partitioned all messages sent from nodes in one partition to nodes in another partition are lost

Connection lost.
Please wait - attempting to reestablish.



CAP vs ACID



Figure: CAP is not made of ACID!



Contents

1 Formal Model

2 Cases

3 Conclusion

4 References



Networks – Asynchronous

It is impossible in the asynchronous network model to implement a read/write data object that guarantees the following properties:

- Availability
- Atomic consistency

in all fair executions (including those in which messages are lost).



Networks – Asynchronous

- We prove this by contradiction
- Assume that the network consists of at least two disjoint nodes $N1$, $N2$
- Assume network meets the three criteria: CAP
- All messages between $N1$ and $N2$ are lost



Networks – Asynchronous




Figure: Best case scenario



Networks – Asynchronous




Figure: Worst case scenario



Networks – Asynchronous




Figure: Case where messages aren't lost



Networks – Partially Synchronous

It is impossible in the partially synchronous network to implement a read/write data object that guarantees the following properties:

- Availability
 - Atomic consistency
- in all fair executions (including those in which messages are lost).



Networks – Partially Synchronous

- We prove this by contradiction
- Assume that the network consists of at least two nodes $G1$, $G2$
- Assume network meets the three criteria: CAP
- All messages between $G1$ and $G2$ are lost



Networks – Partially Synchronous




Figure: Best case scenario



Networks – Partially Synchronous




Figure: Worst case scenario



Partially Synchronous Networks

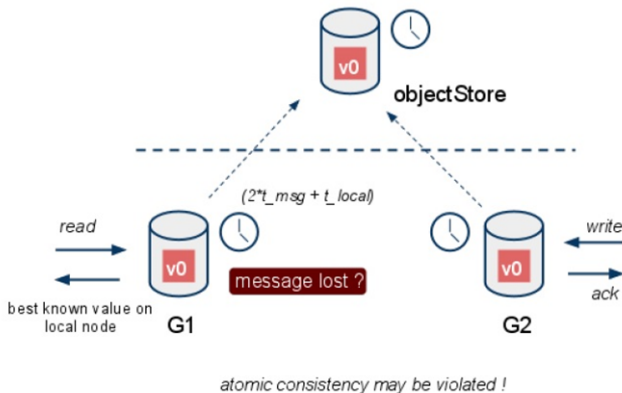


Figure: Centralized Solution



Solution – Take 2



Figure: Pick any 2 and start designing!



Solution – Weaker Consistency Conditions

This guarantee allows for some stale data when messages are lost, but provides a time limit on how long it takes for consistency to return, once the partition heals.



Definition – t-consistent

A timed execution of a read-write object is t-Connected Consistent if two criteria hold. First in executions in which no messages are lost, the execution is atomic. Second, in executions in which messages are lost, there exists a partial order P on the operations.



Centralized algorithm

- What is centralized algorithm?
- Is the modified centralized algorithm is t-Connected consistent?



Solution – Partition Decision

cancel the operation and thus decrease availability,
or
proceed with the operation and thus risk inconsistency



Contents

1 Formal Model

2 Cases

3 Conclusion

4 References



Conclusion

- Proved that CAP is impossible to reliably provide atomic consistent data when there are partitions in the network
- Possible solutions with any two properties of C, A and P



Conclusion



Figure: Multi Player User Game APP



Contents

1 Formal Model

2 Cases

3 Conclusion

4 References



References

- [1] Hagit Attiya, Amotz Bar-Noy, Danny Dolev, Daphne Koller, David Peleg, and Rüdiger Reischuk. *Achievable cases in an asynchronous environment*. 28th Annual Symposium on Foundations of Computer Science, Los Angeles, CA, 1987
- [2] Eric A. Brewer, *Towards robust distributed systems*, Proceedings of the nineteenth annual ACM symposium on Principles of distributed computing, Portland, Oregon, 2000
- [3] Maurice P. Herlihy , Jeannette M. Wing, *Linearizability: a correctness condition for concurrent objects*, ACM Transactions on Programming Languages and Systems, 1990
- [4] Leslie Lamport. *On interprocess communication – parts I and II. Distributed Computing*, 1986.
- [5] Nancy A. Lynch, *Distributed Algorithms*, Morgan Kaufmann Publishers Inc, San Francisco, CA, 1996



Thank you

