

State Machine Replication for the Masses with BFT-SM_AR_T

Hsin-Yang Huang
Chih-shang Chen

Outline

1.Introduction

2.BFT-SMART Design

3.Implementation

4.Alternative Configuration

5.Evaluation

6.Lessons Learned

7.Conclusion

Introduction

Reason:

1. PBFT's architecture does not fully exploit modern hardware
2. UpRight exhibits a performance significantly lower than other systems.

Characteristic:

1. Java-based
2. high-performance and correctness
3. support reconfigurations of the replica

Design principles

- Tunable fault model
 - non-malicious Byzantine-faults
 - malicious Byzantine-faults
 - Simplified SMR protocol
- Simplicity
 - emphasis on protocol correctness
 - avoid optimizations that could bring extra complexity

Design principles

- Modularity
 - uses a well defined consensus primitive in its core
 - easy to implement and reason about

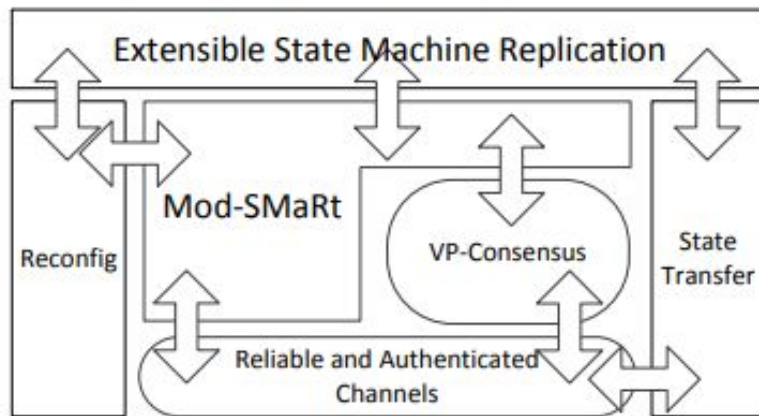


Figure 1. The modularity of BFT-SMART.

Design principles

- Simple and extensible application programming interface
 - Provide simple API such as *invoke(command)* and *execute(command)*
 - Implemented using a set of alternative calls, callbacks or plug-ins (if API not support some methods)
- Multi-core Awareness
 - Take advantage of multi-core architecture of servers

System model

Configuration:

1. $n \geq 3f+1$ to tolerate Byzantine faults
2. $n \geq 2f+1$ to tolerate Crash faults
3. reconfigure replicas at runtime

Links:

1. message authentication code(MAC) over TCP/IP
2. Symmetric keys for replica-replica channel
3. Optional signed request for client-replica channels.

Core protocol

- Total order multicast
 - During normal execution, clients send their requests to all replicas and wait for their replies
 - Total order is achieved through consensus protocol




Figure 2. BFT-SMART normal phase message pattern.

Core protocol (con't)

State Transfer

- to log batches of operations in a single disk
- take snapshots at different points of the execution in different replicas
- perform state transfer in a collaborative way

Core protocol (con't)

Reconfiguration:

- Initiated by *View Manager* client
- Must be signed with a special private key
- *View Manager* sends a special message to the replica that is waiting to be added or removed from the system informing the replica.

Implementation

1. Staged message processing
2. Bounded queue

Netty thread

- Check unordered or ordered request
- Verify client's request

Proposer thread

- Assemble a batch of requests
- Transmitting the PROPOSE message

Sender thread

- Serialize message and produce a MAC
- Send it using TCP sockets

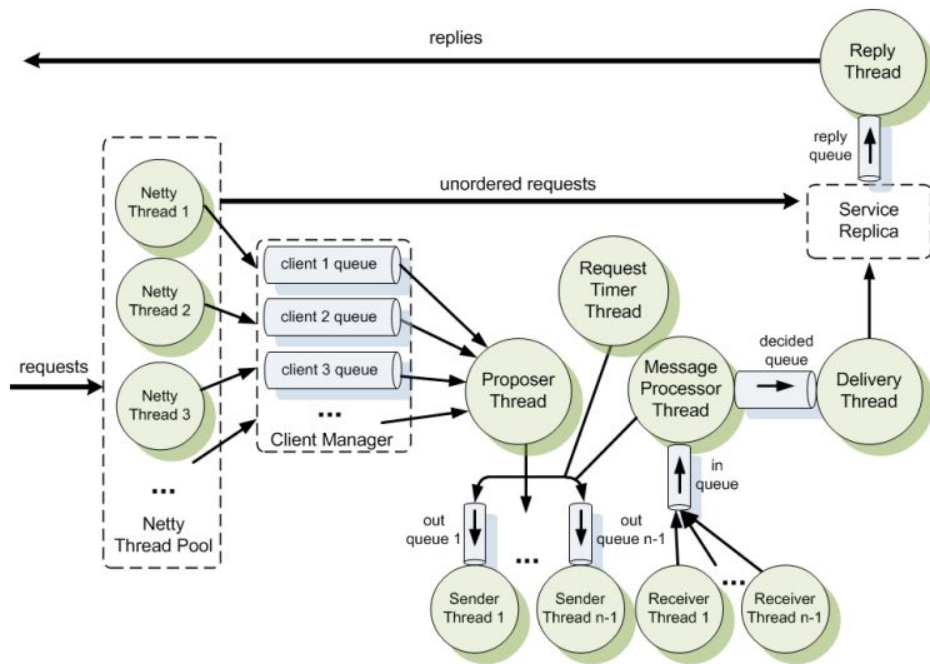


Figure 3. BFT-SMART replica staged message processing.

Implementation

Receiver thread

- Deserialize message
- Put it on the inqueue

Message processor thread

- Fetch messages from the inqueue
- Process message if they belong to current consensus stage
- Put finished decided batch on decide queue

Delivery thread

- Remove request on client queue
- Invoke service replica to generate replies

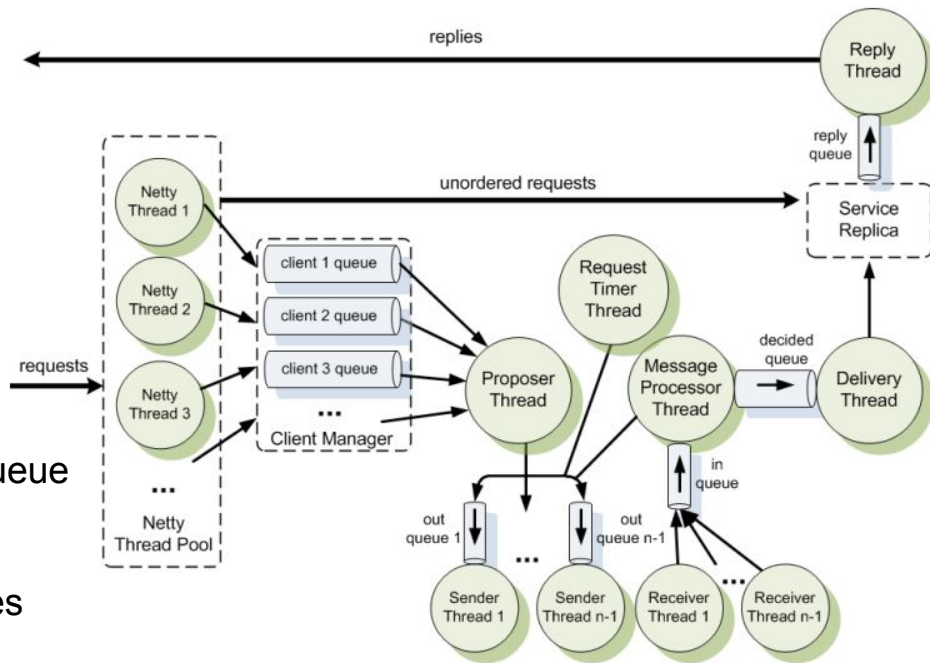


Figure 3. BFT-SMART replica staged message processing.

Implementation

Reply thread

- Fetch request from reply queue
- Send it back to client

Request timer thread

- Activated periodically to verify
If some requests remained more
Than a predefined time.




Figure 3. BFT-SMART replica staged message processing.

Alternative Configurations

1. Crash Fault Tolerance (CFT)

Every node that *do not give a reply* is assumed to be in a crashed state.

Tolerance: $f < n/2$ (simple minority)

Sol => bypass WRITE step

2. Malicious Byzantine Faults


Malicious leader to launch *undetectable attacks*.

Sol => periodic leader changes

Evaluation

1. Raw throughput and Latency
2. Performance in different systems
3. The performance of a BFT-SMART-based system when withstanding faults and reconfiguration.

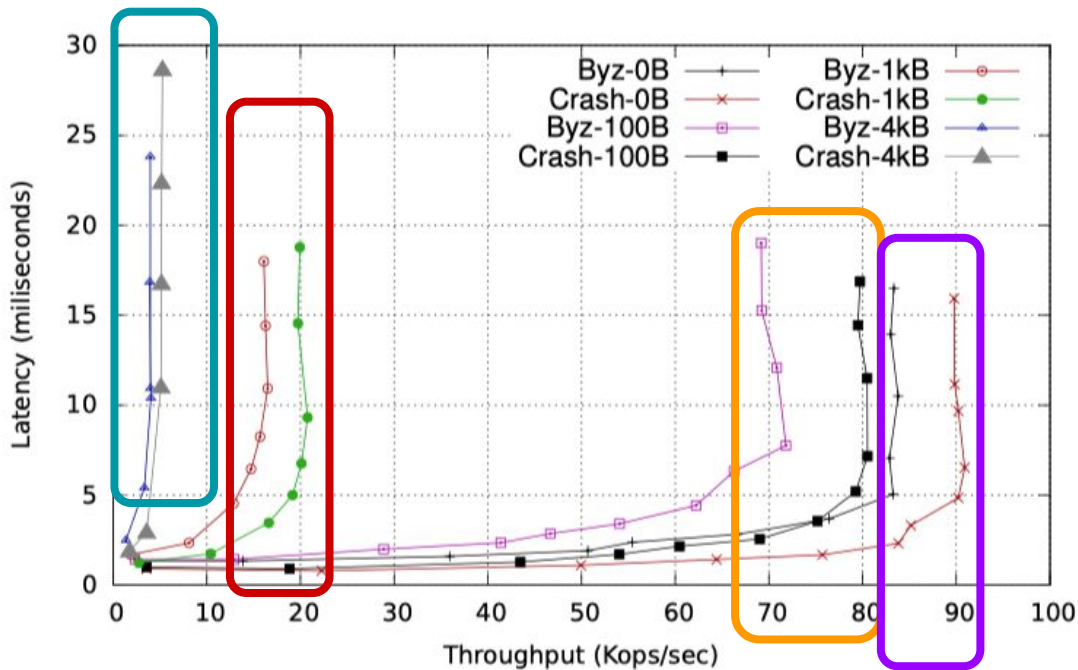
Raw Throughput and Latency



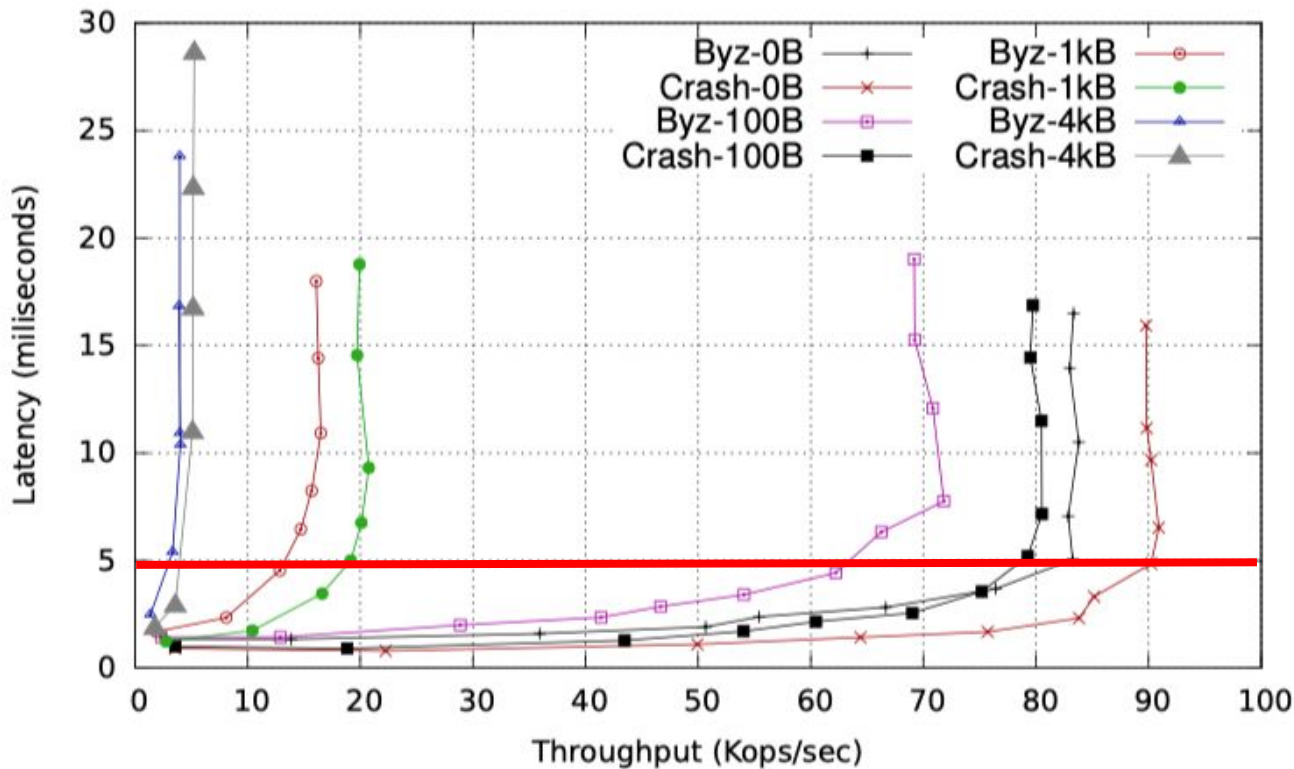
Raw Throughput and Latency

Result 1: CFT setup is always better than BFT


Result 2: Payload size increases -> BFT-SMART performance decreases




Raw Throughput and Latency




Performance in Different System




(a) 0/0



(b) 0/1024




(c) 1024/0




(d) 1024/1024


Performance of BFT-SMART-based System




Performance of BFT-SMART-based System




Performance of BFT-SMART-based System



Performance of BFT-SMART-based System



Performance of BFT-SMART-based System



Lessons Learned

1. BFT in Java
2. How To Test BFT
3. Dealing with Heavy Load
4. Maintenance & Robustness

Lessons Learned

1. BFT in Java

- a. Easy to use
- b. Feasible implementation of secure software

Notice: Need to be used carefully!

2. How To Test BFT

- a. Test on JUnit
- b. Identify the malicious behaviors => carefully analyze
- c. How to inject code for malicious behaviors on replicas => AOP or simple commented code

Lessons Learned

3. Dealing with Heavy Load

- a. Late f replicas in message processing (cuz only needs $n-f$ to progress)
- b. non-Ordered requirements
- c. Thrashing: dropping down throughput under heavy load

4. Maintenance & Robustness

- a. Complex but completed

Core protocol

- Total order multicast
 - During normal execution, clients send their requests to all replicas and wait for their replies
 - Total order is achieved through consensus protocol




Figure 2. BFT-SMART normal phase message pattern.

Lessons Learned

3. Dealing with Heavy Load

- a. Late f replicas in message processing (cuz only needs $n-f$ to progress)
- b. non-Ordered requirements
- c. Thrashing: dropping down throughput under heavy load

4. Maintenance & Robustness

- a. Complex but completed

Conclusions

1. This paper mainly report the process and results in building BFT-SMART library.
2. Describing how to implement the protocol in a safe and efficient way.

Thanks for Listening