

Cerberus

Cerberus: Minimalistic Multi-shard Byzantine
-resilient Transaction Processing. Arxiv'20

By: Jelle Hellings, Daniel P.
Hughes, Joshua Primero, and
Mohammad Sadoghi



Explained By **Prem Rao** and **Greesan Gurumurthy**

Table of Contents

1. What is Cerberus?
2. Heads of Cerberus
3. CCerberus
4. OCerberus
5. PCerberus
6. Strengths & Expected Outcomes of Cerberus

What is Cerberus?

- Cerberus, a set of minimalistic primitives for processing single-shard and multi-shard UTXO-like transactions. Cerberus aims at maximizing parallel processing at shards while minimizing coordination within and between shards.
- **Easy Version:** A sharding protocol bred out of need for fast transaction speed and easy scalability

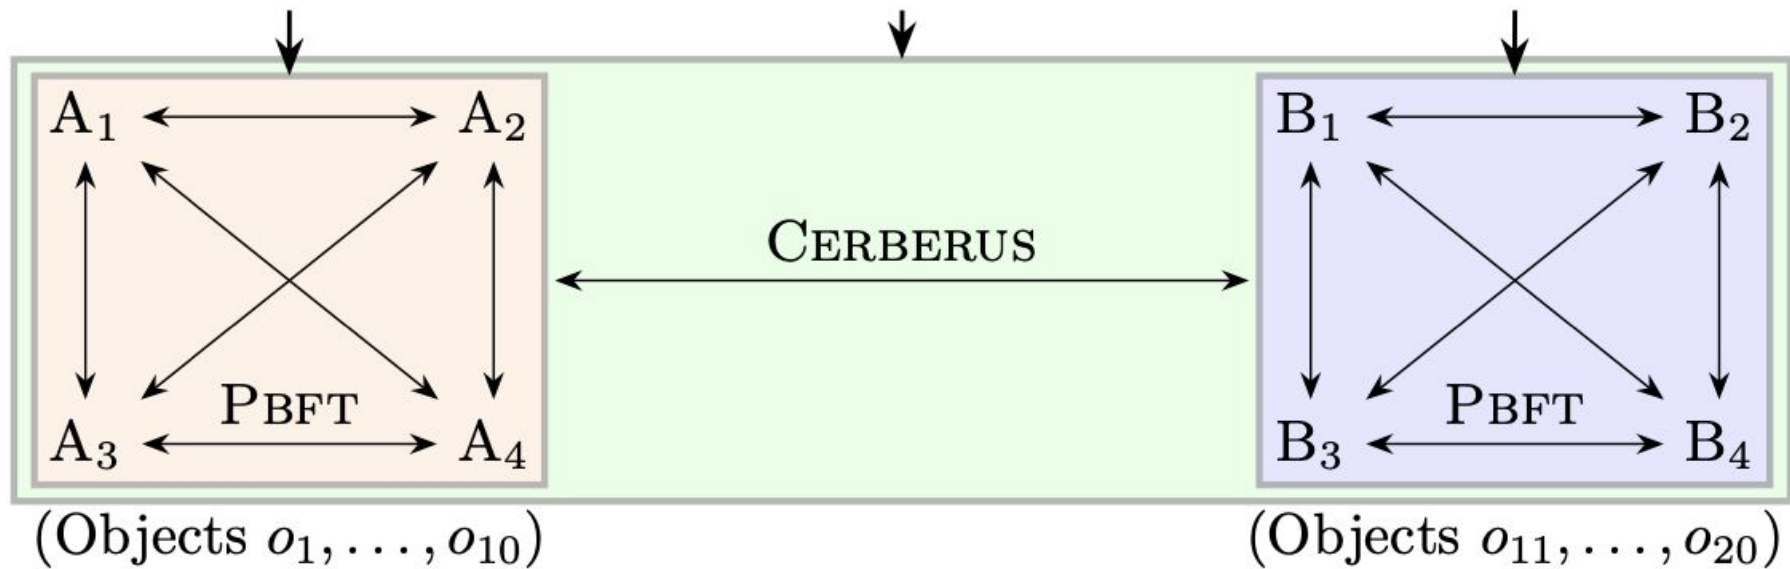
Why Cerberus?

1. Scalability
2. Transaction Speed
3. Comparison to Bitcoin/Mainstream Cryptocurrency
4. Cerberus can reach million transactions per second(STILL THEORETICAL)
5. Permissioned Blockchain vs Public Blockchain

Request on o_3, o_5
(via PBFT)

Request on o_2, o_{14}
(via CERBERUS)

Request on o_{12}, o_{17}
(via PBFT)



Types of Cerberus


1. **Core Cerberus:**uses strict environmental requirements to enable simple yet powerful multi-shard transaction processing
2. **Optimistic Cerberus:**a protocol that does not require any additional coordination phases in the well-behaved optimistic case, while requiring intricate coordination when recovering from attack
3. **Pessimistic Cerberus:**a protocol that adds sufficient coordination to the well-behaved case of Core-Cerberus, allowing it to operate in a general-purpose fault-tolerant environment without significant costs to recover from attack

Shards explained

- Basically a cluster of nodes that hold a unique set of data
- Intra-shard transaction(txn): data impacted or used by txn is located in a single shard, therefore the shard can process the txn and report back
- Cross-shard transaction: data impacted or used by txn located in multiple shards
 - Harder to understand: need to cooperate among shards

Unspent transaction (UTXO) model

Transactions are done by destroying inputs to create specific outputs defined by the transaction message

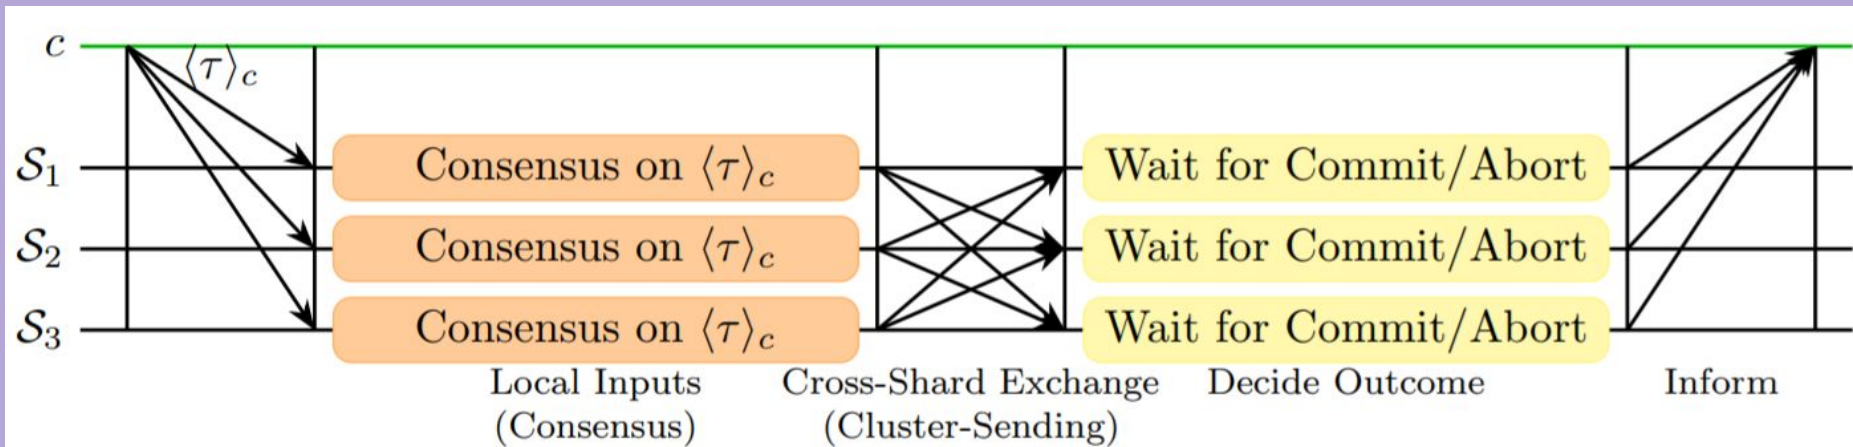


Core Cerberus

Mainly focusing on cross-shard txns

3 main steps:

- Local Inputs: locally decide if shard can contribute inputs to txn
- Cross-Shard Exchange: affected shards exchange/pledge inputs to all other affected shards
- Decide Outcome: if all inputs required by txn have been pledged, shards execute the txn
 - Destroy the inputs to produce the outputs



Message flow of C Cerberus

Notation + Local Inputs:

Txn: $\langle \tau \rangle_c$

Participating Shard: $\mathcal{S} \in \text{shards}(\tau)$

Inputs of S to txn: $I(\mathcal{S}, \tau) = \{o \in \text{Inputs}(\tau) \mid \mathcal{S} = \text{shard}(o)\}$

Currently available inputs of S to txn: $D(\mathcal{S}, \tau)$

If $I(\mathcal{S}, \tau) = D(\mathcal{S}, \tau)$ shard's primary $\mathcal{P}(\mathcal{S})$ pledges inputs to txn, sending

Message: $m(\mathcal{S}, \tau)_\rho = (\langle \tau \rangle_c, I(\mathcal{S}, \tau), D(\mathcal{S}, \tau))$ where ρ refers to the consensus round

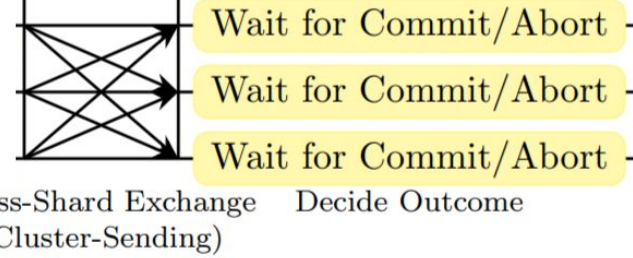
Consensus on $\langle \tau \rangle_c$

Consensus on $\langle \tau \rangle_c$

Consensus on $\langle \tau \rangle_c$

Local Inputs
(Consensus)

X-shard + Deciding Outcome:



Cross-shard exchange: Shards broadcast message $m(\mathcal{S}, \tau)_\rho$ and wait to receive similar messages from all shards involved in txn

Decide Outcome:

If ($I(\mathcal{S}, \tau) = D(\mathcal{S}, \tau) \ \forall \mathcal{S} \in \text{shards}(\tau)$) \rightarrow commit

Else \rightarrow abort

Each replica informs clients of execution outcome, if client receives identical results from $f+1$ replicas from each $\mathcal{S} \in \text{shards}(\tau)$, client can verify that txn has been executed or aborted


Issue: deadlocks

Shards need confirmation from all other involved shards at Xshard exchange

Shards might never reach cross-shard exchange step:

Ex: 2 txns: $Inputs(\tau_1) = Inputs(\tau_2) = \{o_1, o_2\}$ $shard(o_1) = \mathcal{S}_1$ $shard(o_2) = \mathcal{S}_2$

\mathcal{S}_1 Processes τ_1 first, \mathcal{S}_2 processes τ_2 first



Fix:

Internal propagation: shards broadcast msg(txn) to all other involved shards, therefore preventing net issues between client and clusters from stopping communication of the txn

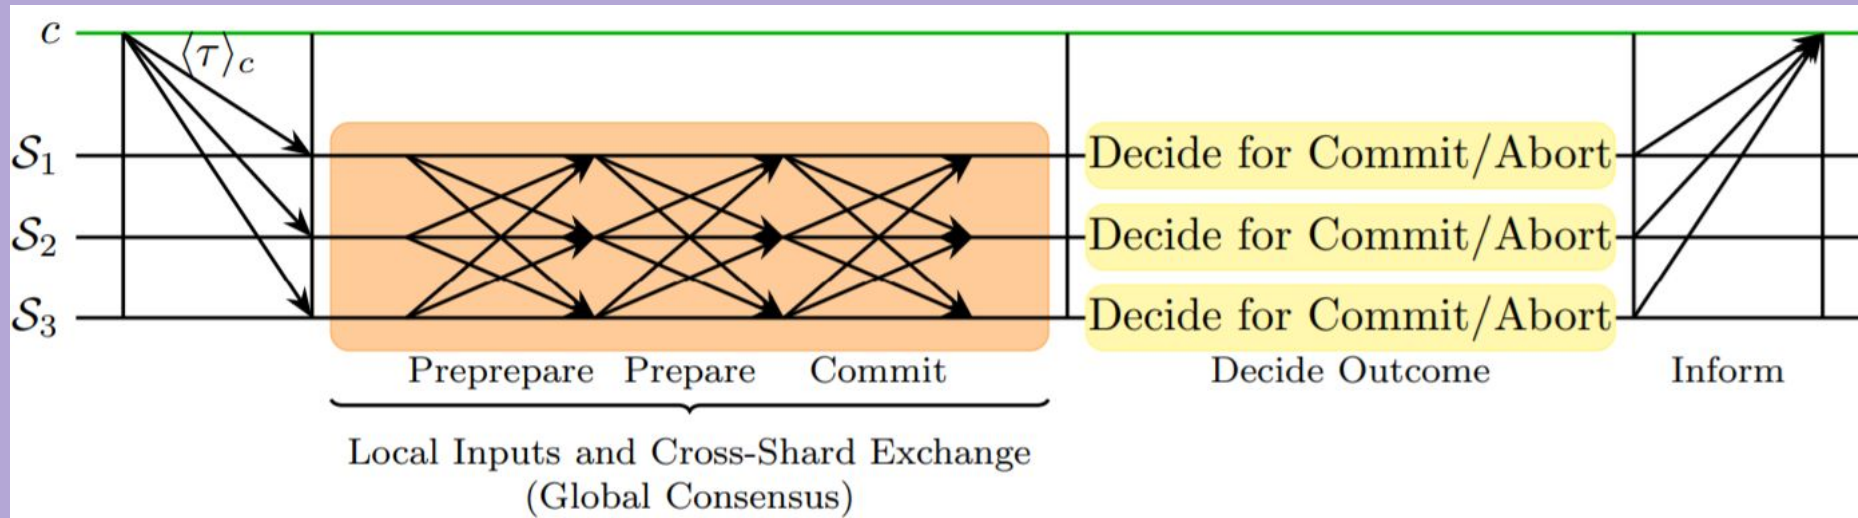
Concurrent resolution: Replicas implement first-pledge and ordered-commit: stops concurrent execution leading to inconsistent state updates

- First Pledge: shard pledges objects made in round p if txn is first transaction proposed after round p requiring those objects
- Ordered-Commit: only commit txn accepted in round p after previous rounds have finished execution



Optimistic Cerberus

Drawbacks with CCerberus that OCerberus attempts to fix:

- Concurrent txns assumed to be result of malicious clients
 - Locks out objects
- Requires consensus than intershard exchange → txn speed decreased
- OCerberus optimized for when no concurrent txns occur
 - Provides recovery for concurrent without locking out objects
 - If malicious entities detected any individual replica can start recovery process
 - Lessens txn processing latency
 - Malicious entity detection and cross-shard coordination spun into one consensus



message flow of OCerberus



Global Preprep:
All primaries of involved
shards send `msg(txn)` to
all replicas of all other
shards

global preprep certificate:
`{msg(txn)}`, ready for next
step

Global Prep:
All replicas of involved
shards send `prep_msg`
to all other replicas in
other involved shards

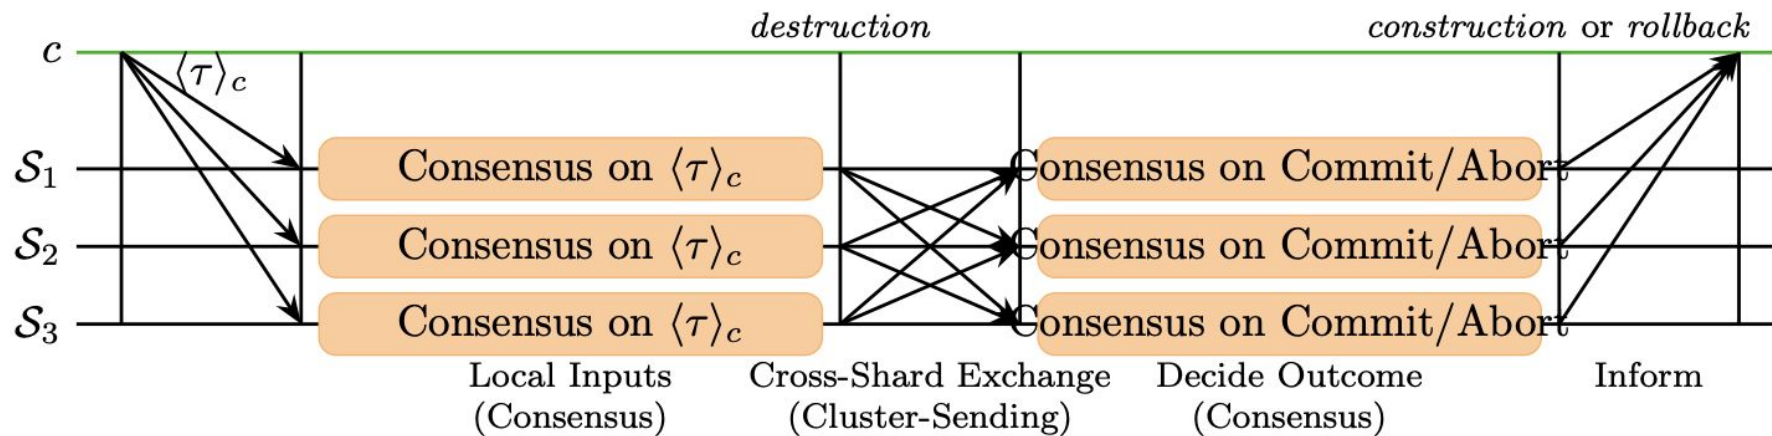
local prep certificate:
`{prep_msg}`
if $\text{length} \geq 2f+1$, ready
for next step

Global Commit:
All replicas of involved
shards send `commit_msg`
to all other replicas in
other involved shards

global commit certificate:
`{commit_msg}`:
if $\text{length} \geq 2f+1$
Ready to decide outcome

Pessimistic Cerberus

- Pessimistic Cerberus is essentially transaction processing under attack.
- We apply a pessimistic approach to CCerberus that is processed to recover from concurrent transactions and made for minimizing the influence of malicious behaviour.
- To better explain PCerberus we can use the following illustration to understand the basic functioning



The message flow of PCerberus for a 3-shard client request that is committed.

The design of PCerberus adds on to the framework of Core Cerberus by adding coordination to the cross-shard exchange and decides outcome steps.

The basic **process** is as follows:

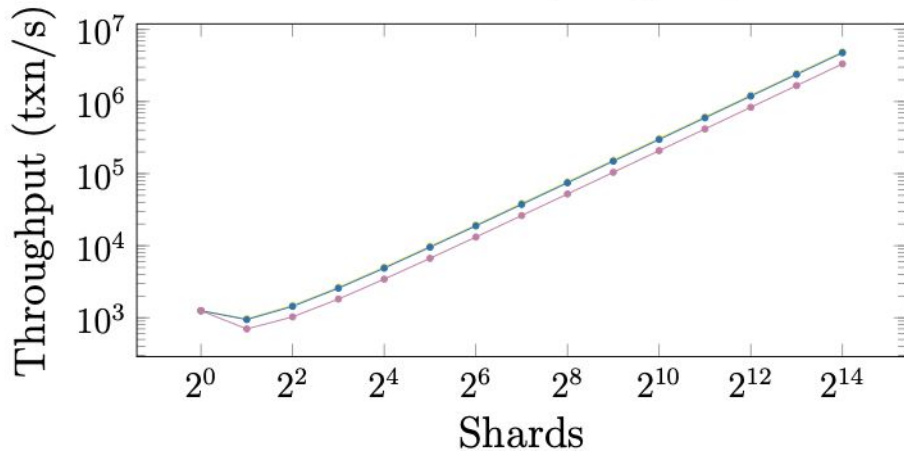
- Acceptance of message by all the replicas -> local inputs step
- Before cross shard exchange, the replicas in the Shard destruct the objects in the readily available inputs(D), thereby fully pledging these objects to τ until the commit or abort decision
- Then, the shard performs cross-shard exchange by broadcasting messages to all the other shards , while the replicas in the shard wait until they receive messages from all other shards
- After cross-shard exchange comes the final decide outcome step.
- We notice that processing a multi-shard transaction via PCerberus requires **two consensus steps** per shard.

Strengths of Cerberus

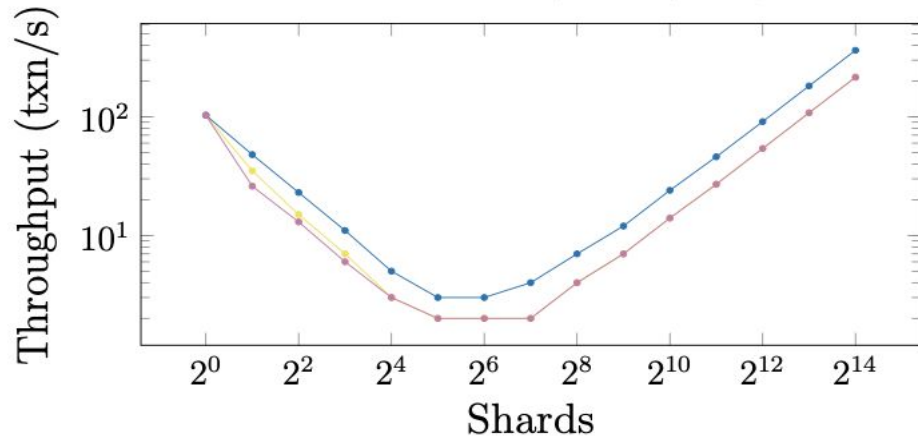
- The main advantages of this protocol through its three mediums are multifaceted:
- Provides Serializable execution
- Maximises per-shard throughput by harnessing out-of-order processing
- Very scalable; one of the reasons for its inception
- Fast transaction speeds, high attainable performance
- Reasonable cost basis; in theory

CCERBERUS OCERBERUS PCERBERUS

Performance (2 obj/txn)

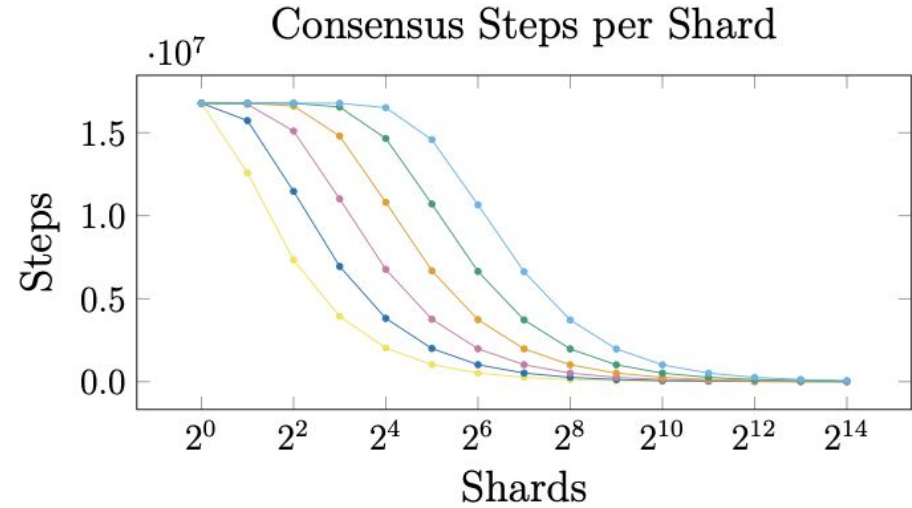
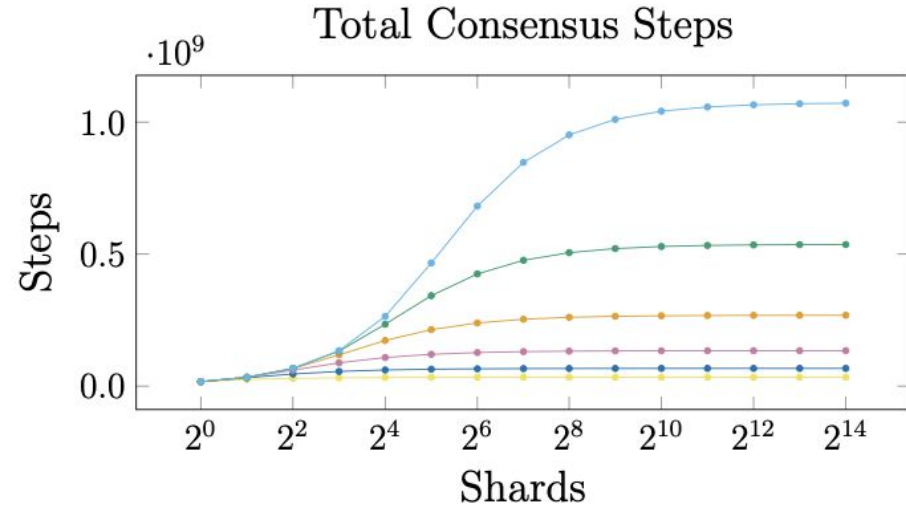


Performance (64 obj/txn)



For every 2^n increase in objects per transaction, the throughput decreases by a factor of 10

—●— 2 obj/txn —●— 4 obj/txn —●— 8 obj/txn —●— 16 obj/txn —●— 32 obj/txn —●— 64 obj/txn



Amount of work, in terms of consensus steps, for the shards involved in processing the transactions

Thank you!