

All about Eve: Execute-Verify Replication for Multi-Core Servers^[1]

OSDI'12

Presenters: Yubo Wang, Chongzhou Fang

Outline

1. Introduction
2. Protocol overview
3. Execution stage
4. Verification stage
5. Experiments
6. Conclusion
7. Reference

Introduction

1. **execute-then-verify** vs. **agree-then-execute**
2. **deterministic execution** vs. **Nondeterministic interleaving of requests**

Protocol Overview

Execution stage:

1. Batching
2. Mixing
3. Executing (in parallel)

Verification stage:

1. Agreement
2. Commit or rollback

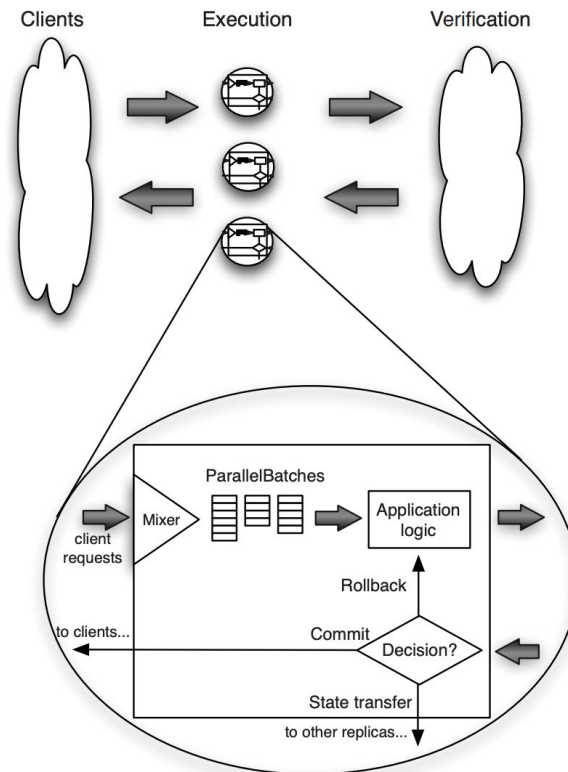



Figure 1: Overview of Eve.

Execution Stage: Mixer Design

1. parallelBatchList
2. readSet
3. writeSet



Execution Stage: Mixer Design

An example for mixer:

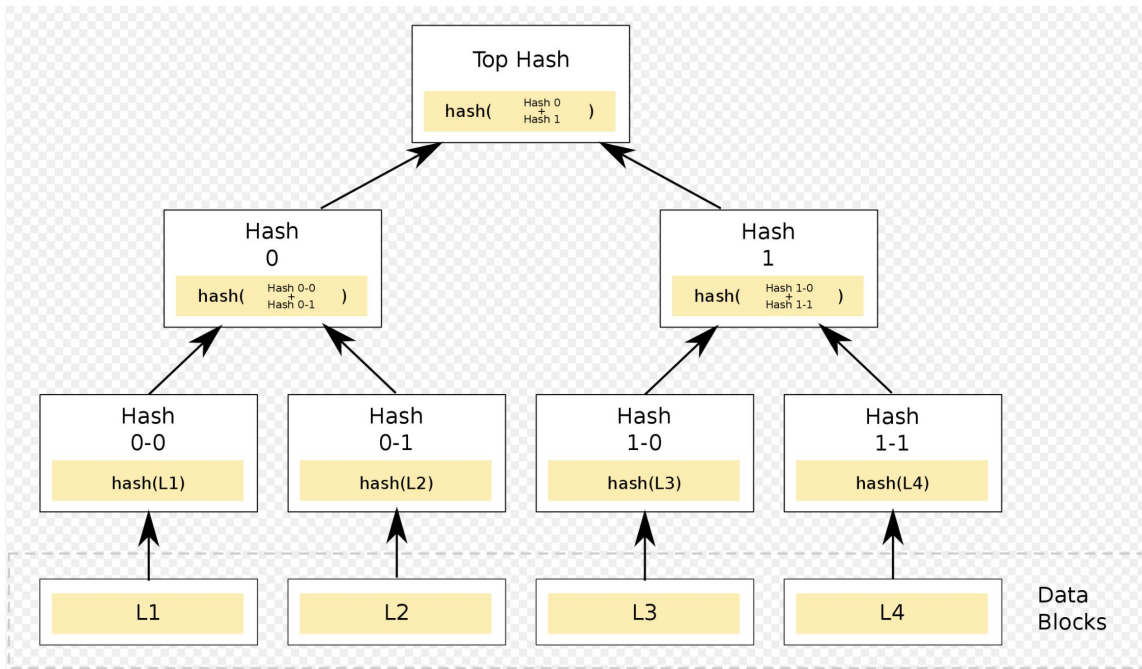
1. parallelBatchList
2. readSet
3. writeSet

Transaction	Read and write keys
getBestSellers	read: item, author, order_line
getRelated	read: item
getMostRecentOrder	read: customer, cc_xacts, address, country, order_line
doCart	read: item write: shopping_cart_line, shopping_cart
doBuyConfirm	read: customer, address write: order_line, item, cc_xacts, shopping_cart_line

Figure 2: The keys used for the 5 most frequent transactions of the TPC-W workload.

Execution Stage: Stage Management

Deterministic Merkle Tree^[2]



Verification Stage

- Goal:
 - Check whether tokens produced by execution replicas match
- Method :
 - Verification Protocol
 - Enough tokens match: Success, commit
 - Not enough: Divergence, roll-back

Verification Stage

- Optimization for Read-Only Requests
 - **First** executed without involving verification stage
 - Enough replies match:
 - Clients receive values
 - Otherwise:
 - Reissued as regular requests


Verification Stage - *Asynchronous BFT*

- Difference between PBFT and Verification Protocol
 1. In PBFT: agree on the output of a single node
 - In Eve: agree on the behavior of execution replicas
 2. In PBFT: agree on the inputs to the state machine
 - In Eve: agree on the outputs of the state machine

Verification Stage - *Asynchronous BFT*

- Verification process

-



Verification Stage - *Asynchronous BFT*

- Upon receiving Verify-Response message
 1. Commit
 - Condition: View number not increased, agreed-upon token matches previously sent one
 - Action: Mark sequence number stable, release requests to clients, etc.

Verification Stage - *Asynchronous BFT*

- Upon receiving Verify-Response message
 1. Commit
 - Condition: View number not increased, agreed-upon token matches previously sent one
 - Action: Make sequence number stable, release requests to clients, etc.
 2. State Transfer
 - Condition: View number not increased, tokens doesn't match
 - Action: Issues a state-transfer request to other replicas

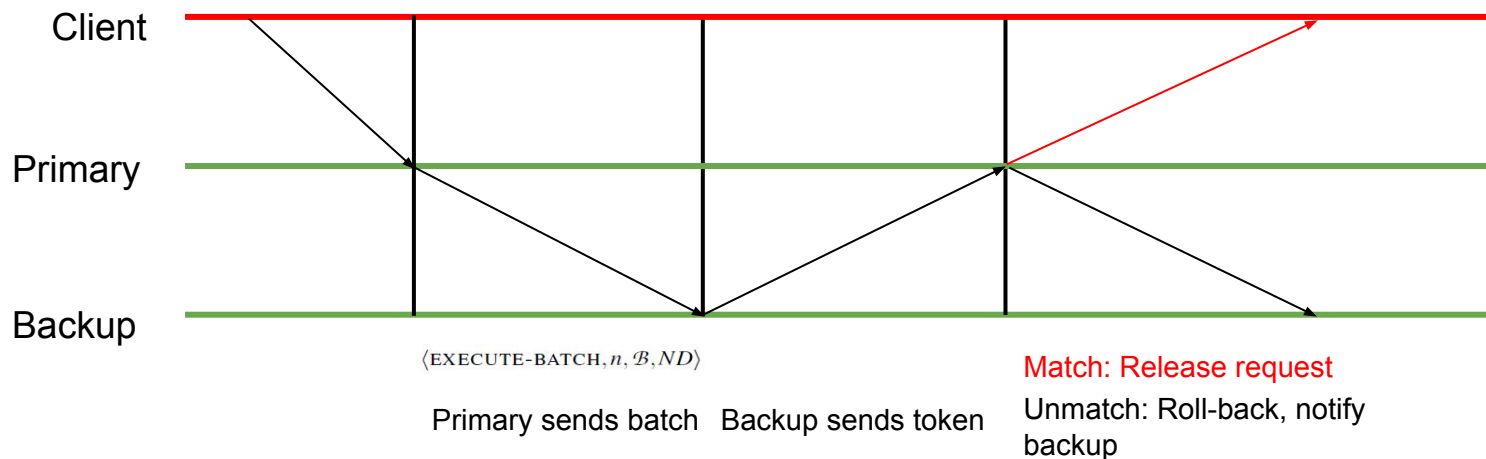
Verification Stage - *Asynchronous BFT*

- Upon receiving Verify-Response message
 1. Commit
 - Condition: View number not increased, agreed-upon token matches previously sent one
 - Action: Make sequence number stable, release requests to clients, etc.
 2. State Transfer
 - Condition: View number not increased, tokens doesn't match
 - Action: Issues a state-transfer request to other replicas
 3. Roll-back
 - Condition: View number increased
 - Action: Roll back state, execute batch sequentially, etc.

Verification Stage - *Synchronous Primary-Backup*

- System Settings

-



Verification Stage - *Tolerating Concurrency Bugs*

- Eve provides protection over concurrency bugs
- Fix concurrency faults: roll-back and sequential execution

Verification Stage - *Tolerating Concurrency Bugs*

- Asynchronous Case
 - When configured with $n_{\text{exec}} = 2u+1$ and $r = 0$, asynchronous Eve is safe, live, and correct despite up to u concurrency or omission faults.

Verification Stage - *Tolerating Concurrency Bugs*

- Asynchronous Case

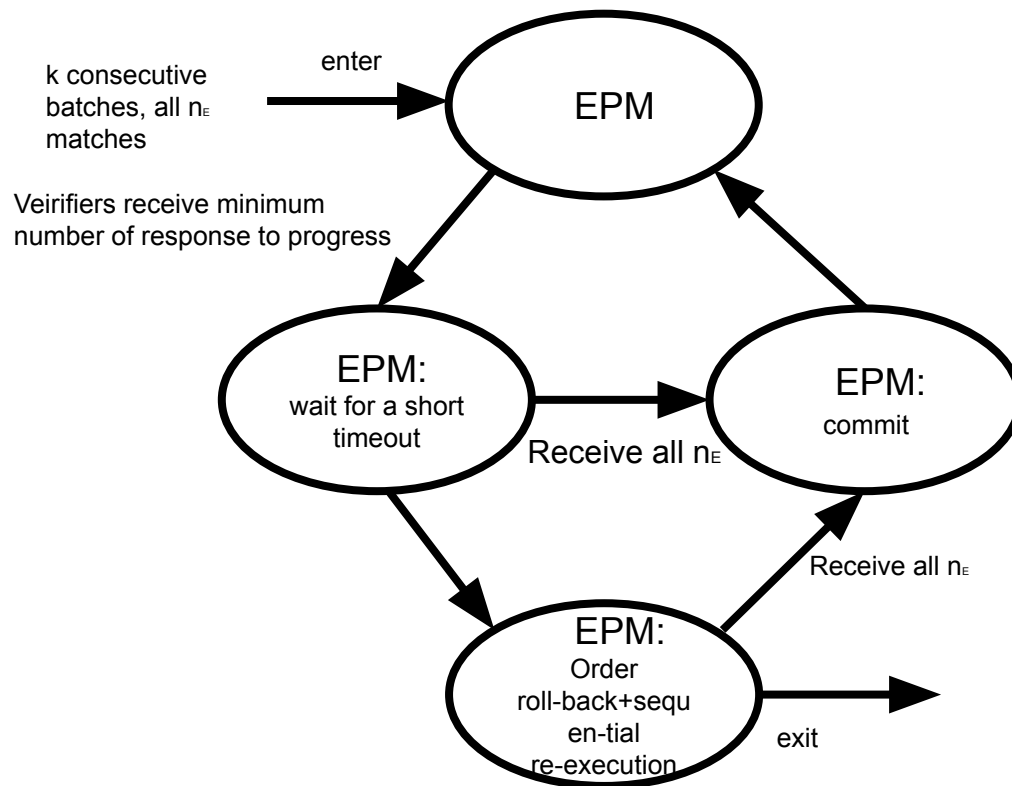
- When configured with $n_{\text{exec}} = 2u+1$ and $r = 0$, asynchronous Eve is safe, live, and correct despite up to u concurrency or omission faults.

- Synchronous Case

- When configured with just $u+1$ execution replicas, Eve can continue to operate with 1 replica if u replicas fail by omission

Verification Stage - *Tolerating Concurrency Bugs*

- Extra protection during good intervals




Evaluation

1. Throughput gain
2. Influence of mixer
3. Currency bug mask

Key-value store application, H2 Database Engine


Evaluation

- Throughput




Evaluation

- Varying CPU demand



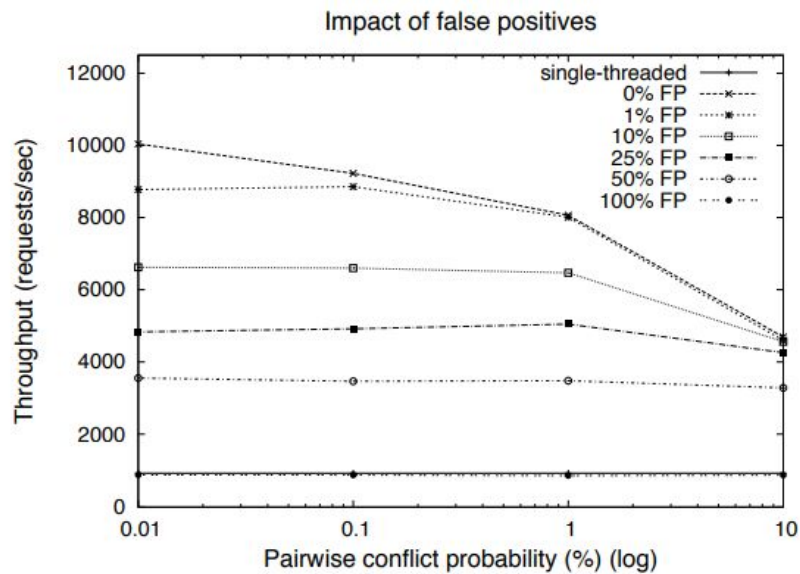
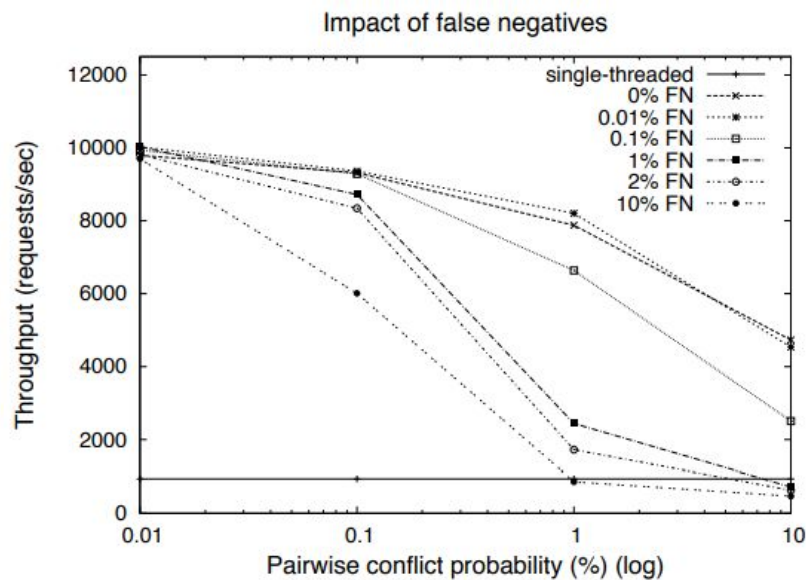
Evaluation

- Varying object size




Evaluation

- Varying conflict probability



Evaluation

- Failure and Recovery





Evaluation

- Concurrency faults

	Group all	1% FN	0.5% FN	0.1% FN	Original Mixer
Times bug manifested	73	51	29	4	0
Fixed with rollback	60	38	18	3	0
All identical (not masked)	13	13	11	1	0
Throughput	1104	1233	1240	1299	1322

Evaluation

- Comparison with Remus



Conclusion

- Eve
 - New Execute-Verify architecture
 - Allow state machine replication to scale to multi-core servers
 - For the first time: allow interleaving requests nondeterministically and execute independently
 - Tolerate omission/commission faults in both asynchronous and synchronous
 - Protects against concurrency bugs

Reference:

- [1] Kapritsos, Manos, et al. "All about Eve: execute-verify replication for multi-core servers." Presented as part of the 10th {USENIX} Symposium on Operating Systems Design and Implementation ({OSDI} 12). 2012.
- [2] Becker, Georg. "Merkle signature schemes, merkle trees and their cryptanalysis." Ruhr-University Bochum, Tech. Rep (2008).