


HotStuff

BFT Consensus in the Lens of Blockchain

The Problem

- View-changes are buggy and time-consuming in PBFT



Step 1: When faulty primary is detected, replica sends VIEW-CHANGE message

Step 2: The next primary in line sends a NEW-VIEW message to everyone

What HotStuff Offers

- Quicker view-changes
 - Achieved linearly, $O(n)$ messages
 - Cost is small enough to where it can change views after every protocol
- Optimistic Responsiveness
 - New leader only needs $n-f$ responses to know progress can be made

Model

Network assumptions


- Synchrony: Known upper bound on the message delays.
- Asynchrony (asynchrony): No known upper bound.
- **Partial synchrony:** The system has an uncertain GST (global stable time) and a Δ , so that the system is in a synchronized state within Δ after the end of GST.

HotStuff works in a Partially Synchronous model!


Transition

- 1-Phase PBFT?
- 2-Phase PBFT
- 2-Phase PBFT Without View Change?
- 2-Phase HotStuff
- 2-Phase HotStuff with Optimistic Responsiveness?
- 3-Phase HotStuff (Basic HotStuff)
- Chained HotStuff


1-Phase PBFT?



1-Phase PBFT – Problem




1-Phase PBFT – Problem




When someone gets stuck, add a phase!

2-Phase PBFT



2-Phase PBFT without View Change?




Definitions

Quorum Certificate

Combines a collection of signatures for the same tuple $\langle type, viewNumber, node \rangle$ signed by $(n - f)$ replicas

Tree and branches


Each command is wrapped in a node that additionally contains a parent link which could be a hash digest of the parent node. Also, in practice, a replica who falls behind can catch up by fetching missing nodes from other replicas



Leader Designation

HotStuff works in a succession of views numbered with monotonically increasing view numbers. Each *viewNumber* has a unique dedicated leader known to all

2-Phase HotStuff



Messages are stored in the form of nodes.

Leader


1. Waits for maximum network delay
2. Proposes new block only if $2f+1$ same lockedQC received.
3. Proposes same prepareQC otherwise.

Replica

Accepts new proposal if **either**:

1. $2f+1$ same lockedQC threshold sign is true, **and**, proposed.node extends from local.lockedQC.node
2. Have the same prepareQC as local prepareQC.

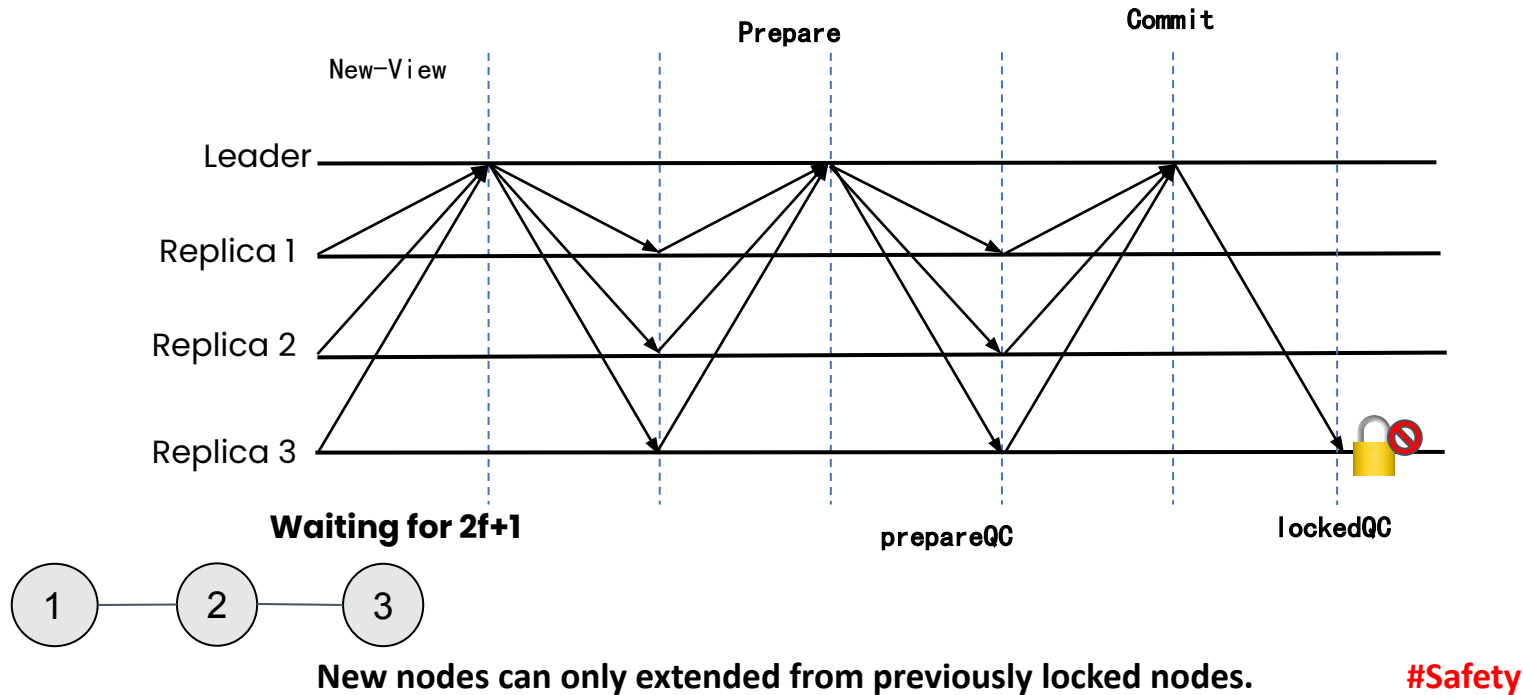
2-Phase HotStuff with optimistic responsiveness



To optimise,

1. The primary only waits for **first 2f+1 messages**.
2. Replica accepts new proposal if **either**:
 - a. *proposed.node* extends from *local.lockedQC.node*
 - b. *proposed.LockedQC.viewNumber* > *local.LockedQC.viewNumber*


Liveness Issue




Liveness Issue



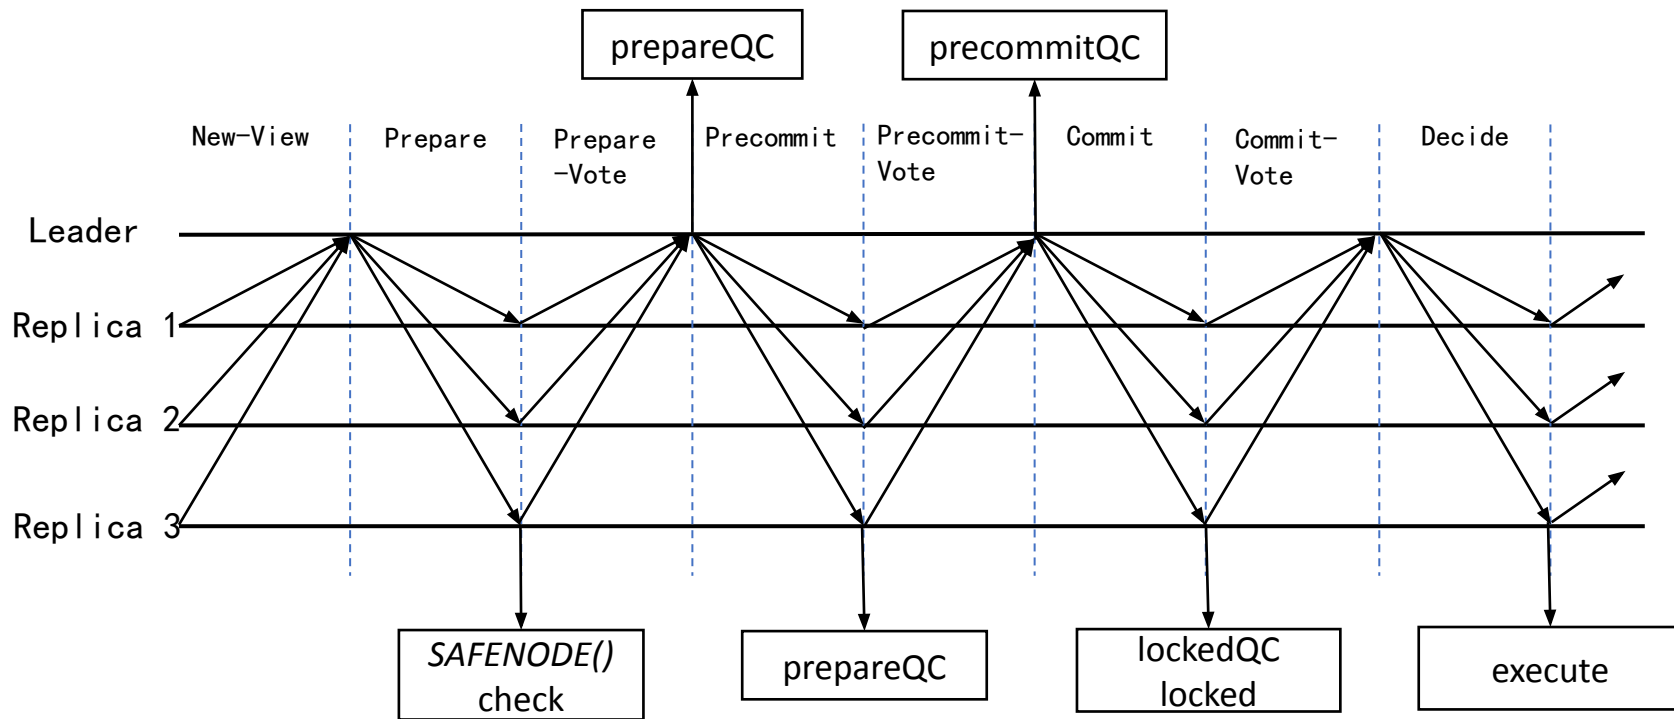
R3 Stays offline for a while, 3f nodes continue.




1. 3f nodes commit on b'.
2. f nodes drop off.
3. 2f locked on b', 1 locked on b.
4. It's a deadlock!



Basic HotStuff 3-Phase Process




New View



Replica:


- NEXTVIEW interrupt (after global timeout in any phase)
- Send $MSG(NEW-VIEW, \perp, prepareQC)$ to $Leader(curView + 1)$

Prepare Phase



Leader:

- Wait for $(n - f)$ *NewView* messages
- Choose *prepareQC* with the highest *viewNumber* as *highQC*
- Create *leaf* on the node with *highQC*
- Broadcast $MSG(PREPARE, curProposal, highQC)$



Replica:

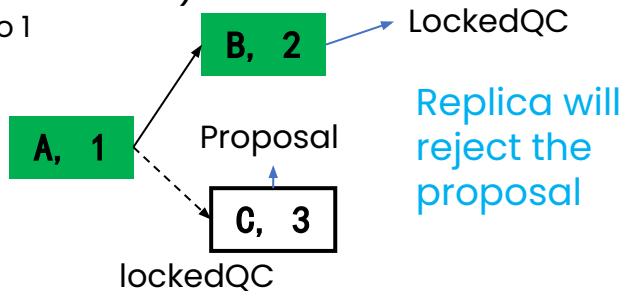
- Wait for $MSG(PREPARE, curProposal, highQC)$ from Leader
- Do *SAFENODE* check
- Send $VOTE_MSG(PREPARE, m.node, \perp)$ to Leader

Prepare Phase – Replica

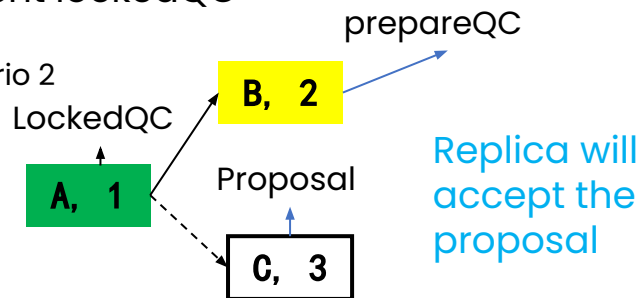
SAFENODE check rules (true if **either** of two rules holds):

- the branch of m.node extends from the currently locked node (Safety rule)
- m.justify has a higher *viewNumber* than the current lockedQC (Liveness rule)


Scenario 1



Scenario 2



Scenario 3




Replica will accept the prepareQC and be unlocked

Why?

Majority accepted the proposal, which means they did not receive the lockedQC and thus there was no commit on view 2

Precommit Phase




Leader:

- Wait for $(n - f)$ *PREPARE* votes
- Combine votes to *prepareQC*
- Broadcast *MSG(PRE-COMMIT, ⊥, prepareQC)*

Replica:

- Wait for *MATCHING_QC(M.JUSTIFY, PREPARE, curView)* from Leader
- Assign *m.justify* to local *prepareQC*
- Send *VOTE_MSG(PRE-COMMIT, m.justify.node, ⊥)* to Leader

Commit Phase




Leader:

- Wait for $(n - f)$ *PRE-COMMIT* votes
- Combine votes to *precommitQC*
- Broadcast *MSG(COMMIT, \perp , precommitQC)*

Replica:

- Wait for *MATCHING_QC(M.JUSTIFY, PRE-COMMIT, curView)* from Leader
- Assign *m.justify* to local *lockedQC*
- Send *VOTE_MSG(COMMIT, m.justify.node, \perp)* to Leader

Decide Phase



Leader:


- Wait for $(n - f)$ *COMMIT* votes
- Combine votes to *commitQC*
- Broadcast *MSG(DECIDE, \perp , commitQC)*

Replica:

- Wait for *MATCHING_QC*(*M.JUSTIFY*, *COMMIT*, *curView*) from Leader
- Execute the command
- Respond to the client


Chained HotStuff

- How we made this flow diagram



Chained HotStuff


- How we made this flow diagram



- All the phases do the the same computation - broadcast messages, replicas partial sign, leader aggregates them. Only the underlying data is different.
- When the state machine is in a phase 'x', all other phases are idle
- Can we take advantage of this, and make all phases do useful work?


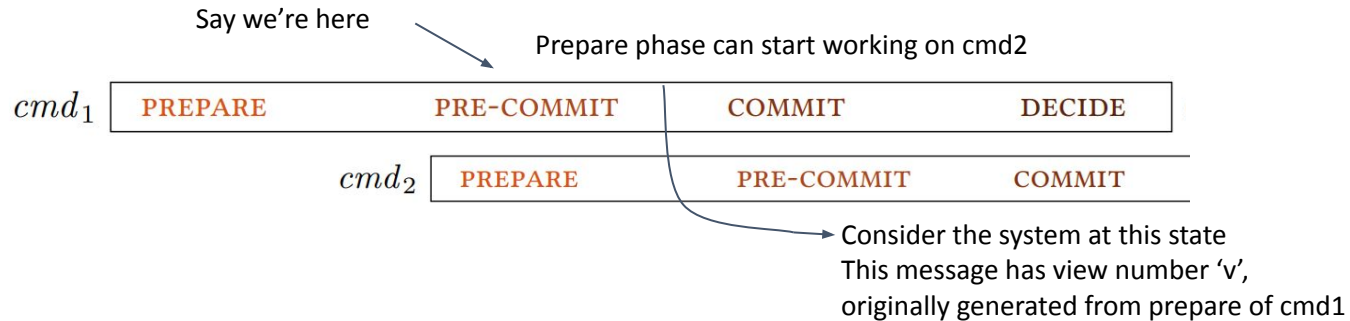
Chained HotStuff

- Generalizing the phases prepare, pre-commit, commit into a “general” message
- How does this effect QC for the phases?



- Read this as “when **cmd1 is in pre-commit**, the system starts processing **prepare for cmd2**”
- Assume cmd_1 starts with view ‘v’, cmd_2 starts with view ‘v+1’

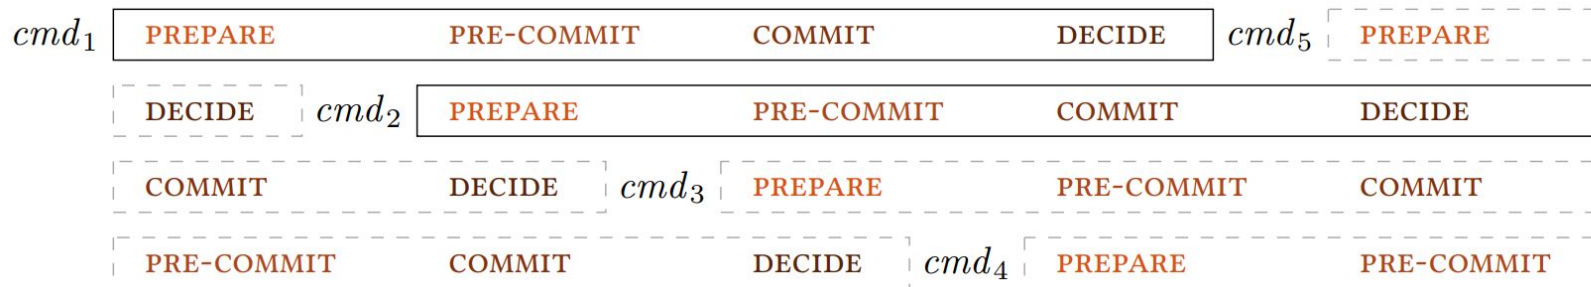
Chained HotStuff



General message with view v in phase $v+1$ implies precommit phase for phase v , prepare phase for $v+1$

Chained HotStuff

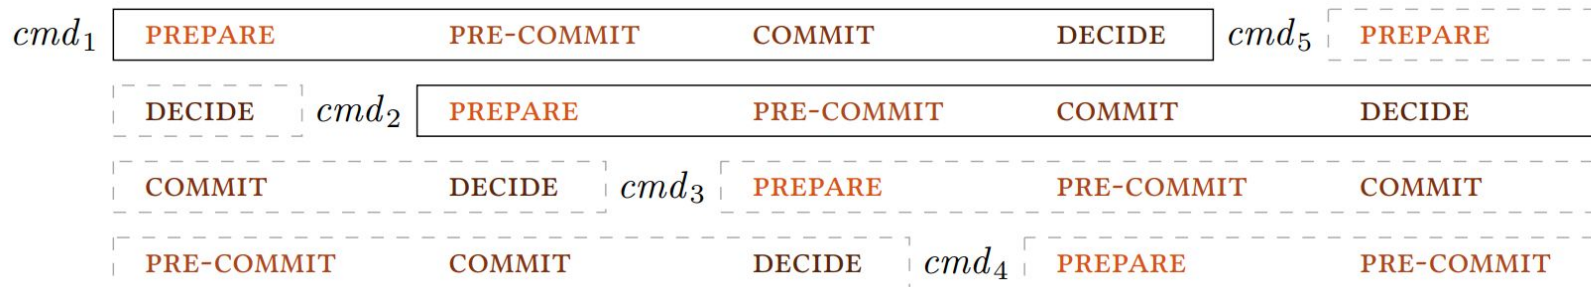
- Other commands are processed simultaneously
- System is 'pipelined'
- Increased throughput, as previously, latency of one command was 3 phases, now it's just one



- Back to the question, how does this effect QC?
- What if QC wasn't reached?
- Can the pipeline stall?
- How is proof (justify) for a command to be committed prepared, when the previous command didn't have QC?

Chained HotStuff

- Other commands are processed simultaneously
- System is 'pipelined'
- Increased throughput, as previously, latency of one command was 3 phases, now it's just one



- Back to the question, how does this effect QC?
- What if QC wasn't reached?
- Can the pipeline stall?
- How is proof (justify) for a command to be committed prepared, when the previous command didn't have QC?

Chained HotStuff

- Each node uses the previous node's QC as justify (proof)

