

RCC: Resilient Concurrent Consensus for High-Throughput Secure Transaction Processing

Presenter: Justin Xu, Shuxian Zhang, Wenchang Liu, Zhixuan Li

Content

- Motivation of RCC
- High level overview and normal case example of RCC
- Dealing with failures
- Client Interactions
- Evaluation of RCC

Limitations of Traditional Consensus

Traditional primary-backup consensus protocols underutilize network resources and thus prevents maximization of transaction throughput.

- Transaction throughput is determined by outgoing bandwidth **B** of primary.

$$T_{max} = \frac{B}{((n - 1)\mathbf{st})}$$

st is the size of each transaction

$$T_{PBFT} = \frac{B}{((n - 1)(\mathbf{st} + 3\mathbf{sm}))}$$

sm is the size of each message

Limitations of Traditional Consensus (cont.)

$$T_{PBFT} \approx T_{max} \text{ when } \mathbf{st} \gg \mathbf{sm}$$

- Replicas are underutilized: primaries must send $(n-1)\mathbf{st}$ while replicas only have to send and receive \mathbf{st} bytes roughly, given that $\mathbf{st} \gg \mathbf{sm}$.
- Underutilization of non-primary replicas in comparison to primaries!

Promise of Concurrent Consensus

Democracy - Give all the replicas the power to be the primary.

Parallelism - Run multiple parallel instances of a BFT protocol.

Decentralization - Always there will be a set of ordered client requests.

Promise of Concurrent Consensus

- HotStuff balances load by consistently switching primaries, but does not address core issue of underutilization of resources.
- Concurrent consensus involves proposing at least nf transactions at a time.

$$T_{cmax} = nf \frac{B}{((n - 1)st + (nf - 1)st)}$$

Towards Resilient Concurrent Consensus

- Concurrent consensus can achieve higher levels of throughput by efficiently utilizing all available replicas.
- RCC is a paradigm for transforming any primary-backup consensus protocol into a concurrent consensus protocol with increased throughput.

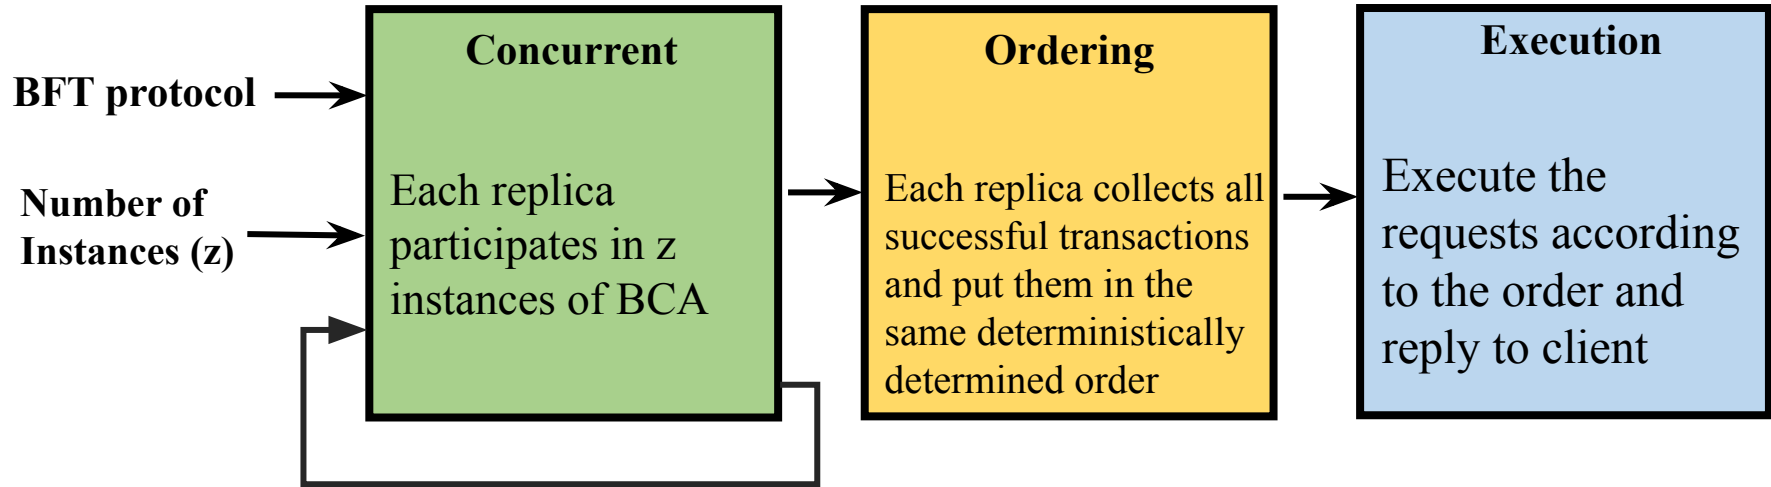
Basic Idea and Design Goals of RCC

- **Idea: Increase throughput by concurrency**
 - Making every replica a primary node
 - Primary nodes propose transactions simultaneously
- **Design Goals:**
 - Provide consensus among replicas on client transactions
 - RCC is a paradigm that can be applied to any other protocols
 - Non-faulty primaries can reach maximum throughput without being affected by faulty behaviors from other replicas
 - Dealing with faulty primaries does not interfere with the operations of other consensus-instances
 - Clients can interact with RCC to force execution of their transactions and learn the outcome of execution

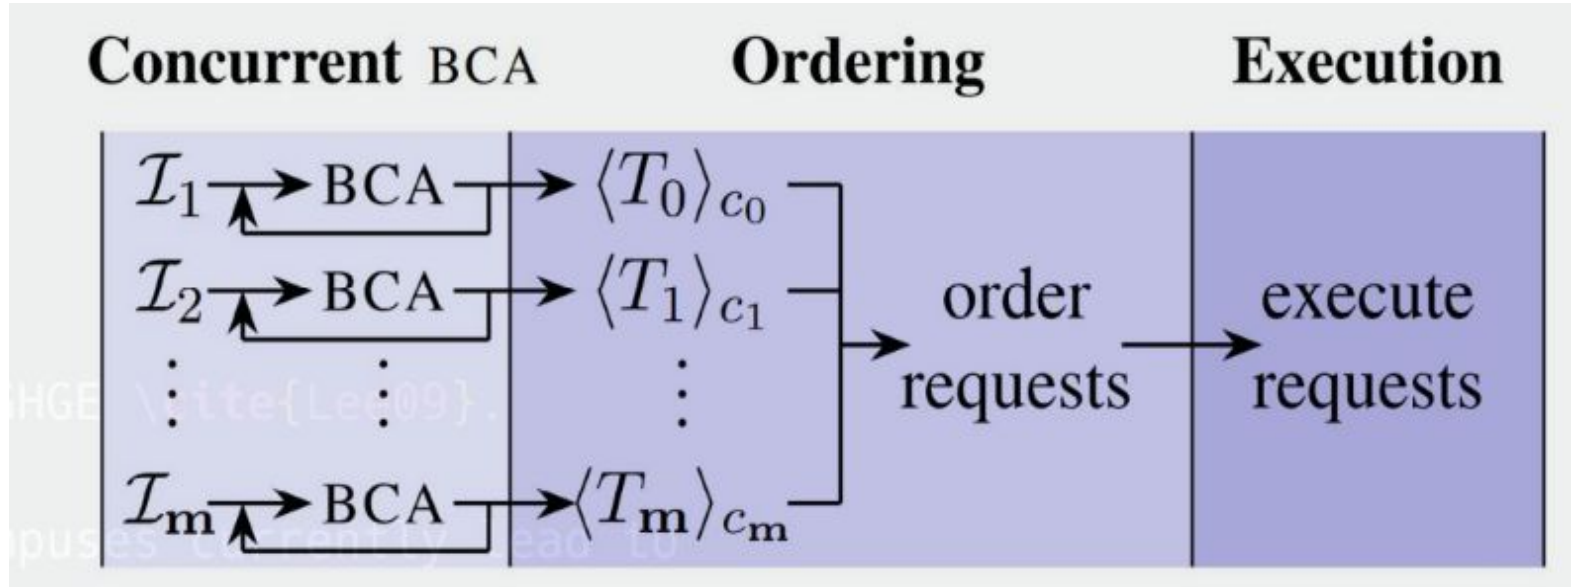
Background / Terminology

- **BCA:**
 - Byzantine commit algorithm(e.g. PBFT, Zyzzzyva, HotStuff etc.)
- **Deterministic:**
 - identical inputs -> identical outputs

RCC Paradigm (Each round)




RCC Paradigm (High Level)




If we have m transactions and m replicas:

- Each replica in each round will have **m** instances participate in **m** BCA.
- Each replica can have one instance to be the primary of a BCA.

Practical Byzantine Fault Tolerance (PBFT)




Normal Case Example: RCC using PBFT with 2 parallel instances




Note: Instance I_{ij} for Replica i participates in Transaction T_j (i.e. $I_{11}, I_{21}, I_{31}, I_{41} \rightarrow T_1$, $I_{12}, I_{22}, I_{32}, I_{42} \rightarrow T_2$)

Dealing with Failures

- **Detectable**
- **Undetectable**



Dealing with Detectable Failures

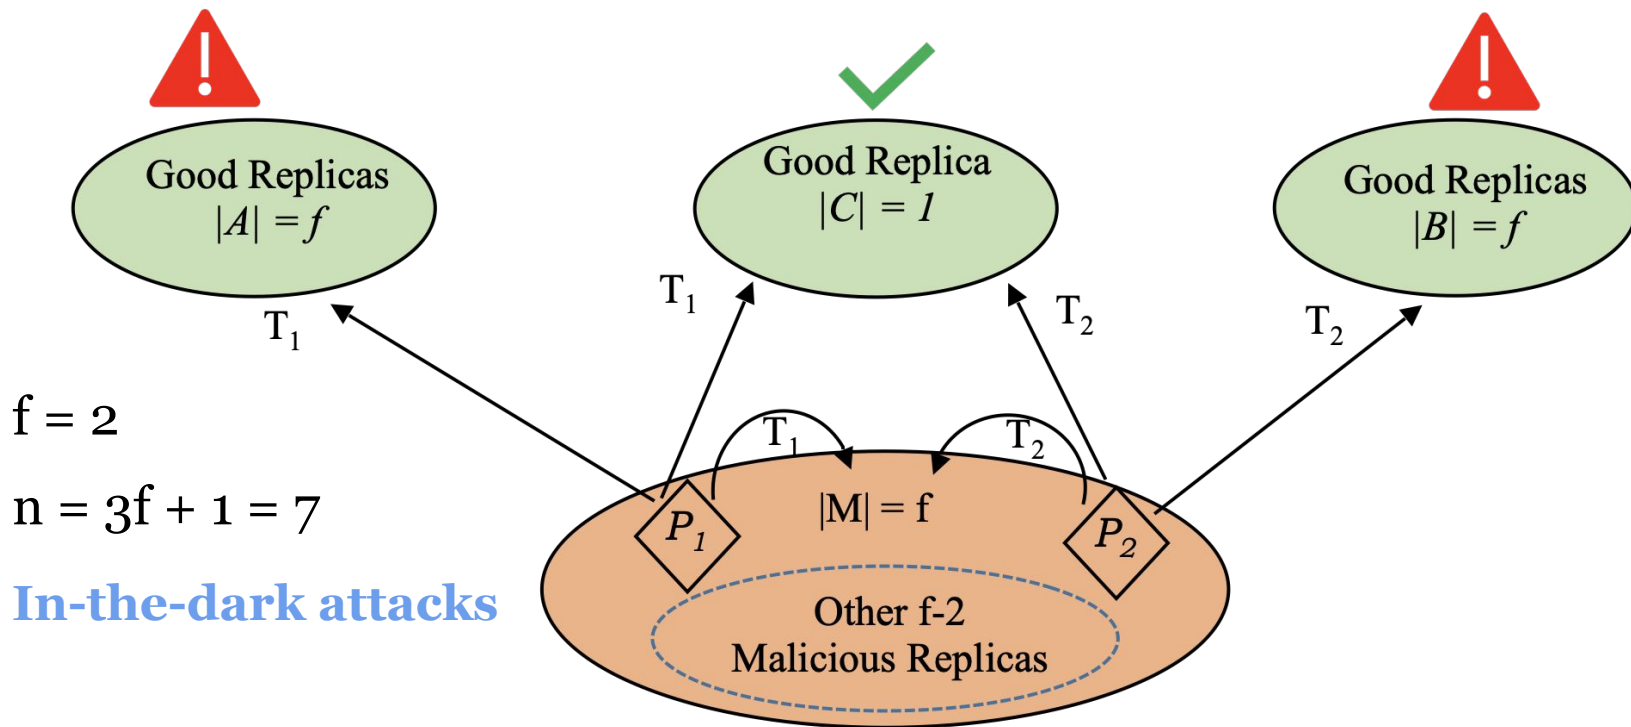


Dealing with Detectable Failures

The **recovery process** in 3 steps:

- All nf replicas need to detect failure
- All nf replicas need to reach agreement on the state of faulty instance
- All nf replicas need to determine which round the faulty primary allowed to resume its operations.


Dealing with Undetectable Failures



Dealing with Undetectable Failures

- A standard checkpoint algorithm
- On a dynamic per-need basis

Dealing with Undetectable Failures



Dealing with Failures


Theorem: *Consider RCC running in a system with n replicas. If $n > 3f$, then RCC provides consensus in periods in which communication is reliable.*

Client Interactions with RCC

- Why is Client Interaction important?
- RCC's optimization
- Problem introduced by optimization
- Solution to the problems


Why is Client Interaction important?

- RCC has multiple clients
- Proposing the same client transaction multiple times reduce throughput



RCC's Optimization

- Assign every client c_i to a single primary P_i such that only the primary P_i can propose c_i 's requests
- Benefit: no need to coordinate clients to assure that they send their transactions to only a single primary




Problems introduced by optimization

- What if primaries do not receive client requests?

Primary proposes a small no-op-request

- What if faulty primaries refuse to propose requests of some clients?




Problems introduced by optimization

- What if faulty primaries refuse to propose requests of some clients?

Incentivize the fault primary to not refuse

Request reassignment



Client Interaction Summary

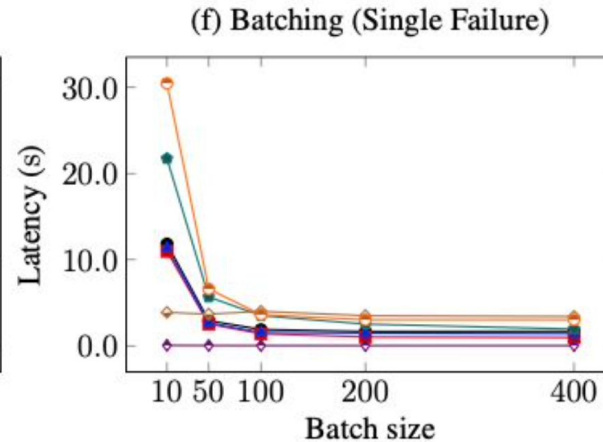

- Each client is assigned to a single primary
- What if primaries do not receive client requests?
 - Primary proposes a small no-op-request instead
- What if faulty primaries refuse to propose requests of some clients?
 - Malicious P: Incentivize malicious primaries to not refuse services
 - Crashed P: reassign primary

Evaluation of RCC

- Varying batch sizes
- Varying numbers of replicas
- Out-of-processing disabled
- Varying BFT protocols

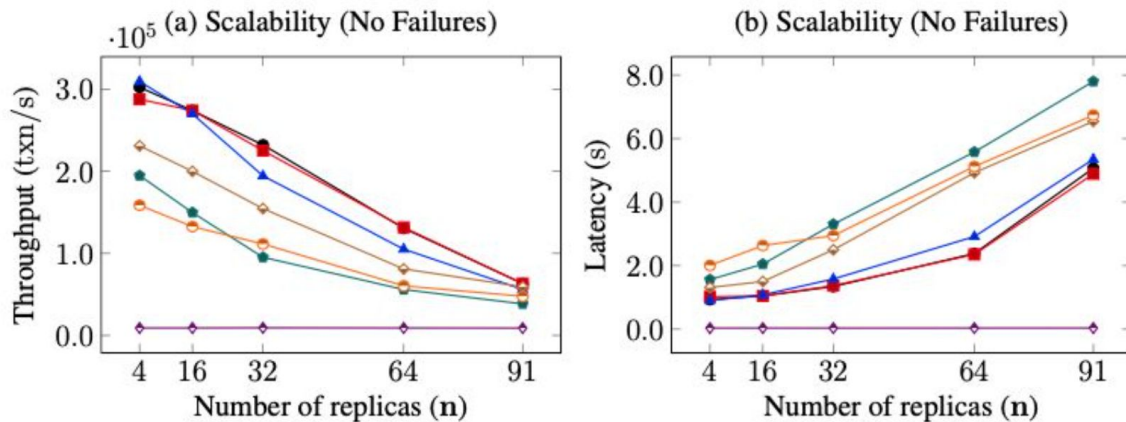
Evaluation - Varying Batch Size

- Performance increases when batch size increases
- Chosen to use 100 txn/batch in all other experiments






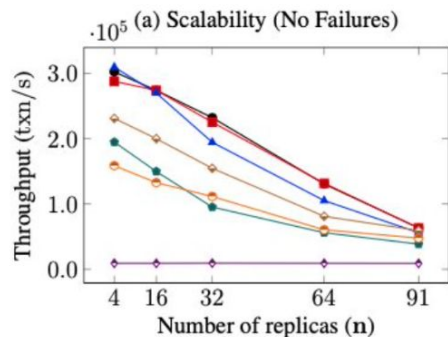
Evaluation - Varying #Replicas (No Failures)

- ZYZZYVA is—indeed—the fastest primary-backup consensus protocol when no failures happens
- RCC easily outperforms ZYZZYVA, even in the best-case scenario of no failures



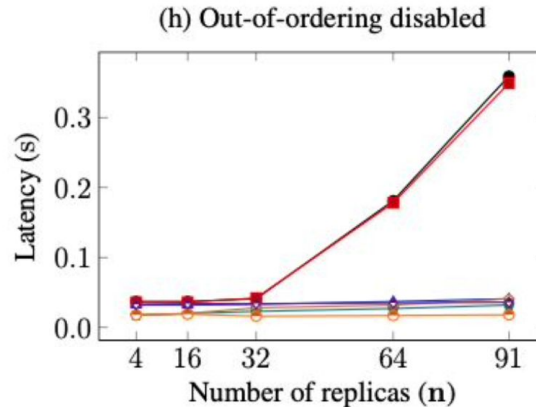
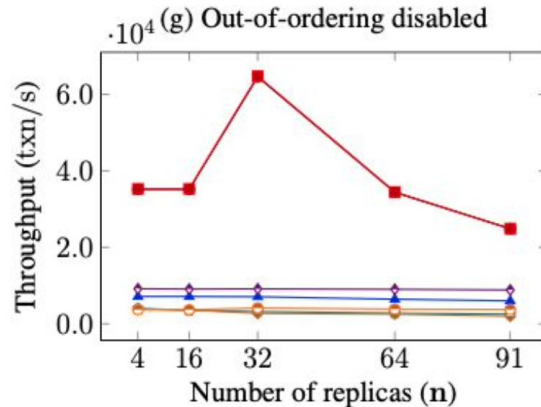
Evaluation - Varying #Replicas

- Three versions of RCC outperform all other protocols
- Performance of RCC with or without failures is comparable
- Adding concurrency by adding more instances improves performance
- On small deployments with $n = 4, \dots, 16$ replicas, the strength of RCC is most evident



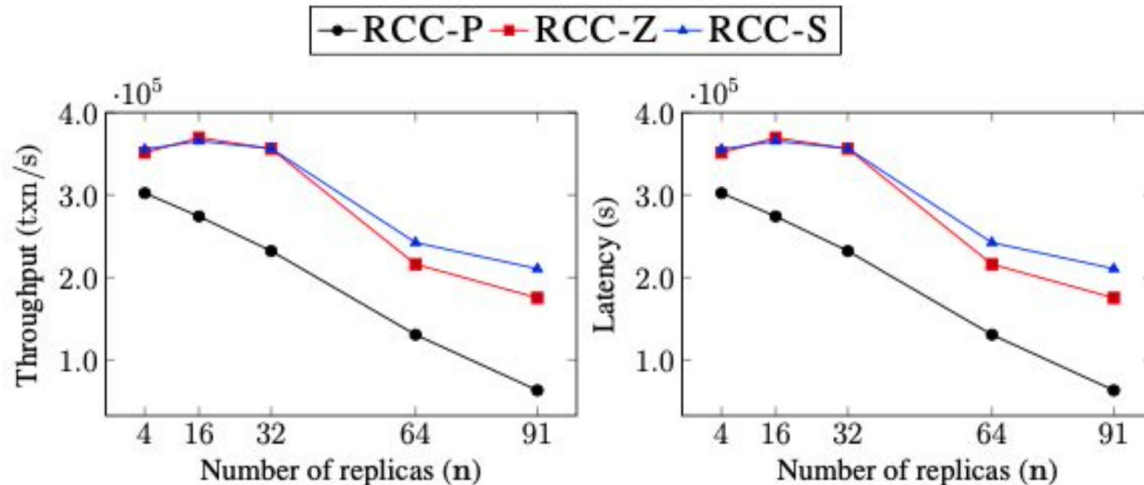
Evaluation - Out-of-ordering disabled

- HOTSTUFF outperforms all other primary-backup consensus protocols
- A non-out-of-order-RCC is still able to greatly outperform HOTSTUFF
- RCC_{f+1} and RCC_n benefit from increasing the number of replicas



Evaluation - Varying BFT Protocol

- RCC-S and RCC-Z achieve higher throughput than RCC-P
- RCC-S consistently attains equal or higher throughput than RCC-Z



Evaluation - Conclusion

- RCC provides more throughput than any primary-backup consensus protocol can provide.
- RCC provides great scalability if throughput is only bounded by the primaries.
- RCC can efficiently deal with failures

Throughput	RCC / SBFT	RCC / PBFT	RCC / HotStuff	RCC / Zyzzyva
Single Failure	2.77x	1.53x	38x	82x
No Failure	2x	1.83x	33x	1.45x

Thank you

Appendix: Ordering (cont.)

Ordering attack example: we have $\text{transfer}(A, B, m)$, A has \$5, B has \$0, C has \$0

- T1 from C1: $\text{transfer}(A, B, \$3)$
- T2 from C2: $\text{transfer}(B, C, \$1)$
- If primary decide T1 \rightarrow T2, then A:\$2, B:\$2, C:\$1
- But if primary decide T2 \rightarrow T1, then A:\$2, B:\$3, C:\$0
- Previously, the only primary can decide the order
- In RCC, all primary nodes decide the deterministic order together, so that no single primary can benefit themselves by deciding the order.

Appendix: Ordering

You can find more details in section: **IV. RCC: IMPROVING RESILIENCE OF CONSENSUS.**

RCC propose a method to deterministically select a different permutation of the order of execution in every round. In such a way that this ordering is partially impossible to predict or influence by faulty replicas.

There is a function that maps an integer h to a unique permutation (if we find a h , we find a permutation, if we find a permutation, we find an order). For sequence S of $k = |S|$ values, there exists $k!$ distinct permutations, so each integer h from $\{0, \dots, k!-1\}$ will correspond to a permutation. Then pick $\mathbf{h} = \mathbf{hash(accepted\ T)mod(k!-1)}$ each round to select that permutation order. Only the concurrent phase ends we know how many are accepted in this round.

Appendix: RCC & GeoBFT

- Similarities:
 - Runs multiple BCAs in parallel
 - Can apply to any other BCAs
 - Aims to increase throughput
 - After BCA, needs to order transactions before execution
- Differences:
 - Different concurrent mechanisms: Physically needs more replicas to run more BCAs for GeoBFT, while RCC logically runs more BCAs with the same amount of replicas
 - Different objectives: GeoBFT aims to resolve latency limitations, RCC aims to make full use of non-primary nodes
 - Different failures dealing mechanisms

References

- [1] S. Gupta, J. Hellings, and M. Sadoghi. RCC: Resilient Concurrent Consensus for High-Throughput Secure Transaction Processing. In *37th IEEE International Conference on Data Engineering*. IEEE, 2021.
- [2] S. Gupta, J. Hellings, and M. Sadoghi. Brief announcement: Revisiting consensus protocols through wait-free parallelization. In *33rd International Symposium on Distributed Computing (DISC)*, 146:44:1–44:3, Schloss Dagstuhl, 2019.
- [3] S. Gupta, J. Hellings, and M. Sadoghi. RCC: Resilient Concurrent Consensus for High-Throughput Secure Transaction <https://www.youtube.com/watch?v=l15M1jyTyvo>