

RingBFT: Resilient Consensus over Sharded Ring Topology

Author: Sajjad Rahnama, Suyash Gupta, Rohan Sogani, Dhruv Krishnan, Mohammad Sadoghi

Presenter: Jiangnan Chen, Fuming Fu, Haochen Yang, Xiaoxi Yu

Date: Oct 25, 2021

Date: Oct 25, 2021

Presenter: Jiangnan Chen, Fuming Fu, Haochen Yang, Xiaoxi Yu

Author: Sajjad Rahnama, Suyash Gupta, Rohan Sogani, Dhruv Krishnan, Mohammad Sadoghi

Roadmap

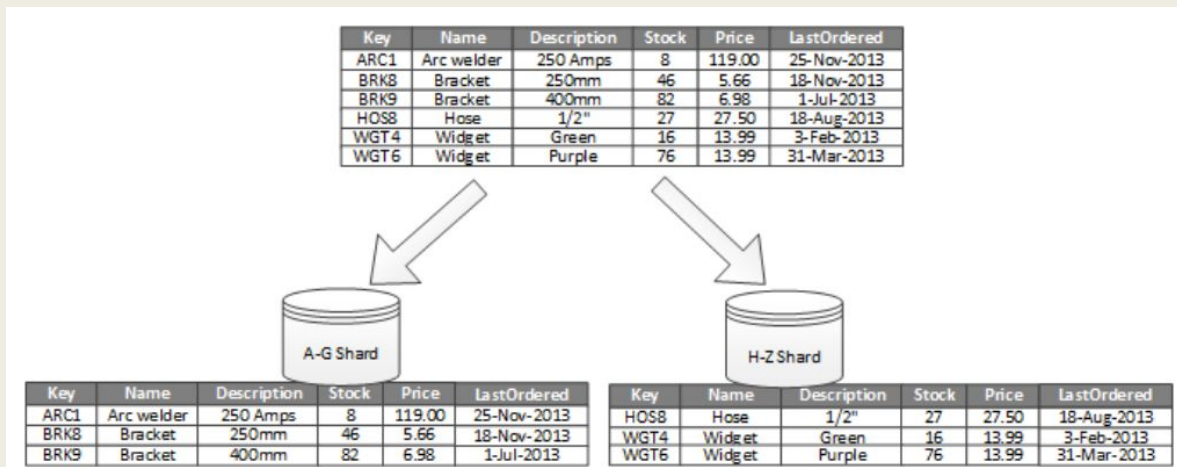
1. Introduction (Yu)
 - a. Something about Sharding in common database system
 - b. Background of cross-shard transactions and why we need ringBFT
2. Single-Shard protocol (Yu)
3. Cross-Shard protocol (Chen, Fu)
4. Uncivil Executions and Attacks (Yang)

Partitioning in Database Systems

- Why data partitioning? Improve scalability.
- Strategies of data partitioning:
 - Horizontal Partitioning - separate by entities
 - Vertical Partitioning - separate by features

Partitioning in Database Systems

- Why data partitioning? Improve scalability.
- Strategies of data partitioning:
Horizontal Partitioning - separate by entities

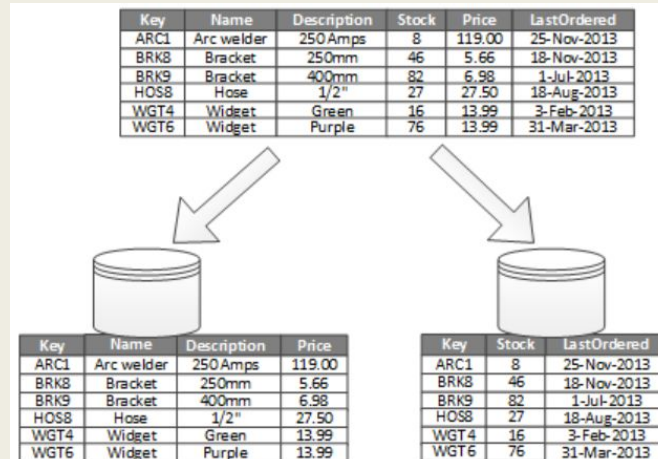


Partitioning in Database Systems

- Why data partitioning? Improve scalability.
- Strategies of data partitioning:

Horizontal Partitioning - separate by entities

Vertical Partitioning - separate by features



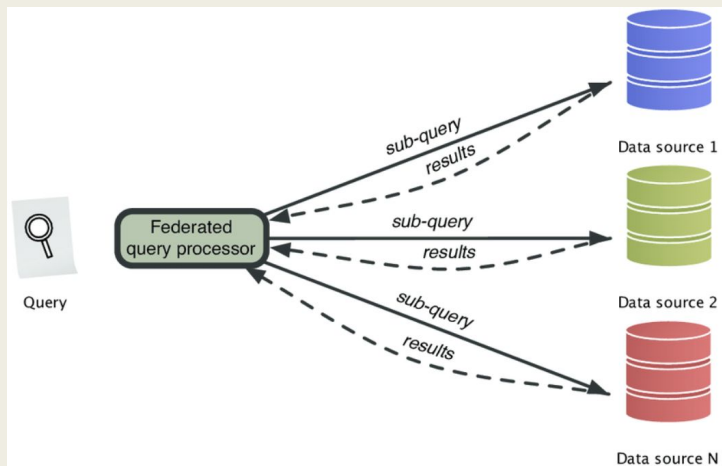
<https://docs.microsoft.com/en-us/azure/architecture/best-practices/data-partitioning>

Partitioning in Database Systems

- Why data partitioning? Improve scalability.
- Strategies of data partitioning:
 - Horizontal Partitioning - separate by entities
 - Vertical Partitioning - separate by features


Shard is usually referred to a **horizontal partition** of data in a database. Each shard is held on a separate database server instance.

Federated Data System



- Correlate data from local tables and remote data sources, as if all the data is stored locally in a single database
- Update data in relational data sources, as if the data is stored in a single database
- Move data to and from relational data sources
- Take advantage of data source's processing strengths, by sending requests to particular data sources for processing
- Compensate for SQL limitations at a data source by having the federated server process parts of a distributed request

Cross-Shard Transactions

- In federated database system, the cross-shard transactions are common.
 - Cross-shard transactions require not only communication between shards, but also their fate depends on the consent of each of the involved shards.
 - Traditional BFT not work.
-
- Designated Committee (AHL)
 - Initiator Shard (Sharper)
- 
- Require all-to-all communication between replicas of each shard

$O(n^2)$

Why RingBFT ?

1. linear communication between replicas
2. secure against byzantine attacks
3. high throughputs
4. low latencies
5. scalable for cross-shard transactions
6. inexpensive when transactions require access to multiple shards

Safety and Liveness

- **Involvement**

Each $S \in \mathfrak{S}$ processes a transaction if $S \in \mathfrak{I}$.

- **Termination**

Each non-faulty replica in \mathfrak{R}_s executes a transaction.

- **Non-divergence**

All non-faulty replicas in \mathfrak{R}_s execute the same transaction.

- **Consistence**

Each non-faulty replica in \mathfrak{S} executes a conflicting transaction in same order.

Traditional Replicated System

- **Involvement**

Each $S \in \mathcal{S}$ processes a transaction if $S \in \mathcal{I}$.

- **Termination** \longrightarrow **liveness**

Each non-faulty replica in \mathcal{R}_s executes a transaction.

- **Non-divergence** \longrightarrow **safety**

All non-faulty replicas in \mathcal{R}_s execute the same transaction.

- **Consistence**

Each non-faulty replica in \mathcal{S} executes a conflicting transaction in same order.

Single-shard Transactions

- **Involvement**

Each $S \in \mathfrak{S}$ processes a transaction if $S \in \mathfrak{I}$.

- **Termination** \longrightarrow **liveness**

Each non-faulty replica in \mathfrak{R}_s executes a transaction.

- **Non-divergence** \longrightarrow **safety**

All non-faulty replicas in \mathfrak{R}_s execute the same transaction.

- **Consistence**

Each non-faulty replica in \mathfrak{S} executes a conflicting transaction in same order.

Cross-shard Transactions

- **Involvement**

Each $S \in \mathcal{S}$ processes a transaction if $S \in \mathcal{I}$.

- **Termination** \longrightarrow **liveness**

Each non-faulty replica in \mathcal{R}_s executes a transaction.

- **Non-divergence** \longrightarrow **safety**


All non-faulty replicas in \mathcal{R}_s execute the same transaction.

- **Consistency**


Each non-faulty replica in \mathcal{S} executes a conflicting transaction in same order.

Single-Shard Protocol


- No communication among the shards.
- Each transaction only need to requires access to data in one shard.



PBFT-request




PBFT-preprepare

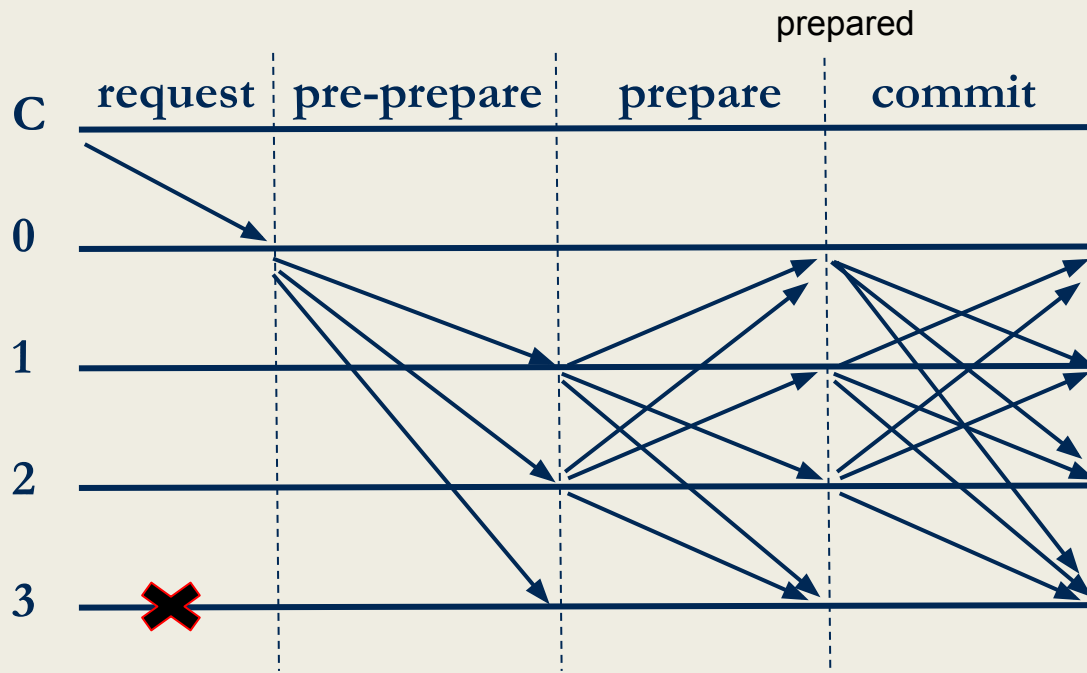


Preprepare message:
(1) sequence number
(2) digest of the client transaction

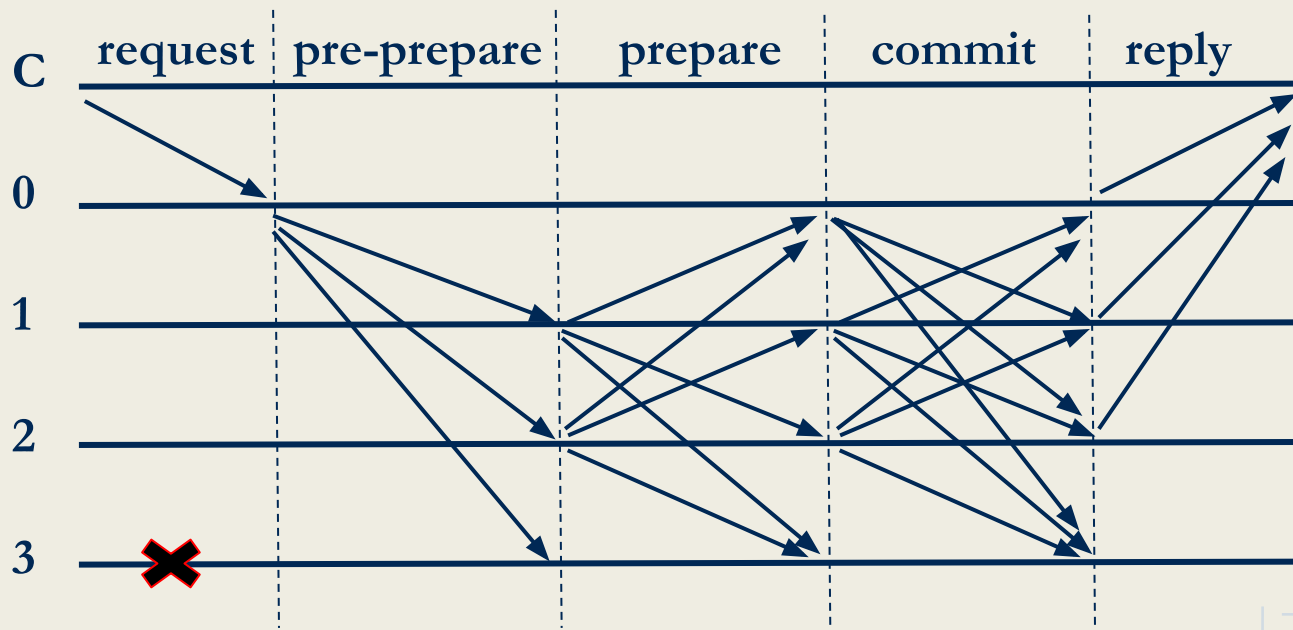
PBFT-prepare



PBFT-commit



PBFT-reply



Cross-Shard Consensus Algorithm

Problem

Cross-shard transactions require not only communication between shards, but also their fate depends on the consent of each of the involved shards.

Prior work

- Designated Committee (AHL)
- Initiator Shard (Sharper)



Require all-to-all communication between replicas of each shard



Improvements

- Linear Communication between replicas
- Lock
- Ring order

Cross-Shard Consensus Algorithm

Network Situations :

- Normal-case (stable / no faulty replicas)
- Uncivil executions



Cross-shard Cases

- Simple Case
- Complex Case




Clarification:

- Simple Case: Each shard can independently run consensus and execute its fragment
- Complex Case : An involved shard may require access to data from other involved shards to execute its fragment.


Cross-Shard Consensus Algorithm

Cross-shard Cases

- Simple Case
- Complex Case



Cross-Shard Consensus Algorithm



Ring order: S-U-V-W

Problem:


Both client C1 and C2 will access shard U and shard V.

Scenarios:(eg. Shard U)

1. Each client requires different data from shard U.

Result: No conflict, each runs independently.


Cross-Shard Consensus Algorithm



Scenarios (eg Shard U):
2. What if two clients require the same data from shard U?



- The first access to data-item a in shard U will lock data-item a so that others can not access the same data.



Cross-Shard Consensus Algorithm


How does the protocol run in normal cases?

What does Linear Communication Primitive mean?

Why shards need ring order to deal with deadlock?





Assumption


1. A system of four shards: S, U, V, and W, the ring order is $S \rightarrow U \rightarrow V \rightarrow W$.
2. The number of replicas in each shard $\geq 3f+1$.
3. Byzantine replicas are unable to impersonate non-faulty replicas.



Client Request

- Specify
- Wait







Client Request Reception

Primary:

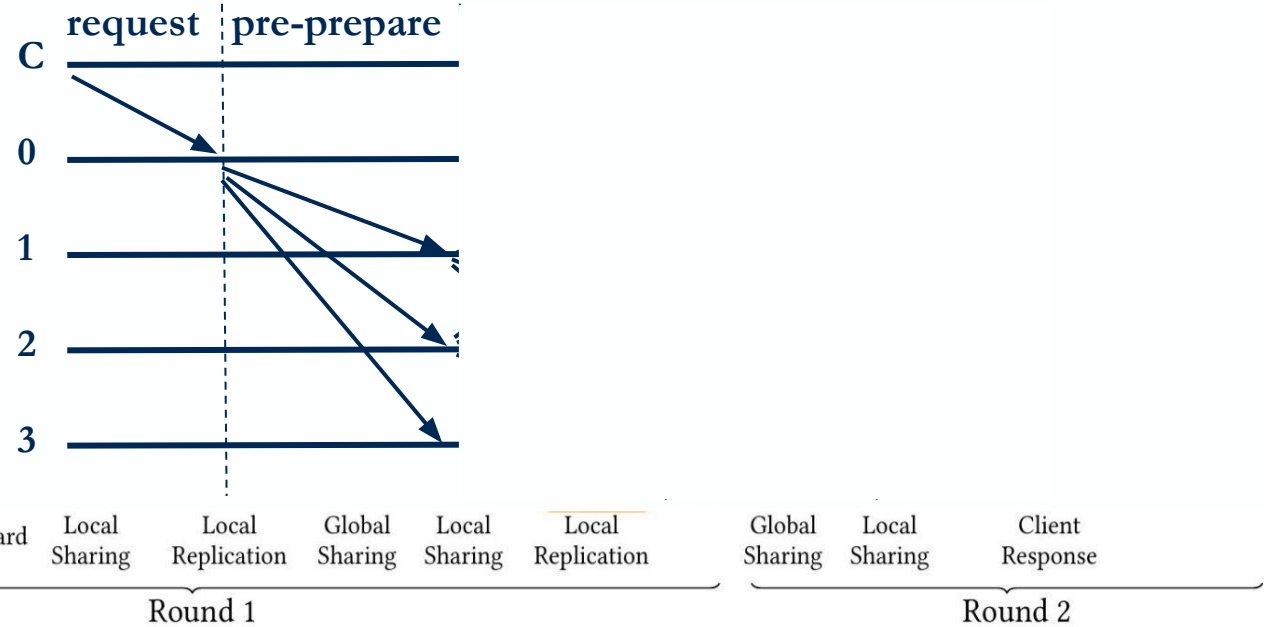
- Check, Validate
- Assign
- Calculate
- Broadcast a preprepare message






Pre-prepare

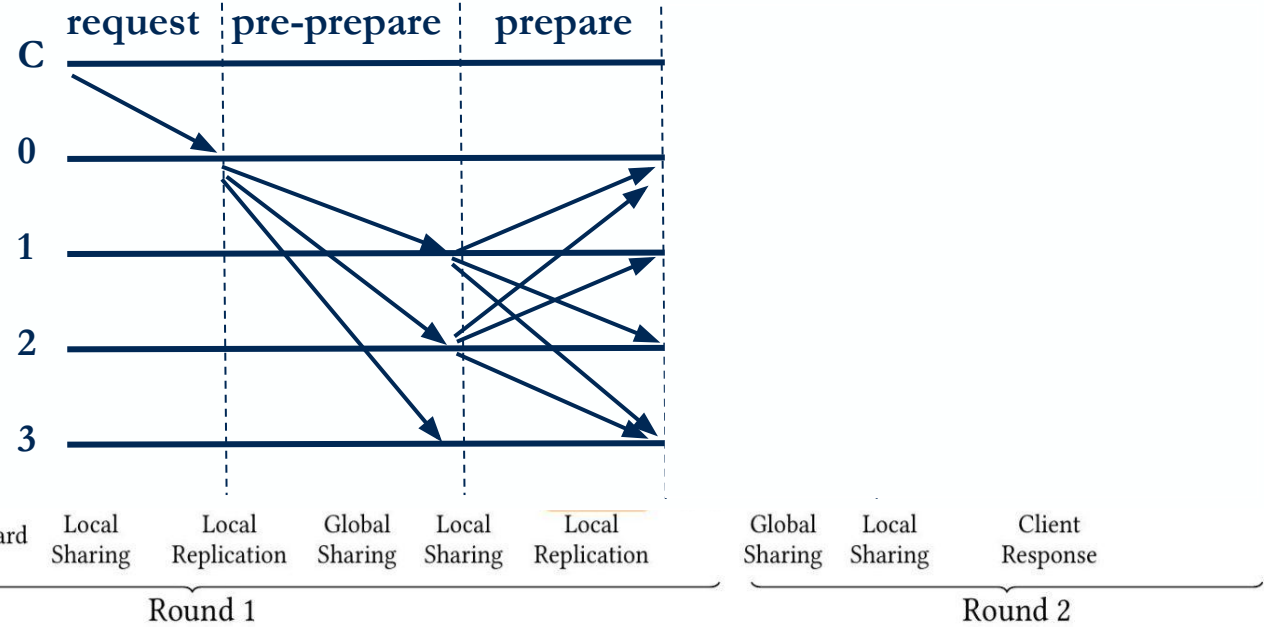
- Check
- Broadcast a prepare message.






Prepare

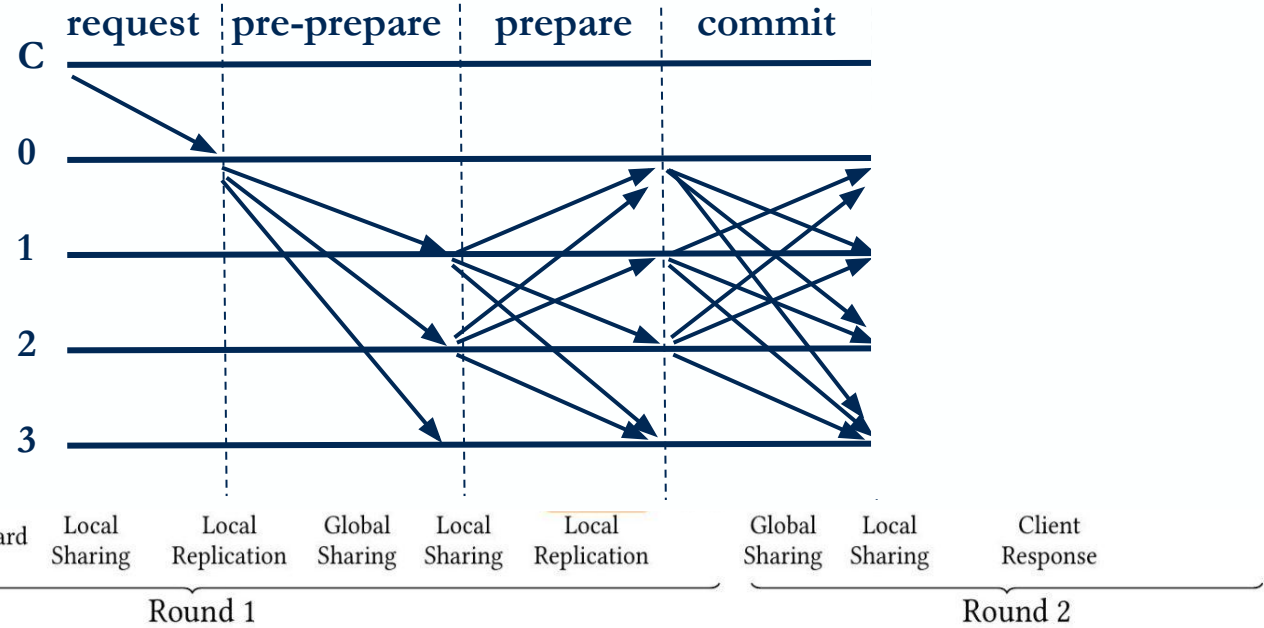
- Receive
- Mark
- Broadcast a commit message






Commit and data lock

- Check and mark
- lock





Forward to next Shard linearly

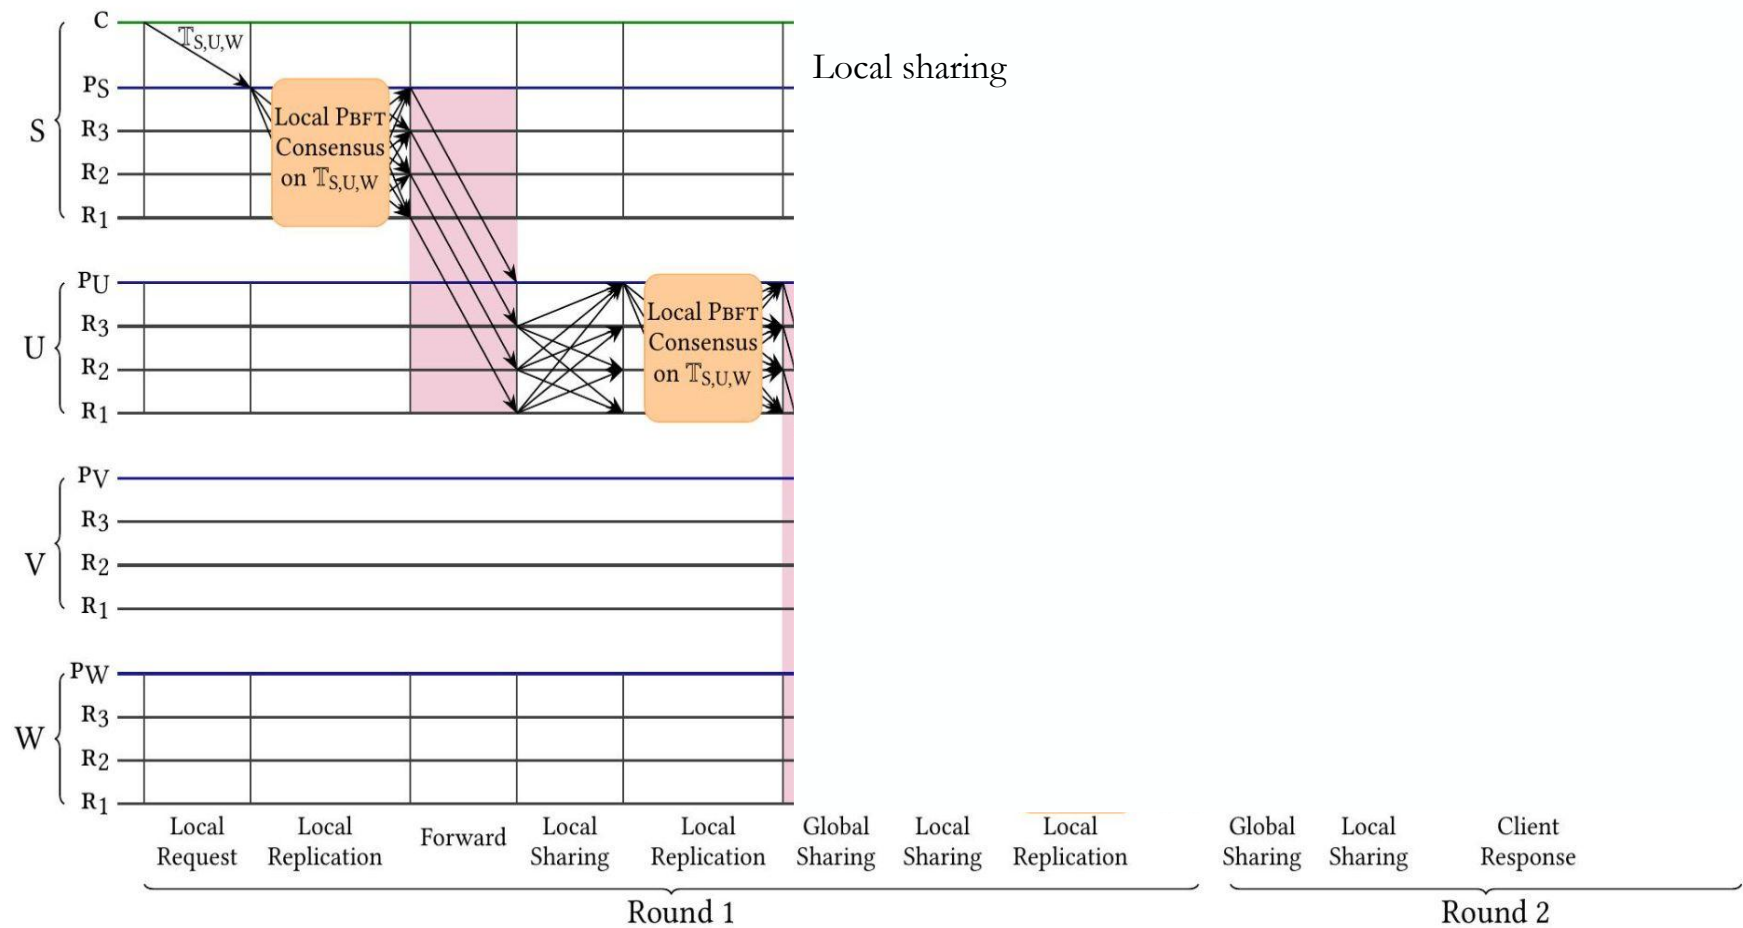
- Same-identifier forwarding
- Commit proof

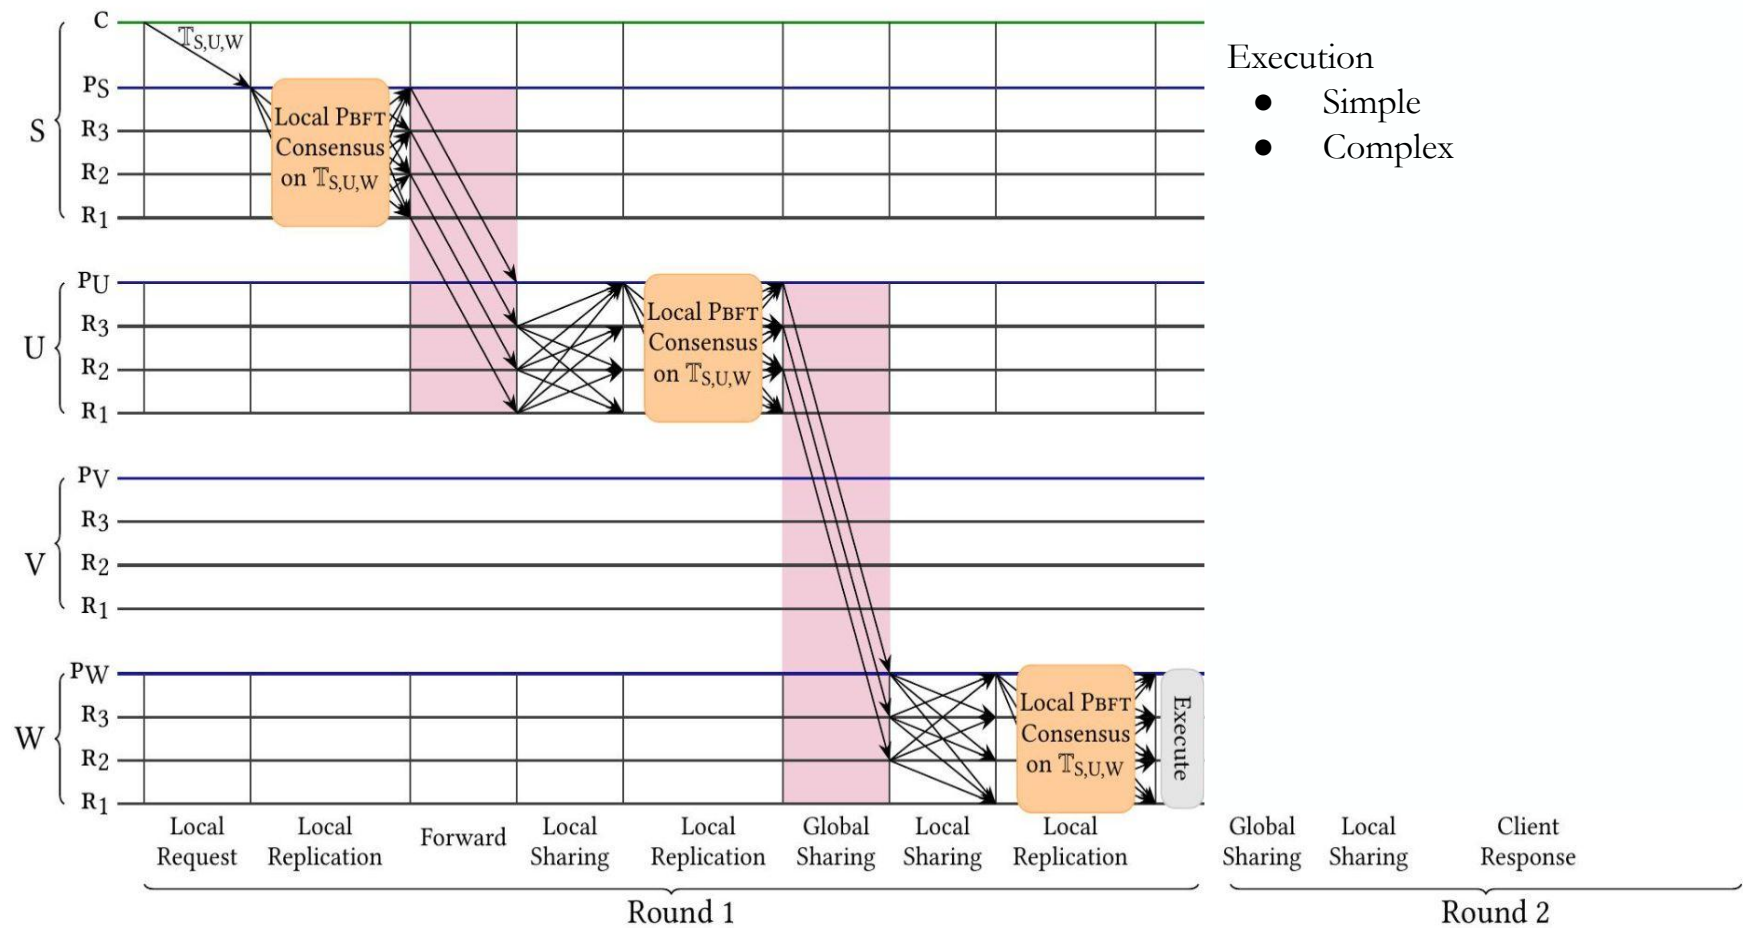
Rightness


- $F+1$ replicas will receive the message

Round 1


Round 2







Deadlock



Uncivil Executions

FLP Impossibility:

“In a fully asynchronous system, is there a deterministic consensus algorithm that can be safe, live, and fault tolerant?”

Premise of RingBFT:

Any BFT protocol can provide safety under asynchronous settings but liveness in the period of synchrony even if up to f replicas are byzantine

Main question:

How can we guarantee liveness during periods of synchrony


Solution:

3 Timers(Replicas) + Recovery Protocols

Uncivil Executions(Single Shard)

Client behavior and attacks:

Situation: Client does not receive sufficient message



Note: Malicious client could always forward transaction to all replicas to blame primary
However, if primary responds to replicas on time, attacks will fail

Uncivil Executions(Single Shard)

Situation: Primary is faulty(Easy) and/or Unreliable network

Solution: View change

Quite similar to the previous example and situation in PBFT


Main reasons for timeout of local timer:

1. Replicas do not receive sufficient commit messages
2. Primary fails to propose a request from client

Uncivil Executions(Single Shard)

Situation: Primary is malicious

Solution: Checkpoint




Primary keeps up to f non-faulty replicas in the dark. Such replicas will send view change messages although will not succeed in the end.

Uncivil Executions(Cross Shard)

Situation 1: No Communication

Solution: Message retransmission(Transmit timer)




Uncivil Executions(Cross Shard)

Situation 2: Partial Communication

Reason: Byzantine Primary of previous shard or unreliable network

Solution: Remote timer of replica in next shard



Evaluations

Benchmark: Yahoo Cloud Serving Benchmark(YCSB)













Included experiments:

1. Scaling number of shards
2. Scaling number of replicas per shard
3. Varying percentage of Cross-shard Txns
4. Varying the Batch size
5. Varying number of involved shards
6. Varying number of clients
7. Impact of Primary Failure
8. Impact of Complex Cross-Shard Transactions

(Please refer to section 8 in the paper for more details)

Evaluations


● RINGBFT ■ SHARPER ▲ AHL




Evaluations

● RINGBFT ■ SHARPER ▲ AHL

(V) Impact of X-Shard Workload Rate (Throughput)



(VI) Impact of X-Shard Workload Rate (Latency)



Conclusions

RingBFT, a novel BFT protocol for sharded blockchains.

For a single-shard transaction, it performs as efficient as any state-of-the-art sharding BFT consensus protocol.

For a cross-shard transaction, it resolves throughput drop by requiring each shard to participate in at most 2 rotations around the ring

RingBFT achieves 25x higher throughput than the most recent sharding protocols and easily scales to nearly 500 globally-distributed nodes.

Reference

1. Michael Abd-El-Malek, Gregory R. Ganger, Garth R. Goodson, Michael K. Reiter, and Jay J. Wylie. 2005. Fault-scalable Byzantine Fault-tolerant Services. In Proceedings of the Twentieth ACM Symposium on Operating Systems Principles. ACM, 59–74. <https://doi.org/10.1145/1095810.1095817>
2. Yair Amir, Claudiu Danilov, Jonathan Kirsch, John Lane, Danny Dolev, Cristina Nita-Rotaru, Josh Olsen, and David Zage. 2006. Scaling Byzantine Fault-Tolerant Replication to Wide Area Networks. In International Conference on Dependable Systems and Networks (DSN'06). 105–114. <https://doi.org/10.1109/DSN.2006.63>
3. Mohammad Javad Amiri, Divyakant Agrawal, and Amr El Abbadi. 2019. SharPer: Sharding Permissioned Blockchains Over Network Clusters. <https://arxiv.org/abs/1910.00765v1>
4. D. Dolev. 1981. Unanimity in an unknown and unreliable environment. In 22nd Annual Symposium on Foundations of Computer Science. IEEE, 159–168. <https://doi.org/10.1109/SFCS.1981.53>
5. Danny Dolev and Rüdiger Reischuk. 1985. Bounds on Information Exchange for Byzantine Agreement. J. ACM 32, 1 (1985), 191–204. <https://doi.org/10.1145/2455.214112>
6. D. Dolev and H. Strong. 1983. Authenticated Algorithms for Byzantine Agreement. SIAM J. Comput. 12, 4 (1983), 656–666. <https://doi.org/10.1137/0212045>

The End