

RapidChain:

Scaling Blockchain via Full Sharding

Mahdi Zamani, Mahnush Movahedi, Mariana Raykova


Presenters: Arindaam Roy & Divjeet Singh Jas

Roadmap

- Limitations of Traditional Methods
- Sharding based Consensus
- RapidChain: Overview
- Bootstrapping
- Consensus
- Reconfiguration
- Evaluation

Traditional Consensus Protocols


- Cannot be used in **Open membership** setting
- Vulnerable to **Sybil Attacks** - [Explained](#)



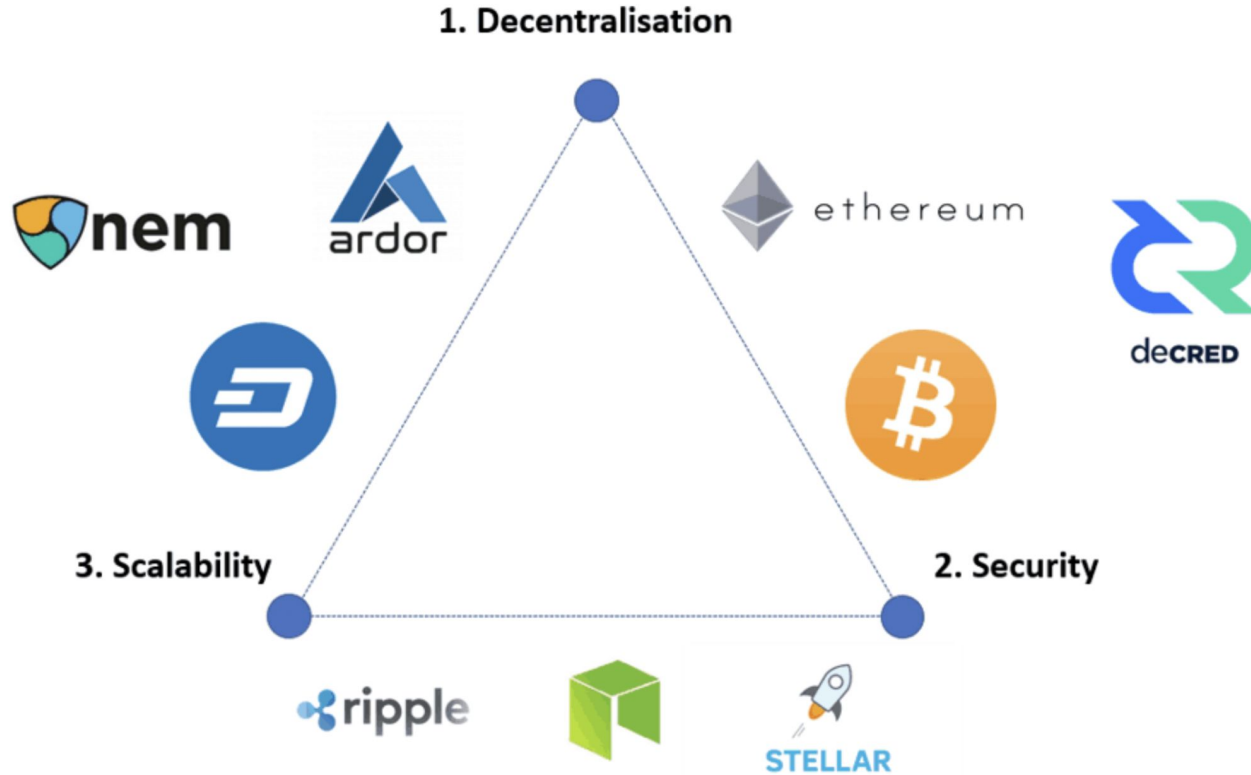
Bitcoin

- Uses **Nakamoto Consensus**
- Inhibits Sybil using **PoW**
- **Full Replication**: Low transaction throughput, high latency, poor scaling
- **Sacrifices** scalability for decentralization






The Blockchain Trilemma




Sharding-Based Consensus

- Electing random small committees
 - Disjoint Blocks Transactions
 - Disjoint Ledgers
- Each committee runs consensus on behalf of all
- And, maintains a disjoint ledger
- Throughput increases linearly with # nodes



Full Sharding

- Computation
- Storage
- Communication



Sharding Challenges

- **Reconfiguration** *to avoid Sybil attack*
- **Cross-shard transactions**
 - Verify transactions located in other committees
- **Decentralized bootstrapping**
 - Creating Initial Random Committees
 - Establish PKI without initial randomness

RapidChain

RapidChain

- Sharding-based public blockchain protocol
- Byzantine faults from up to **1/3** of its participants
- No trusted setup required
- Achieve more than **7,300** tx/sec
- Expected confirmation latency of roughly **8.7** seconds in a network of **4,000** nodes
- A time-to-failure of more than **4,500** years.

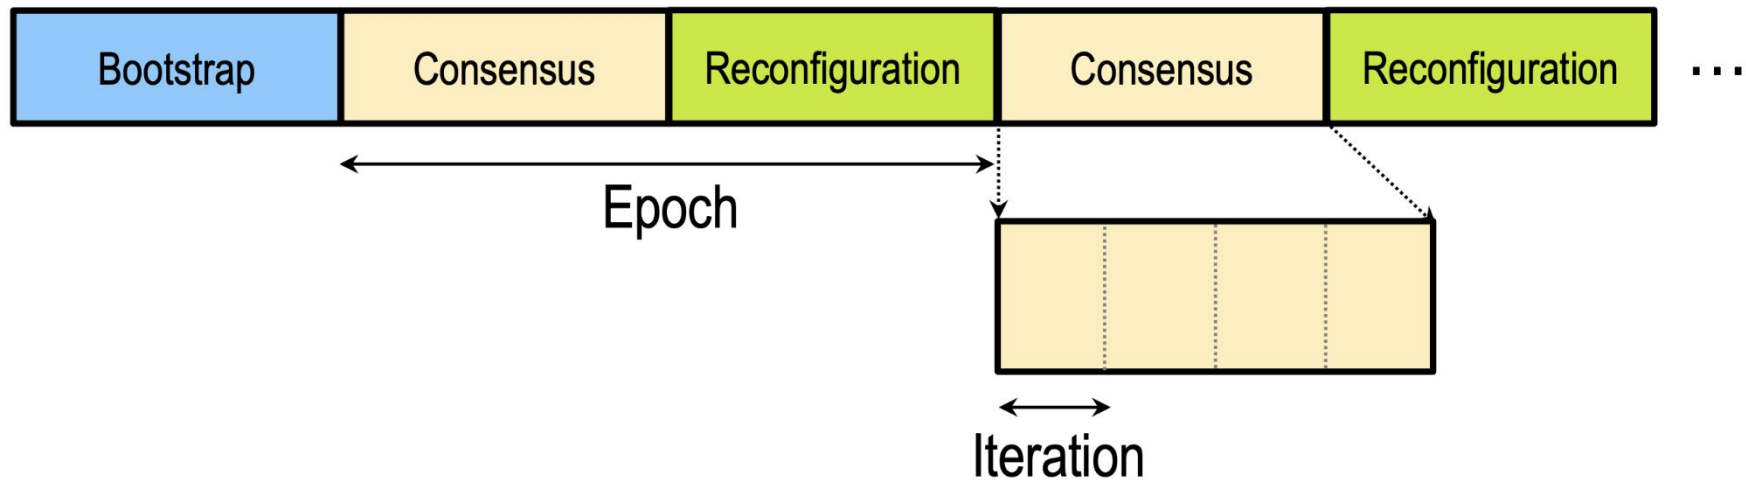
RapidChain Proposes

- A way to bootstrap the initial system
- A way of forming (and re-forming) committees, and allowing nodes to join and leave
- A way of reaching consensus within a committee (shard)
- A way of verifying transactions amongst cross-shards

Model Overview

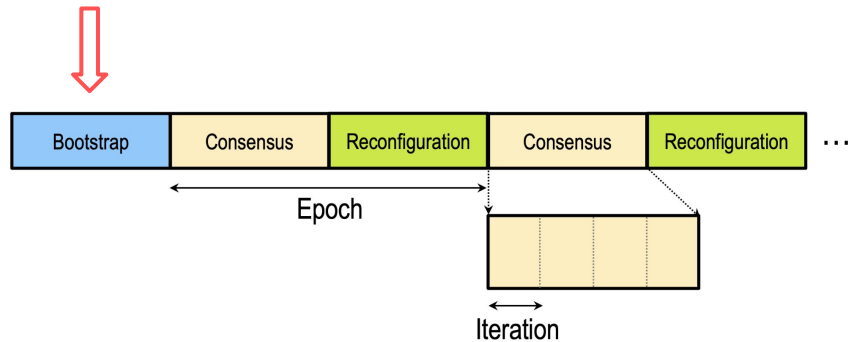
- n nodes with public (pk_i) and secret key (sk_i)
- m committees (or shards), each of size $m = c \cdot \log(n)$, c is a security parameter
- P2P network with gossiping
- Bounded message delay Δ
- $t < n/3$, where t is byzantine nodes in the entire network
- Less than 5% churn in every epoch

Top Level Diagram



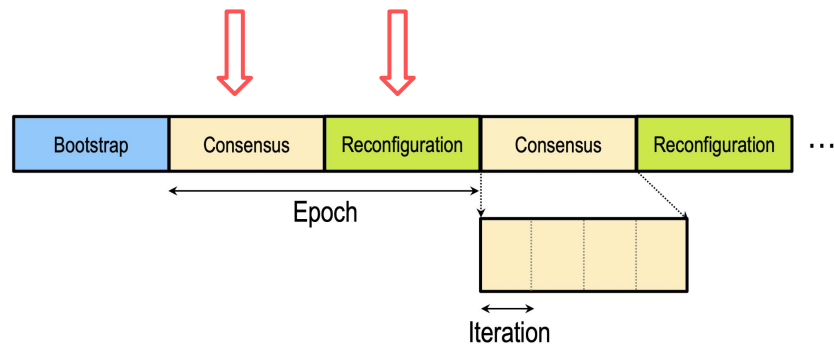
RapidChain Overview

- Proceed in *epochs*
- Bootstrapping Phase - Establishing a reference committee
 - Epoch Randomness
 - Samples sharding committees
 - Challenges for new nodes joining the system
 - Reconfiguration Block



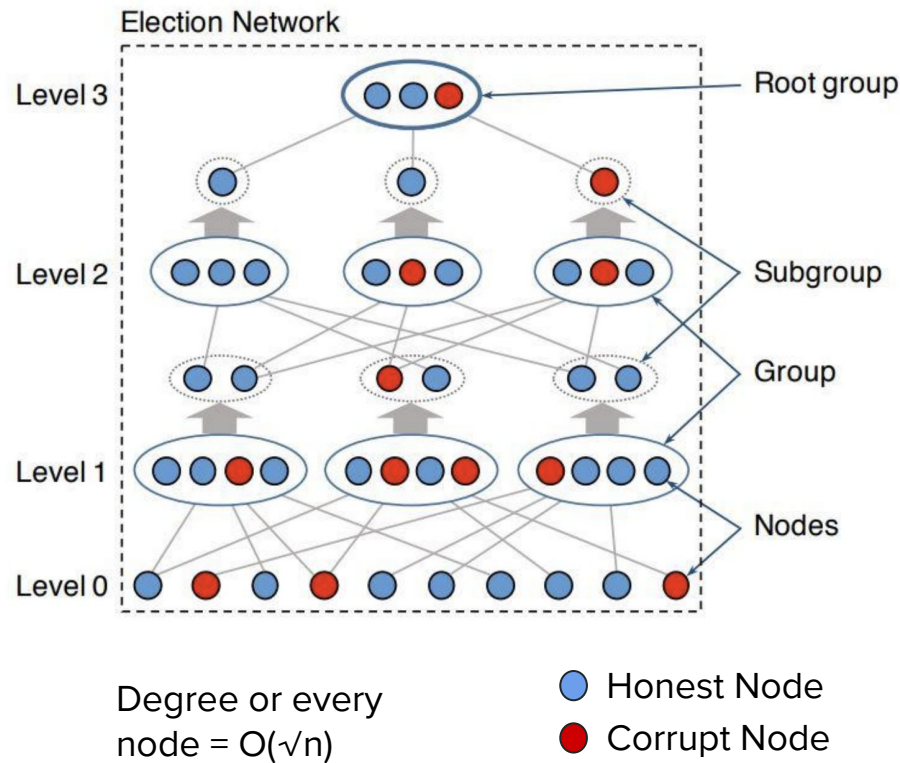
RapidChain Overview (cont'd)

- Consensus Phase
 - Each tx is sent to a random node
 - Tx sent to the output committee, C_{out}
 - Members of C_{out} verifies and adds tx to block
- Reconfiguration Phase
 - Reconfiguration block generated
 - Fresh epoch randomness
 - New list of participants
 - Cuckoo rule used to reconfigure existing committees




Bootstrapping

- Runs an election committee protocol
- Initial nodes have hardcoded seed (s), ID and knows initial network size
- Group formation
- Running subgroup election protocol
 - $h = H(s \parallel \text{ID})$
 - $h \leq 2^{256-e}$
- Subgroup Peer Discovery
- Committee Formation
 - root group of size $O(\sqrt{n})$
 - First shard or reference committee



Bootstrapping (cont'd)

- Reference Committee formation
 - Root group generating a sequence of random bits to establish a reference committee of size $O(\log n)$.
- Establish Committees
 - Reference committee are responsible to create other committees- size $O(\log n)$





Consensus within committees

Idea: Gossip the block, then agree on the hash of the block

Consists of two parts:

- A gossiping protocol to propagate the messages (such as transactions and blocks) within a committee
- A synchronous consensus protocol to agree on the header of the block



Let's Gossip

Visualisation : <https://flopezluis.github.io/gossip-simulator/>



Gossiping Large Blocks

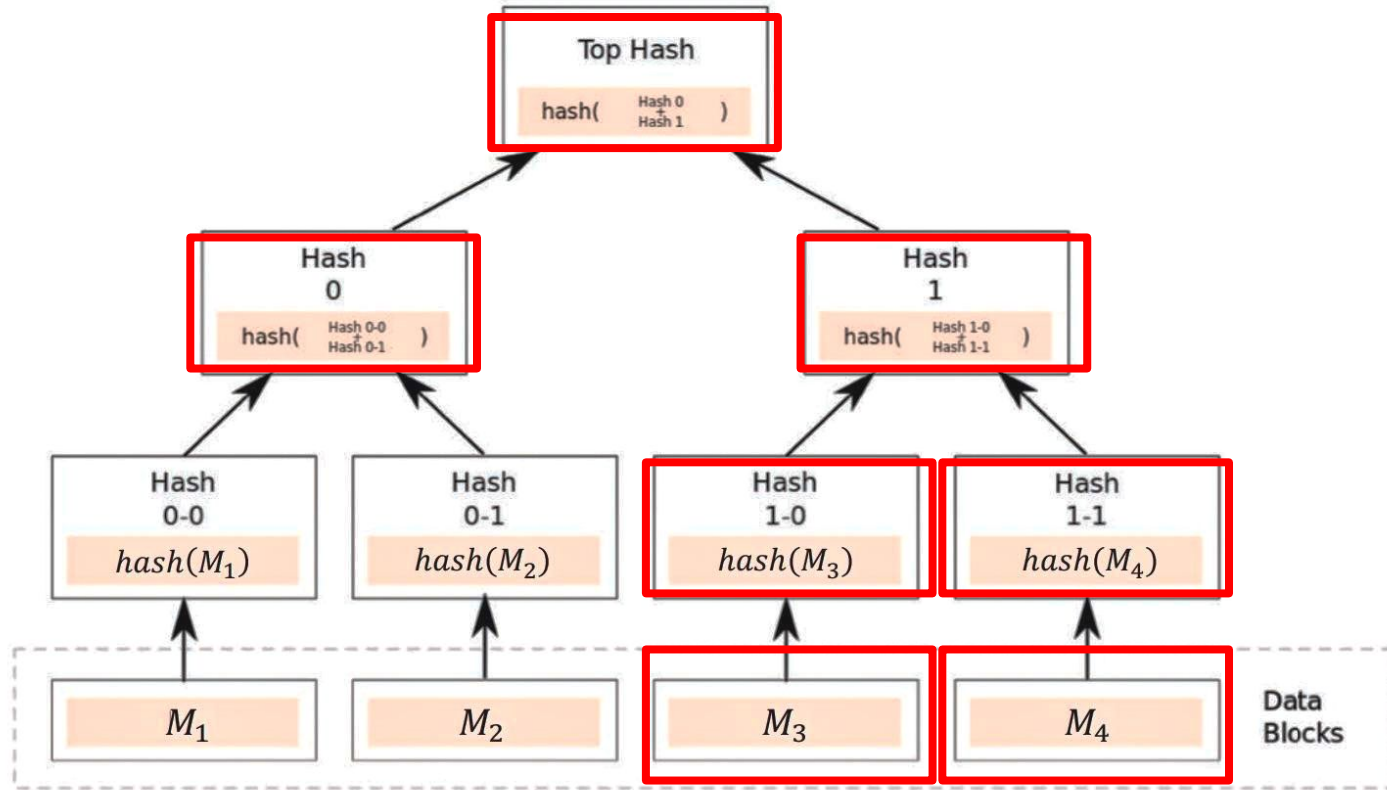
Information dispersal algorithm (IDA)

- Sender divides message M into $(1-\phi)k$ -equal sized chunks $M_1, M_2, \dots, M_{(1-\phi)k}$
- Apply **erasure coding mechanism** for additional ϕk parity chunk to obtain M_1, M_2, \dots, M_k
- Leader node computes a Merkle tree with leaves M_1, \dots, M_k .
- Give each neighbor k/d chunks (gossips M_i and its Merkle proof)
- The message can be reconstructed from any set of $(1 - \phi)k$ valid chunks

Φ - fraction of corrupt neighbors
 K - no. of chunks of message M
 d - no. of neighbors

Gossiping Merkle Hash

- Compute a Merkle hash tree over message chunks M_1, \dots, M_k
- Send Merkle proof along with message chunk to neighbors
- Each node verifies the message using the Merkle proof and the Merkle root.



Consensus Protocol

- Can tolerate $f < \frac{1}{2} m$
- It's a synchronous consensus protocol with a constant delay Δ
- Each committee picks a leader randomly using the epoch randomness
- The leader gathers all the transactions it has received (from users or other committees) in a block B_i
- Leader gossips the block using IDA-gossip and creates the block header H_i that contains the iteration number as well as the root of the Merkle tree from IDA-Gossip.

Consensus Protocol (cont'd)

1

Leader Gossips:

$H_i + \text{propose tag}$

2

Nodes echo received H_i :

$H_i + \text{echo tag}$

3

If nodes see different versions of H_i

Leader is malicious and gossip:

$H_i^j + \text{gossip message with } \textit{pending tag}$

4




If nodes see $mf+1$ echoes and there is only 1 version of H_i

Accept the header and gossip:

$H_i + \text{accept tag} + \text{proof}$

Cross-Shard Transactions

- Tx has a unique id, list of inputs, list of outputs
- Node verifies
 - Input is unspent
 - $\text{Sum of the outputs} < \text{sum of the inputs}$



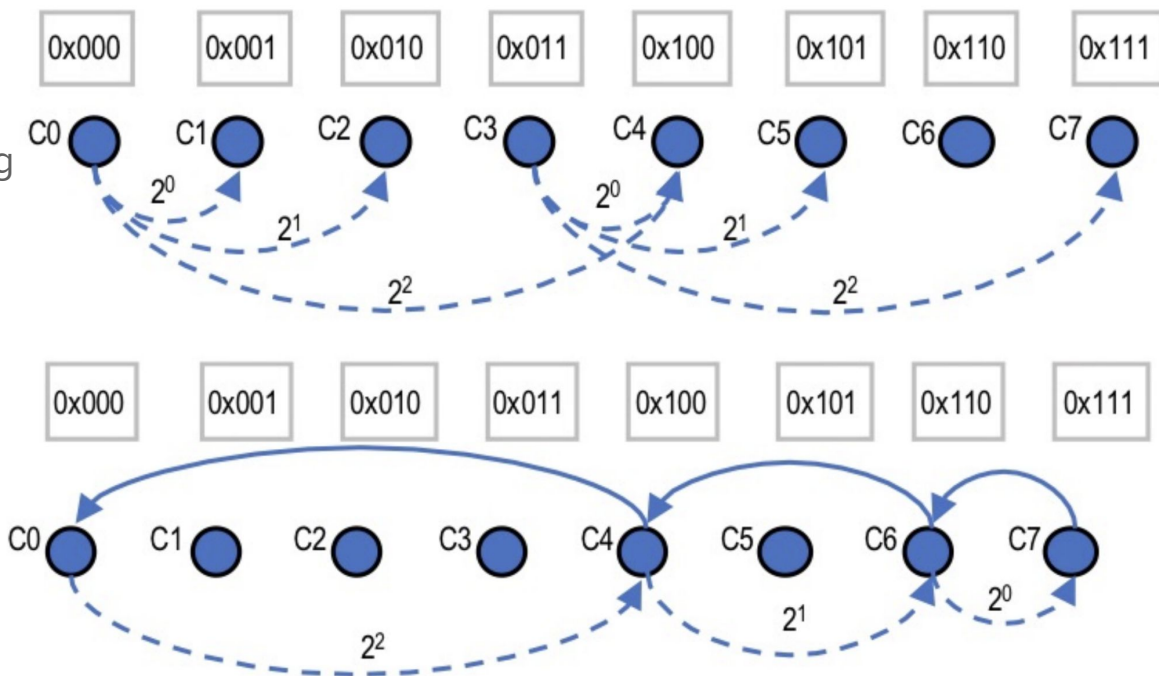
Inter-Committee Routing

Kademlia routing algorithm

Each committee maintains a routing table of $\log n$ records that point to $\log n$ different committees

Distance of 2^i $0 \leq i \leq \log n - 1$

Committee C0 wants to locate committee C7 (via C4 and C6) responsible for transactions with prefix 0x111.



Problem with Committees

- Join/Leave attacks: Corrupt nodes could strategically rejoin the network to take control of a committee.
- Malicious nodes can corrupt the good nodes

RapidChain's Defense


- A 'pay-to-play' scheme (used by bitcoin)
- Selective random shuffling based on the Cuckoo rule

Reconfiguration

- **Offline PoW**
 - Rely on Pow to protect against Sybil attacks
 - Reference committee is responsible to verify PoW result
 - $O = H(\text{timestamp} \parallel PK \parallel r_i \parallel x)$
- **Randomness Generation**
 - Reference Committee run a Distributed random generation protocol
- **Cuckoo Rule**
 - Randomly assign new node
 - Assign a number of members in the committee to another committee


Cuckoo Rule

- New node assigned a random shard
- Evict k nodes from the shard, not including the new node
- Assign these k nodes to another committee



Cuckoo Rule


- New node assigned a random shard
- Evict k nodes from the shard, not including the new node
- Assign these k nodes to another committee



Bounded Cuckoo Rule

Cuckoo Rule

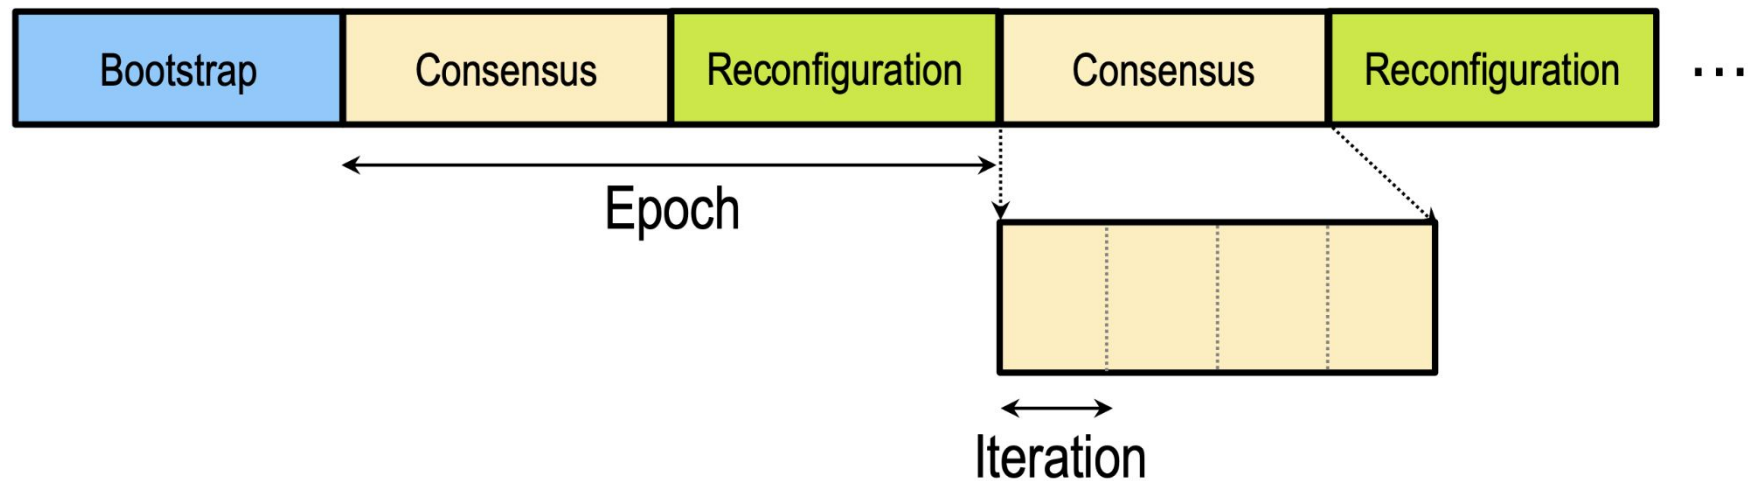
- New node assigned a random shard
- Evict k nodes from the shard, not including the new node
- Assign these k nodes to another committee




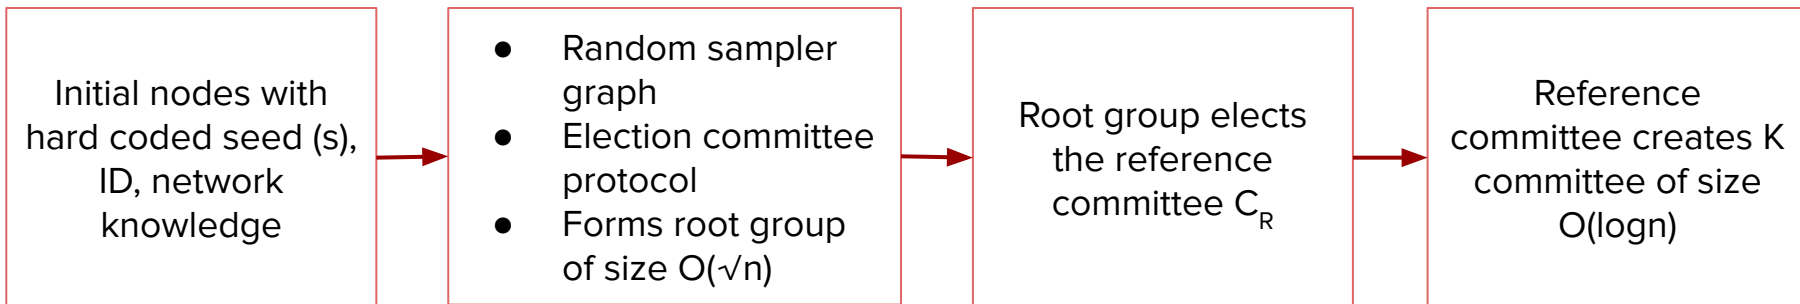
Bounded Cuckoo Rule

What we have learned?

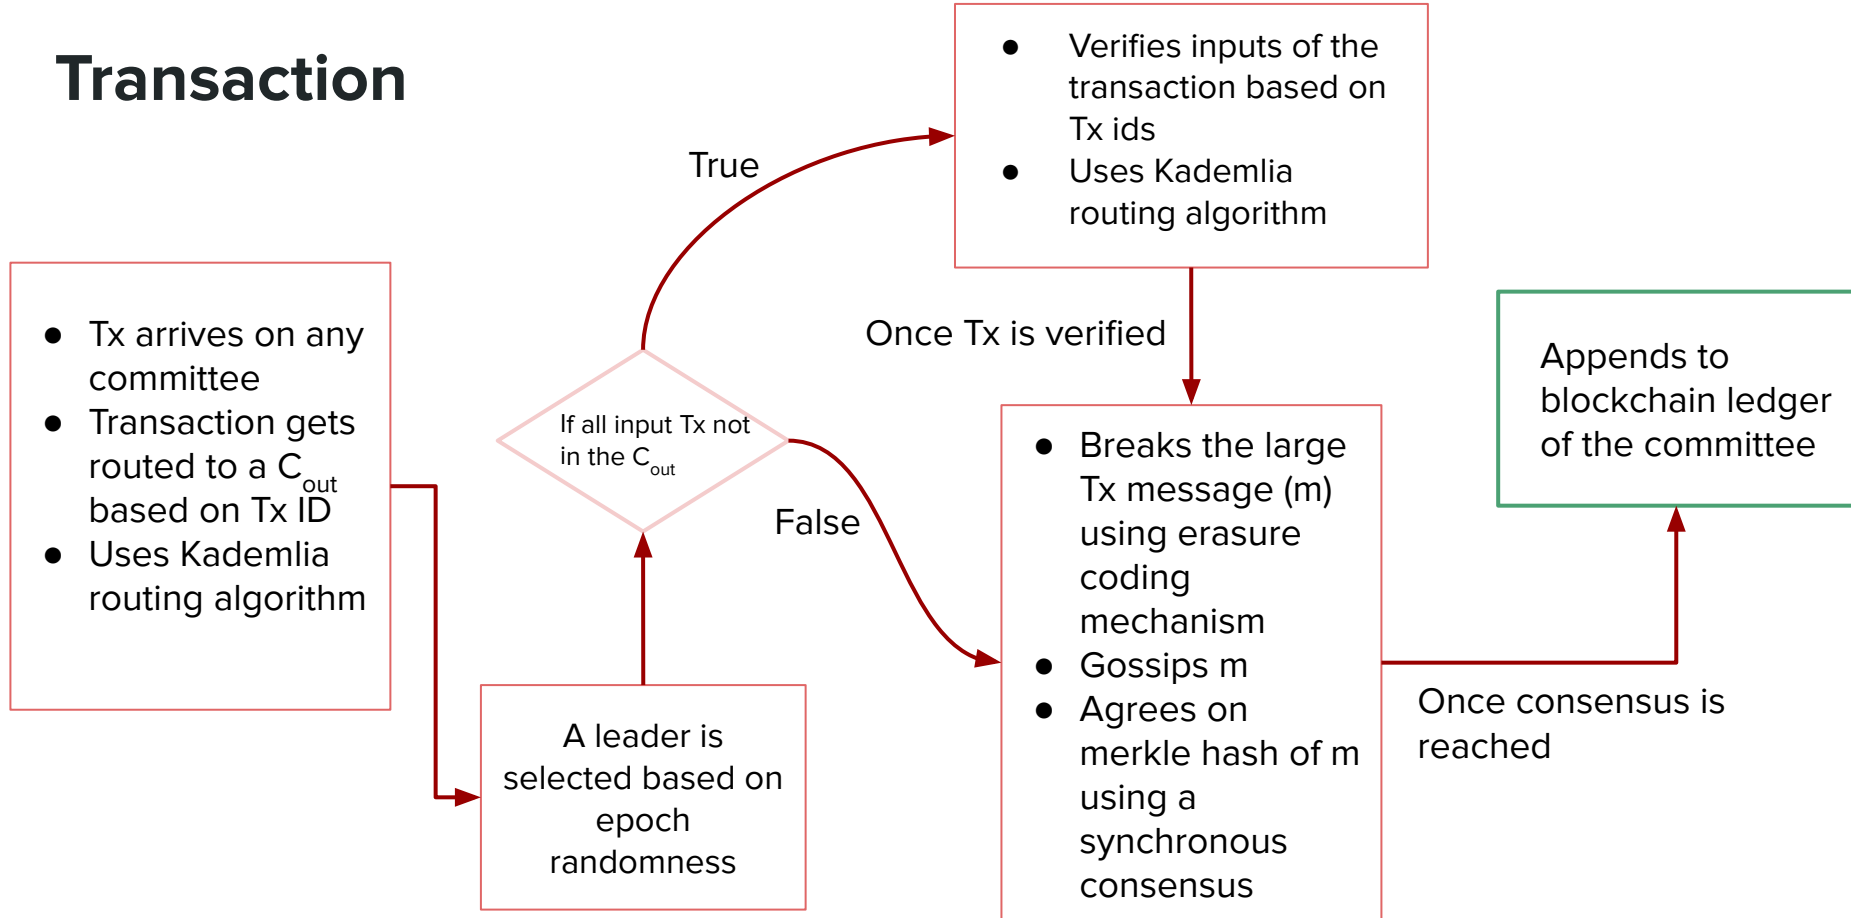
Top Level Diagram



Bootstrap



Transaction



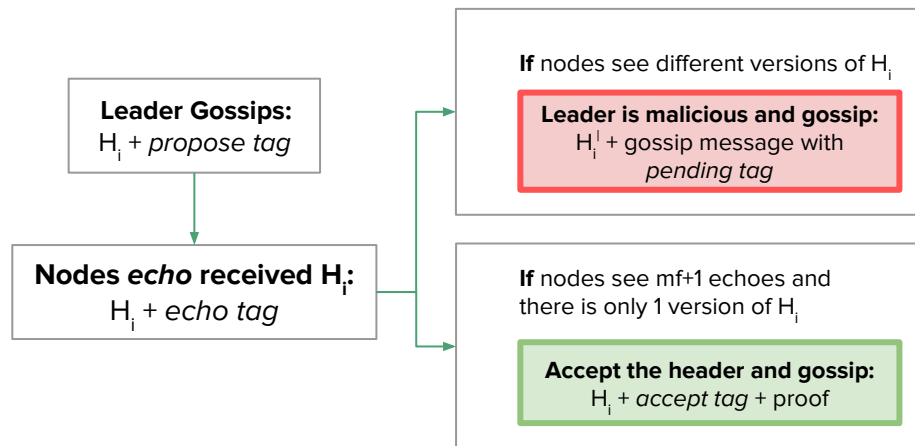
Consensus

1

A **gossiping protocol** to propagate the messages (such as transactions and blocks) within a committee

2

A **synchronous consensus protocol** to agree on the header of the block



Reconfiguration

Offline PoW

- At the end of each epoch
- Rely on Pow to protect against Sybil attacks
- Reference committee is responsible to verify PoW result
- $O = H(\text{timestamp} \parallel \text{PK} \parallel r_i \parallel x)$

Randomness Generation

Reference Committee run a Distributed random generation protocol


Cuckoo Rule

- Randomly assign new node
- Assign some members in the committee to another committee


Every committee gets a new block with updated members

Evaluation


Throughput vs n



Latency vs n



Impact of Block Size



Where does RapidChain stand?

Protocol	# Nodes	Resiliency	TPS	Latency	Storage	Shard Size	Time to Failure
Elastico	1,600	$n/4$	40	800 <i>sec</i>	1x	100	1 <i>hour</i>
OmniLedger	1,800	$n/4$	500	14 <i>sec</i>	1/3x	600	230 <i>years</i>
OmniLedger	1,800	$n/4$	3,500	63 <i>sec</i>	1/3x	600	230 <i>years</i>
RapidChain	1,800	$n/3$	4,220	8.5 <i>sec</i>	1/9x	200	1,950 <i>years</i>
RapidChain	4,000	$n/3$	7,380	8.7 <i>sec</i>	1/16x	250	4,580 <i>years</i>

References

1. <https://eprint.iacr.org/2018/460.pdf>
2. <https://cbr.stanford.edu/seminarTalks/zamani.pdf>
3. <https://blog.acolyer.org/2018/12/07/rapidchain-scaling-blockchain-via-full-sharding>

THANK YOU!
