

# It's not easy to relax:

## liveness in chained BFT protocols

Presenter:

Dakai Kang, Musheng He, Zizhong Li, Xiaoxing Chen, Piaopiao Long




# Part 1


- **Basic HotStuff & Chained HotStuff**
- **Liveness Concern**
- **Siesta**
- **HotStuff vs Siesta**




# Basic HotStuff



# Chained HotStuff



# Liveness Concern



## Consecutive-leader requirement:

- One to propose a block
- Another to collect votes and form QC

Hotstuff needs **4 non-faulty consecutive leaders** to commit!

## Result:

- Averagely, 12 views to commit
- worst-case latency of 129 views

# Siesta

- No protocol can commit in **non-consecutive views** in the presence of Byzantine faults (proof in Appendix C, terrible)
- In the presence of omission faults (crash fault or bad network condition), Siesta commits in **non-consecutive views** as long as one can prove that **no QC** for a conflicting block could have formed in between views.
- Detects **equivocation** and **slashes** a Byzantine leader.

# HotStuff vs Siesta

- $N = 3f+1$
- Partial Synchrony
- Threshold Signature
- 4 consecutive non-faulty leaders
- No slashing

- $N = 3f+1$
- Partial Synchrony
- Digital Signature
- 3 non-consecutive non-faulty leaders
- Slashing

# Hotstuff vs Siesta

## 1. Local State Variables

- **Hotstuff:**

- ***viewNumber** : current view number*
- ***lockedQC** : highest precommittedQC*
- ***prepareQC** : highest preparedQC*

- **Siesta:**

- ***vr** : current view number*
- ***QC<sub>r</sub>** : the QC received with the highest view*
- ***vr<sub>r</sub>** : highest (by view) VOTE-REQ message received*
- ***vp<sub>r</sub>** : highest (by view) VOTE-RESP message sent*



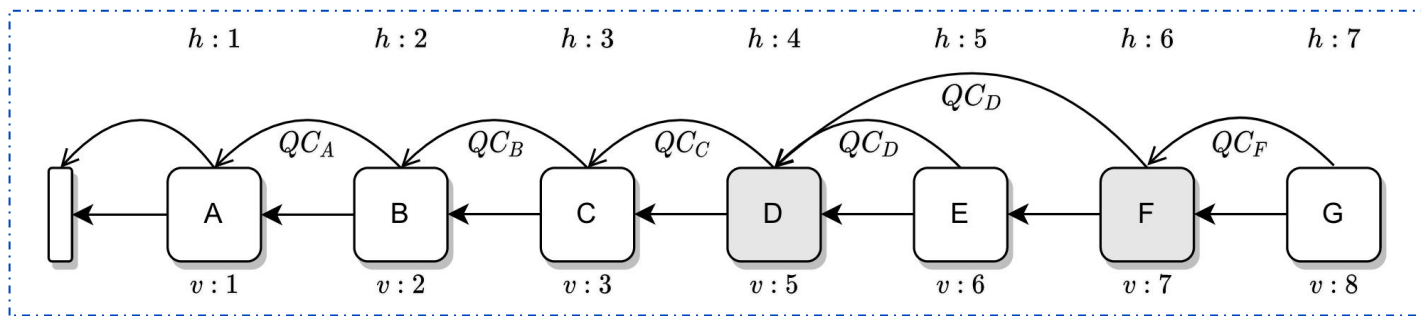
# Hotstuff vs Siesta

## 2. Block

- $h$ : the height of the block
- $H(B_{h-1})$ : hash pointer to parent block
- $b_h$ : batch of client transactions
- $v$ : view in which block was proposed
- $QC$ : QC attesting validity of parent block
- $\sigma_l$ : leader signature

### Additional state for a slow block:

- $S$ : set of NEWVIEW messages
- $\pi$ : optional equivocation proof



# Hotstuff vs Siesta

## 3. *GENERIC* vs *VOTE-REQ*

### *VOTE-REQ for fast blocks*

- *$b_h$* : set of client transactions
- *$h$* : height
- *$v$* : current view
- *$H$* : hash pointer to the parent block
- *$QC$* : QC attesting the validity of the parent block

### *VOTE-REQ for slow blocks*

- *$S$* : set of *NEW-VIEW* Message
- *$\pi$* : optional equivocation proof

# Hotstuff vs Siesta

## 4. NewView & VOTE-RESP

- **VOTE-RESP in Siesta:**
  - same as the NewView message in HotStuff
  - A vote for the block in the VOTE-REQ message received
- **NewView in Siesta:**
  - $\langle \text{NEWVIEW}, v, vr_r, vpr, QC_r \rangle_{\sigma_r}$
  - the highest received VOTE-REQ  $vr_r$
  - the highest sent VOTE-RESP  $vpr$

# Part 2

- **Fast View Change**
- **Slow View Change**



# View Change

## Fast view change:

- As views are contiguous, the new leader is guaranteed to learn the latest certified block and can thus easily extend the chain.

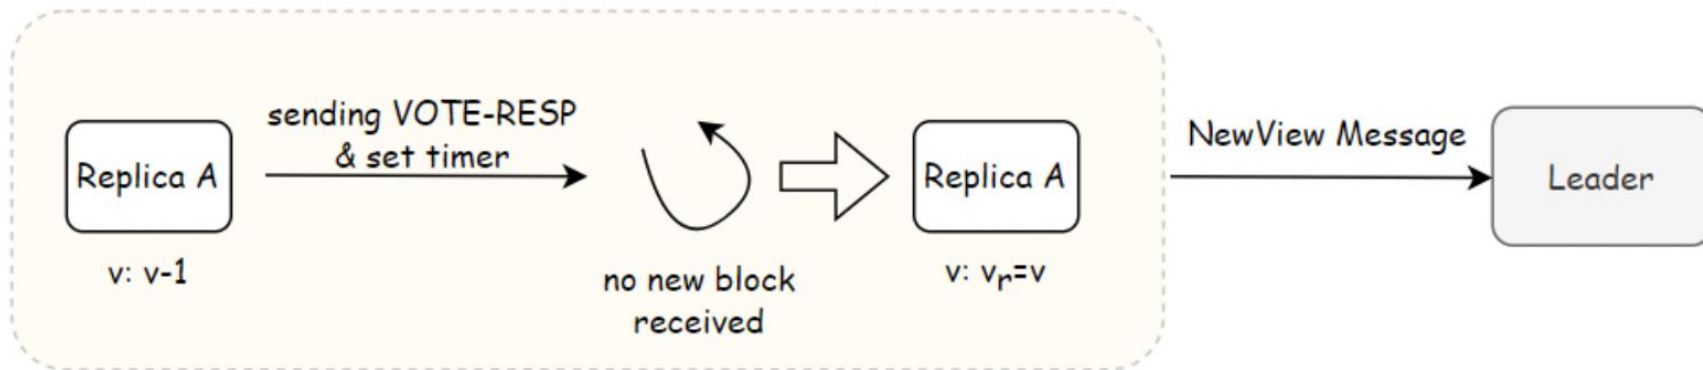
## Slow view change:

- Triggered when sufficiently many  $(2f+1)$  replicas fail to receive a new block proposal before a specific timeout value
- Replicas move on to the next view and transfer necessary state information to new leader
- So it is no longer possible for the new leaders to simply extend a block, as conflicting blocks could have been proposed during views that did not complete.

# Trigger a NewView

Step 1:  $R$  to  $L_v$

Send  $\langle \text{NEWVIEW}, v, QC_r, vr_r, vp_r \rangle_{\sigma_R}$



# New Leader Logic:

## *Decide the QC in VOTE-REQ*

- In HotStuff, new leaders only take existing QCs into consideration
- In Siesta, new leaders attempt to generate new QCs from NEW-VIEW messages received
- We represent VOTE-RESP in NEW-VIEW in the form of a tuple:

**$\langle v, w, hv, hw \rangle$**


$\langle \text{current view, view of parent block, block hash, parent block hash} \rangle$

# New Leader Logic

<3, 2, h3, h2>

<3, 2, h3, h2>

<3, 2, h3, h2>



<3, 2, h3, h2>

$QC_i$

$N = 4, f = 1$

**upon event**  $S = 2f+1 \langle \text{NEWVIEW}, v, vr_r, vp_r, QC_r \rangle$  **do**

**if**  $2f + 1$  matching  $vp_r$  **then**

$QC_i = \text{set of matching } vp_r$

**else**

$B_{anc} = \text{earliest ancestor block of } 2f + 1 \text{ matching } vp_r$

**if**  $B_{anc}$  uncommitted **then**

$QC_{anc} = \text{set of matching } vp_r$

**else**

$QC_{high} = \max(QC_r)_{r \in R}$

$QC_{choice} = \max(QC_i, QC_{anc}, QC_{high})$

$highVote = \max(vr_r)_{r \in R}$

**if**  $QC_{choice} \leftarrow highVote.B$  **then**  $\triangleright highVote$  is the correct block to extend

**else**

**if**  $highVote.B.S$  contains  $n - f$  NEWVIEW messages that do not contain  $QC_{choice}$  **then**  $\triangleright QC_{choice}$  could not have committed

$QC_{choice} = highVote.QC$

**else**

$highVote = vr_r$  that contains  $QC_{choice}$




# New Leader Logic

$\langle 3, 2, h3, h2 \rangle$

$\langle 3, 2, h3, h2 \rangle$

$\langle 2, 1, h2, h1 \rangle$



$\langle 2, 1, h2, h1 \rangle$

$QC_{ans}$

$N = 4, f = 1$

**upon event**  $S = 2f+1 \langle \text{NEWVIEW}, v, vr_r, vp_r, QC_r \rangle$  **do**

**if**  $2f + 1$  matching  $vp_r$  **then**

$QC_i = \text{set of matching } vp_r$

**else**

$B_{anc} = \text{earliest ancestor block of } 2f + 1 \text{ matching } vp_r$

**if**  $B_{anc}$  uncommitted **then**

$QC_{anc} = \text{set of matching } vp_r$

**else**

$QC_{high} = \max(QC_r)_{r \in R}$

$QC_{choice} = \max(QC_i, QC_{anc}, QC_{high})$

$highVote = \max(vr_r)_{r \in R}$

**if**  $QC_{choice} \leftarrow highVote.B$  **then**  $\triangleright highVote$  is the correct block to extend

**else**

**if**  $highVote.B.S$  contains  $n - f$  NEWVIEW messages that do not contain

$QC_{choice}$  **then**  $\triangleright QC_{choice}$  could not have committed

$QC_{choice} = highVote.QC$

**else**


$highVote = vr_r$  that contains  $QC_{choice}$

# New Leader Logic

- $QC_{high}$  is highest QC in the NEW-VIEW messages
- $QC_{choice} = \max(QC_i, QC_{ans}, QC_{high})$

```
upon event  $S = 2f+1 \langle \text{NEWVIEW}, v, vr_r, vp_r, QC_r \rangle$  do
  if  $2f + 1$  matching  $vp_r$  then
     $QC_i$  = set of matching  $vp_r$ 
  else
     $B_{anc}$  = earliest ancestor block of  $2f + 1$  matching  $vp_r$ 
    if  $B_{anc}$  uncommitted then
       $QC_{anc}$  = set of matching  $vp_r$ 
    else
       $QC_{high} = \max(QC_r)_{r \in R}$ 
       $QC_{choice} = \max(QC_i, QC_{anc}, QC_{high})$ 
       $highVote = \max(vr_r)_{r \in R}$ 
      if  $QC_{choice} \leftarrow highVote.B$  then  $\triangleright$  highVote is the correct block to extend
      else
        if  $highVote.B.S$  contains  $n - f$  NEWVIEW messages that do not contain
         $QC_{choice}$  then  $\triangleright QC_{choice}$  could not have committed
           $QC_{choice} = highVote.QC$ 
        else
           $highVote = vr_r$  that contains  $QC_{choice}$ 
```

# New Leader Logic



```

upon event  $S = 2f+1 \langle \text{NEWVIEW}, v, vr_r, vp_r, QC_r \rangle$  do
  if  $2f + 1$  matching  $vp_r$  then
     $QC_i = \text{set of matching } vp_r$ 
  else
     $B_{anc} = \text{earliest ancestor block of } 2f + 1 \text{ matching } vp_r$ 
    if  $B_{anc}$  uncommitted then
       $QC_{anc} = \text{set of matching } vp_r$ 
    else
       $QC_{high} = \max(QC_r)_{r \in R}$ 
       $QC_{choice} = \max(QC_i, QC_{anc}, QC_{high})$ 
       $highVote = \max(vr_r)_{r \in R}$ 
      if  $QC_{choice} \leftarrow highVote.B$  then  $\triangleright$  highVote is the correct block to extend
      else
        if  $highVote.B.S$  contains  $n - f$  NEWVIEW messages that do not contain
         $QC_{choice}$  then  $\triangleright QC_{choice}$  could not have committed
           $QC_{choice} = highVote.QC$ 
        else
           $highVote = vr_r$  that contains  $QC_{choice}$ 
    
```

# New Leader Logic

Propose a new block extending from the block of  $QC_{choice}$ .

$b_h$  = batch of transactions

$vote_{reqs}$  = set of  $vr_r$  messages with the same view  $vr_r.B.v$  but different  $vr_r.B$

**if**  $vote_{reqs} \neq \perp$  **then**

▷ Equivocation proof

$\pi = vote_{reqs}$

$B = (h, H(highVote.B), b_h, v, QC_{choice}, S, \pi)$

$\sigma_{L_v}$  = signature on  $B$

Multicast  $\langle VOTE-REQ, B \rangle_{\sigma_{L_v}}$  to all replicas

# Replica Logic

```
upon event  $\langle \text{VOTE-REQ}, B \rangle$  do
  if  $v == B.v$  then
    if  $B.\pi \neq \perp$  then
      Ignore messages from identified faulty leader
    goto line 4 using  $B.S$            ▷ Check if the leader did the view change correctly
    if result matches  $B$  then
       $vr_r = \langle \text{VOTE-REQ}, B \rangle$ 
       $QC_r = B.QC$ 
      Send  $\langle \text{VOTE-RESP}, B \rangle_{\sigma_r}$  to  $L_{v+1}$ 
       $vp_r = \langle \text{VOTE-RESP}, B \rangle_{\sigma_r}$ 
       $v_r = v_r + 1$ 
    else
      Ignore message
  else
    Ignore message
end event
```

# Part 3


- **Commit Rule**
  - **Contiguous QCs**
  - **Non-contiguous QCs: No-QC proof**
- **Slashing**
- **Advantage**



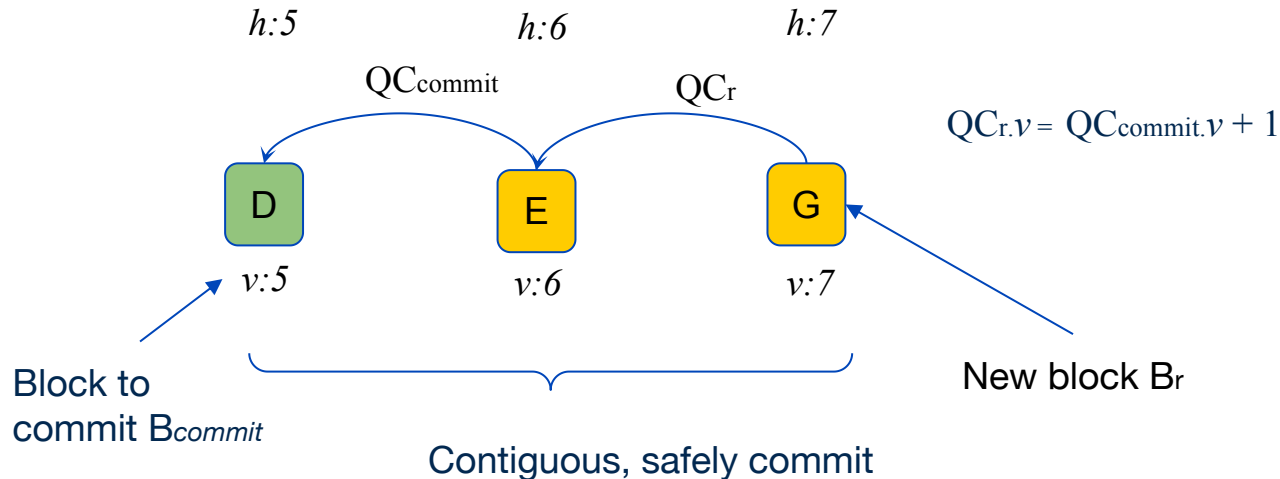
# Commit Rule

## Traditional way: contiguous QC


- **1st QC:** Certifies Block and achieve agreements across replicas
- **2nd QC:** achieves persistence across views




## Siesta Support



# Non-Contiguous QC



**Conflict**




**No conflict**

$$QC_r.v > QC_{commit}.v + 1$$



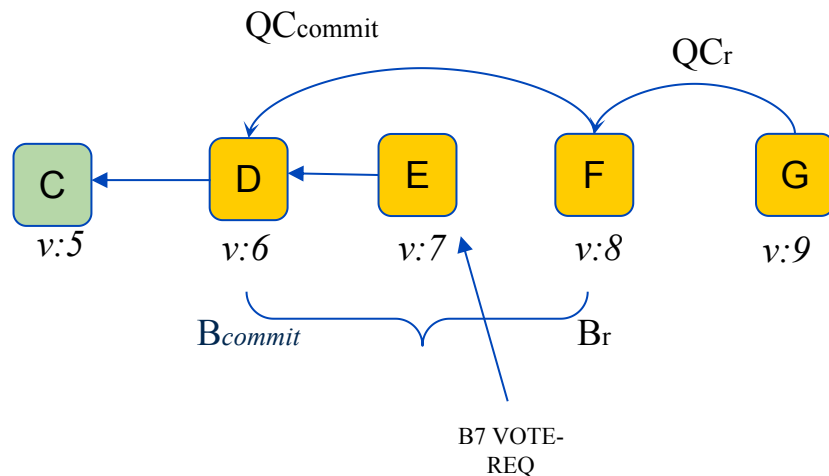
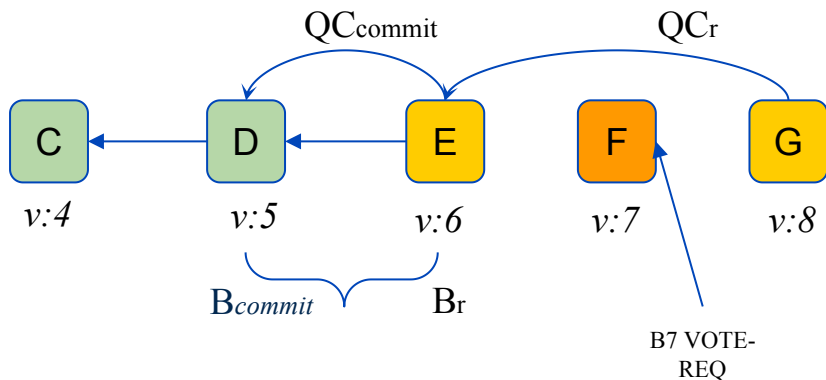
# No-QC Proof in Detail



- It is safe to commit a block in non-consecutive views as long as one can prove that no QC for a conflicting block could have formed in between views. Siesta carefully designs such proofs to be efficient.
- If a QC for  $B_{\text{conflict}}$  forms, all subsequent leaders will observe **at least one Vote-req message** for  $B_{\text{conflict}}$
- Check every block on the chain between  $B_r$  (included) and  $B_{\text{commit}}$  to determine whether such a conflicting Vote-req exists.

# How No-QC Works

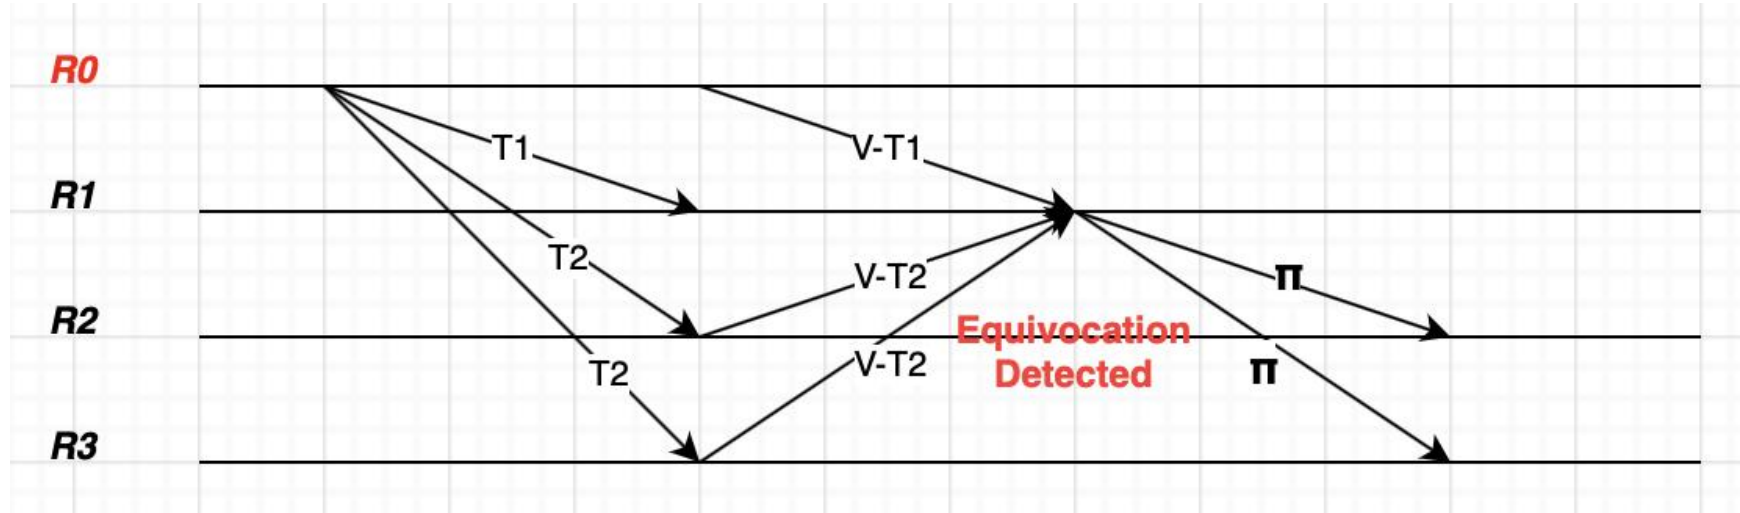
## Conflicting Detected



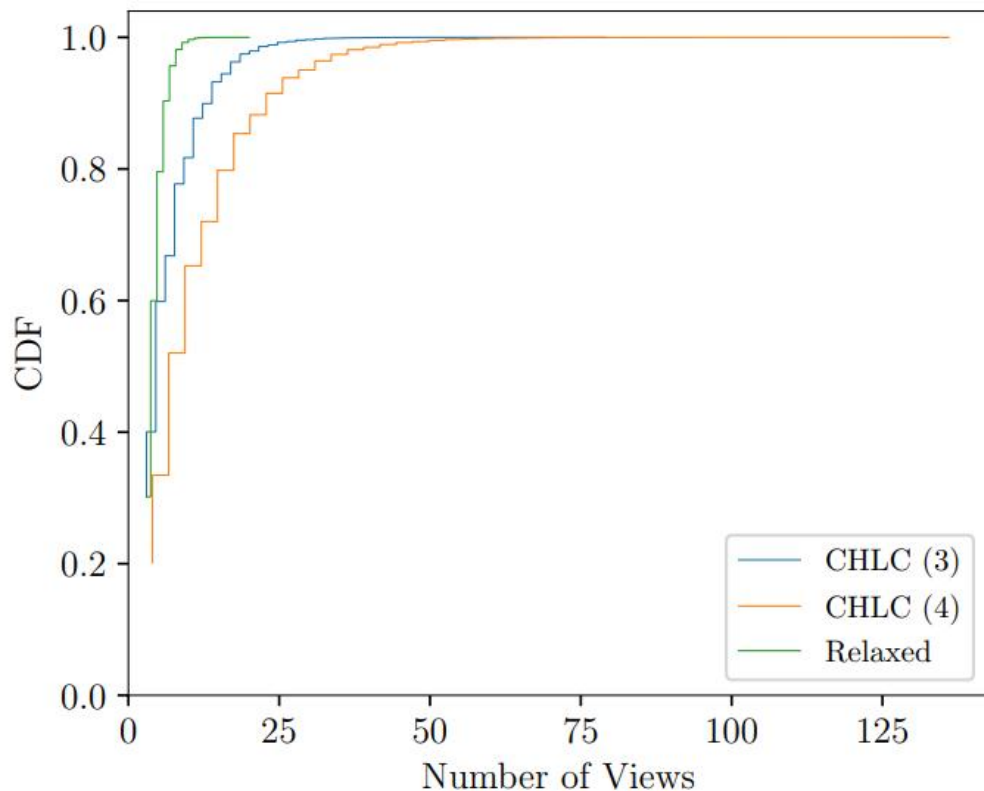
- $r$  checks every block on the chain between  $B_r$  and  $B_{commit}$  to determine whether such a conflicting VOTE-REQ exists.

No  $B'$  VOTE-REQ  $\Rightarrow$  No Conflicting

# Equivocation Proof & Slashing



# The Advantage of Siesta



## Average commit latency:

- Siesta: 4.5
- CHLC(3): about 7 rounds
- CHLC(4): 12 rounds

## Worst-case commit latency:

- Siesta: 18
- CHLC(3): 76
- CHLC(4): 129

# It's not easy to relax: liveness in chained BFT protocols

**Q & A**

Presenter:

Dakai Kang, Musheng He, Zizhong Li, Xiaoxing Chen, Piaopiao Long

