# Making Byzantine Consensus Live

By Manuel Bravo, Gregory Chockler, Alexey Gotsman

Presented by Che-Yu Chang, Yan-Da Chen, Ling-Yuan Chen

October 24th Fall 2022
Department of Computer Science
UC Davis



#### Outline

Liveness

- Synchronizer Specification
- FastSync
  - Switching view
  - Relaying Wishes



#### **Byzantine Consensus**

- Several replicas need to agree on a single value
- Some replicas may be malicious:  $\leq$  f out of n = 3f + 1
- Byzantine consensus: decision finality, but permissioned system fixed set of participants
- Blockchain consensus (proof-of-work, proof-of-stake): no hard finality guarantees, but permissionless system - anyone can participate

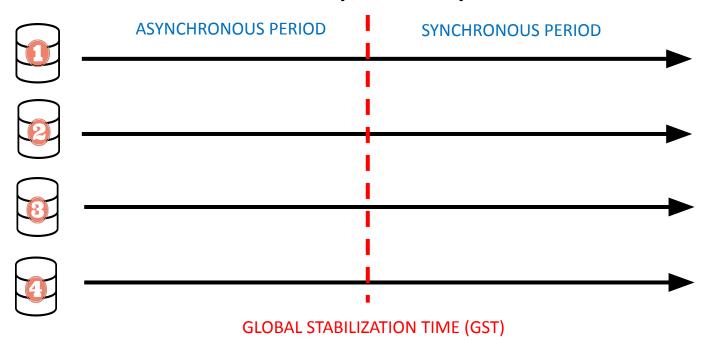


#### The Liveness Problem

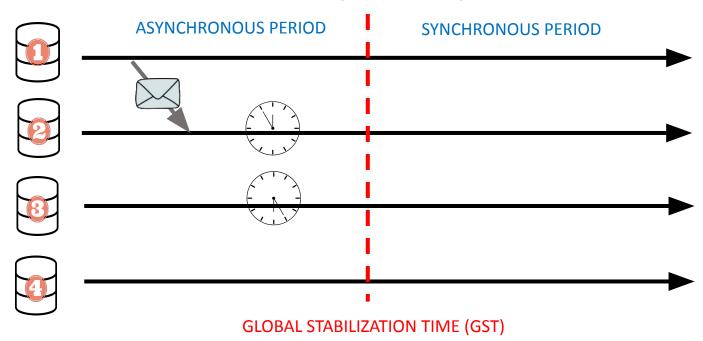
- Many complications arise from ensuring liveness: protocol has to be responsive make a decision
- Can not make consensus both safe and live in an asynchronous network = no bound on message delivery time
   Can not tell a crashed replica from a slow one
- Provide safety and liveness only under synchrony



- Synchronous model: there exists some known finite time bound  $\delta$ . For any message sent, the adversary can delay its delivery by at most  $\delta$
- Asynchronous model: for any message sent, the adversary can delay its delivery by any finite amount of time.
- Partial synchrony model: middle ground between these two models. The assumption is that there exists some finite time bound  $\delta$  and a special event called GST (Global Stabilization Time).

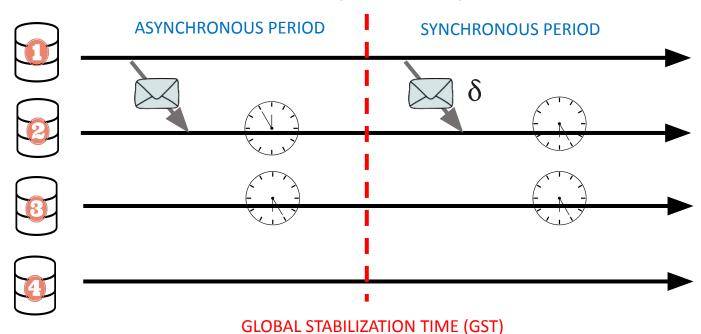






- Messages delayed or lost
- Replica clocks out of sync: synchronized using messages



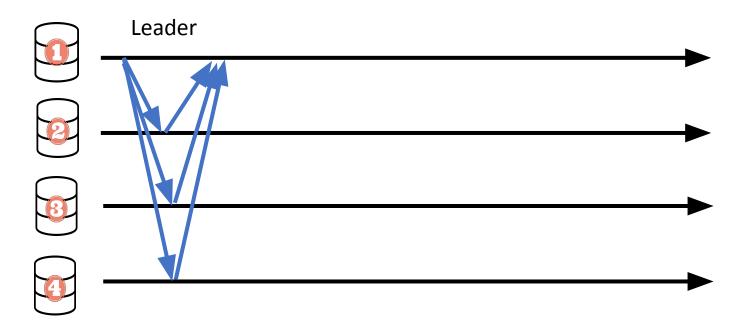


- Messages delayed or lost
- Replica clocks out of sync: synchronized using messages

- Messages between non-faulty replicas delivered within and unknown time  $\boldsymbol{\delta}$
- Replica clocks track real time.



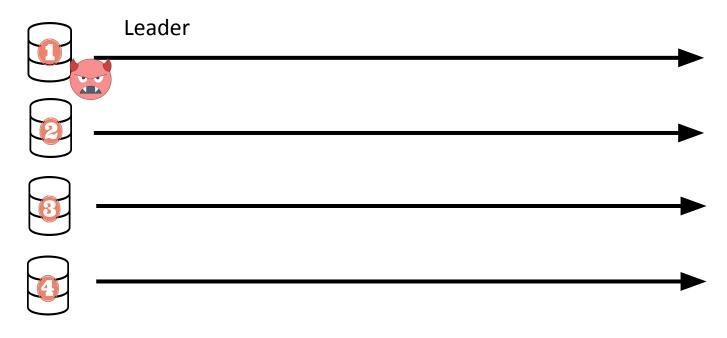
#### Leader-driven Consensus



- Replicas: can vote to accept a value
- Leader: proposes a value to vote on
- Decision: 2f + 1 votes out of 3f + 1



#### **Failed Votes**



- Votes may fail: faulty leader or asynchronous network
- Can not tell the difference between the two: have to be able to change the leader



# Views Leader Leader Leader View 1 View 2 View 3

• Divide the execution into views, each with a fixed leader: view % n

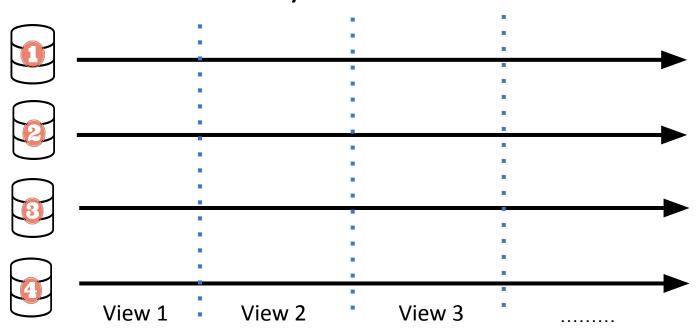


# **Views** Leader Leader Leader View 1 View 2 View 3 **GST**

- Divide the execution into views, each with a fixed leader: view % n
- Eventually hit a correct leader and after GST messages get through,
   then we can terminate

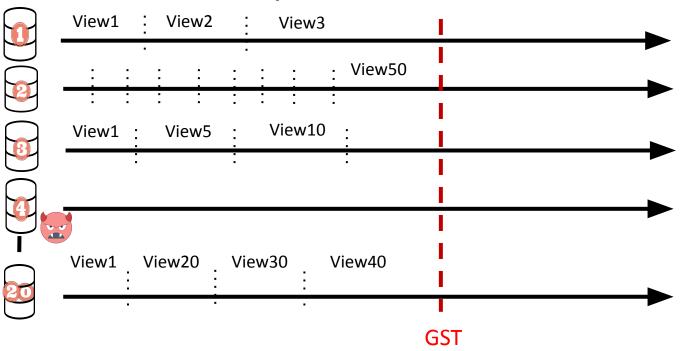


### View Synchronization



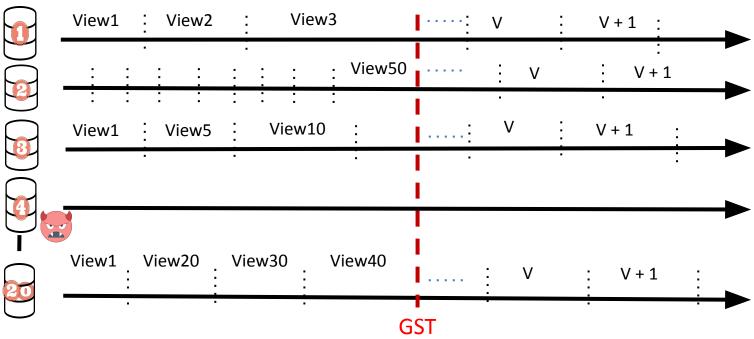


#### View Synchronization



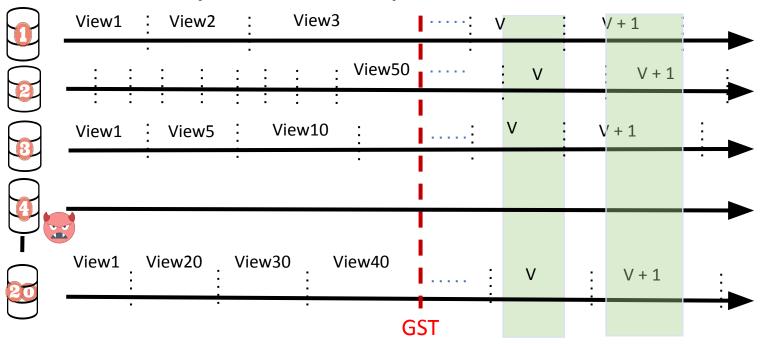
- Before GST: clocks out of sync, messages delayed or lost
- After GST: bring all non-faulty replicas into the same view
- Despite malicious replicas trying to disrupt this





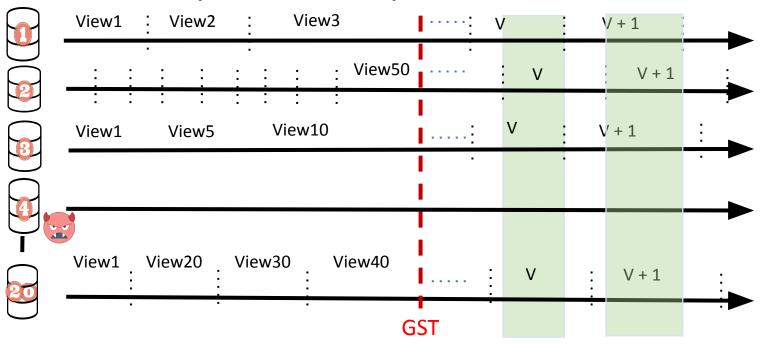
- J view V: all non-faulty replicas enter every view >= V after GST and overlap in it for a non-zero time
- Leaders rotate round-robin ⇒ keep hitting correct leaders





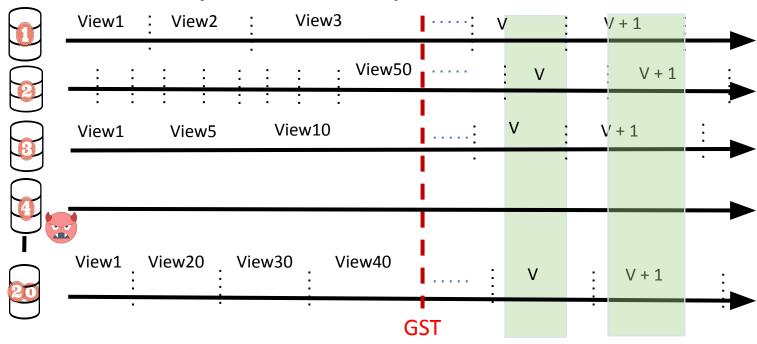
- view V: all non-faulty replicas enter every view >= V after GST and overlap in it for a non-zero time
- Leaders rotate round-robin ⇒ keep hitting correct leaders





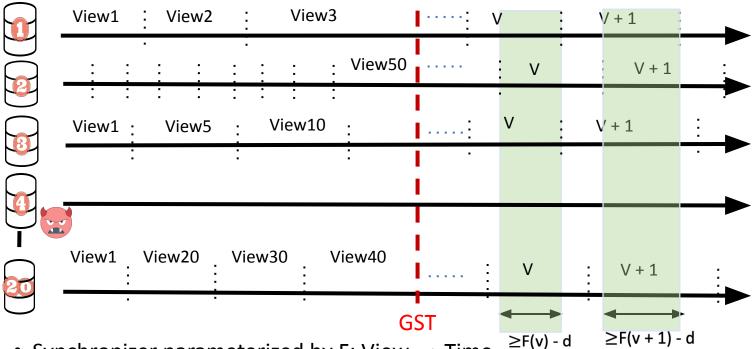
- To decide, replicas need to stay in the view long enough to complete voting
- Don't know  $\delta \Rightarrow$  keep increasing view duration until it's long enough





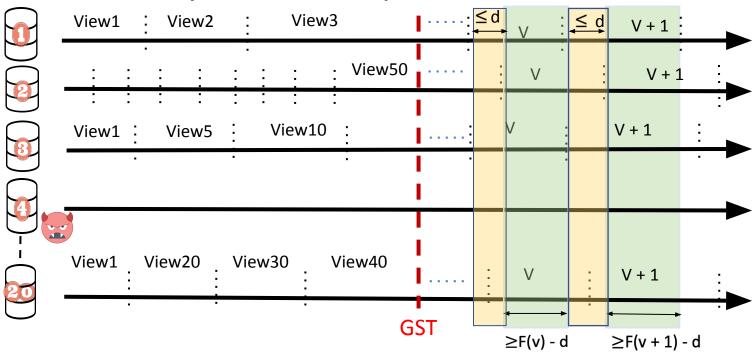
- Synchronizer parameterized by F: View  $\rightarrow$  Time
- Monotonic and unboundedly increasing:  $F(v) = c^{v}$  or  $F(v) = c^{*}v$





- Synchronizer parameterized by F: View → Time
- Monotonic and unboundedly increasing:  $F(v) = c^v$  or  $F(v) = c^*v$
- Replicas overlap in view v for at least F(v) d (e.g.,  $d = 2\delta$ )
- Eventually hit a long enough view with a correct leader ⇒ decide

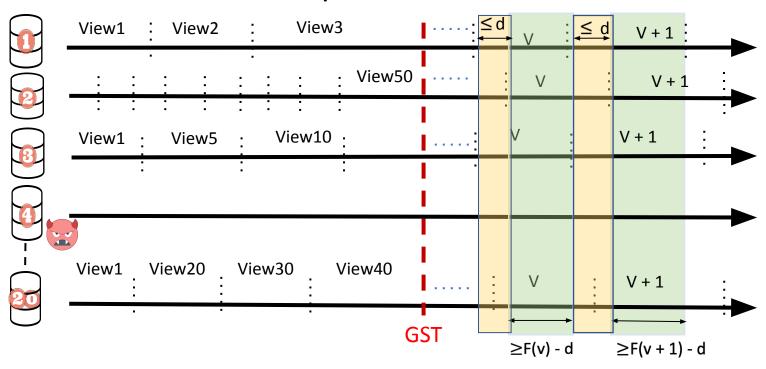




• Non-faulty replicas enter each view within d (e.g.,  $d = 2\delta$ )



#### **View Specification**



Proving liveness of several protocols: Hotstuff, Tendermint, PBFT



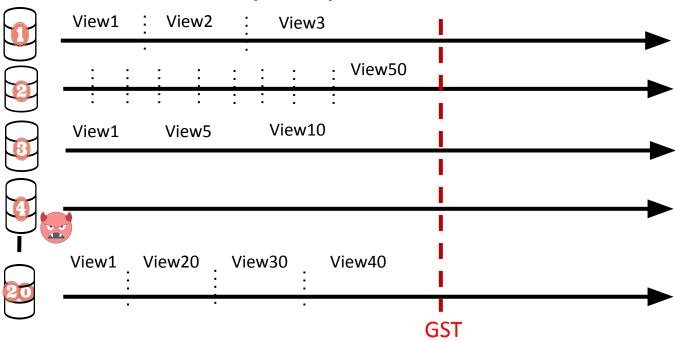
### FastSync Synchronizer



- Upon entering a view, set a timer for F(v)
- Try to switch views when the timer expires



### FastSync Synchronizer



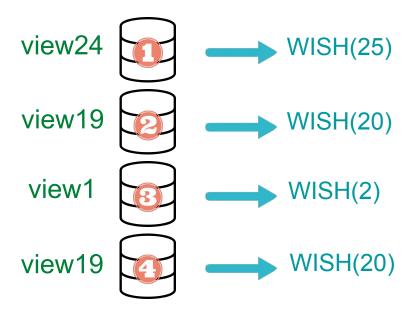
- At GST replicas may be in different views
- Timers are not enough: replicas will have to communicate



- When a timer for view v expires, try to enter v+1
  - ⇒ broadcast WISH(v+1)

- Enter view v+1 when enough replicas express a similar wish
   E.g., WISH(v+1)
- Enough replicas = 2f+1 out of 3f+1 replicas
  - ⇒ guard against disruption by faulty replicas



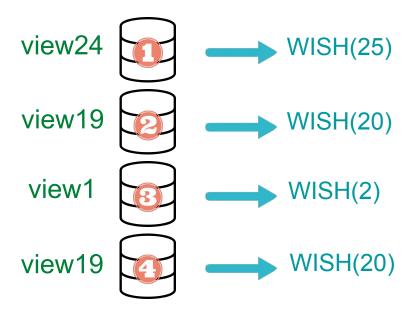




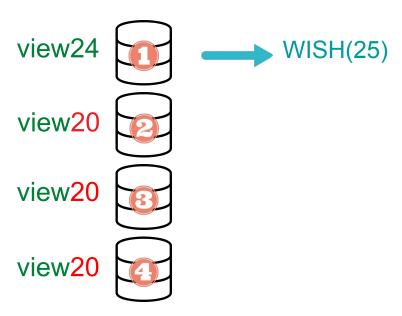


- Bounded space: maintain an array with the highest WISH received from each replicas
- When received 2f+1 WISHes for views > your (view),
   enter directly to the minimal WISH view in them

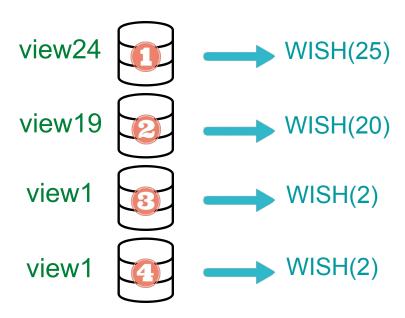




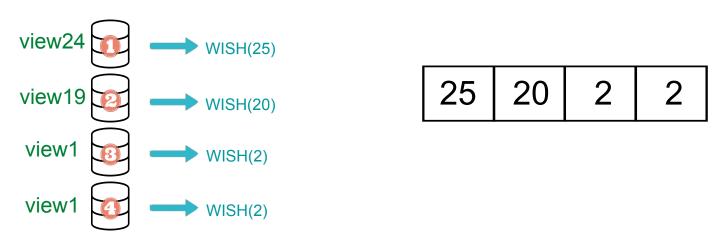






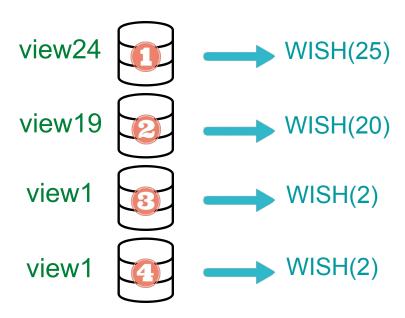




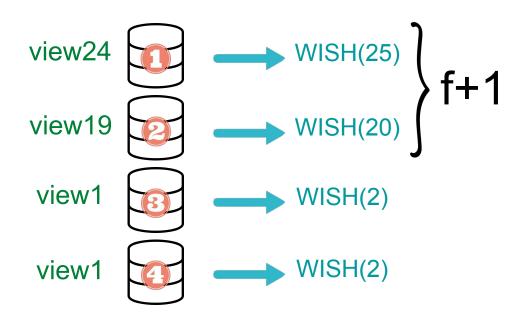


- f+1 WISHes > the view you're trying to enter (WISH)
  - ⇒trying to enter the minimal view in them and relay the corresponding WISH

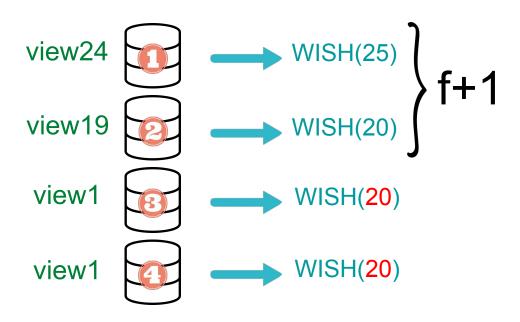




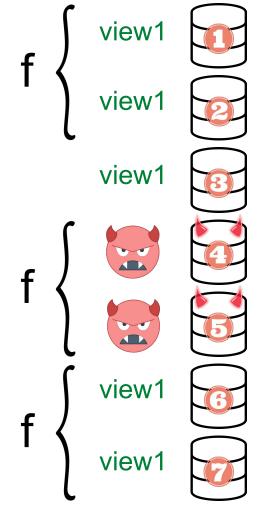




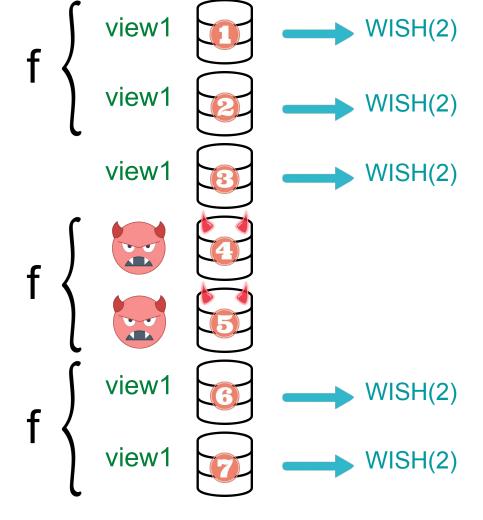




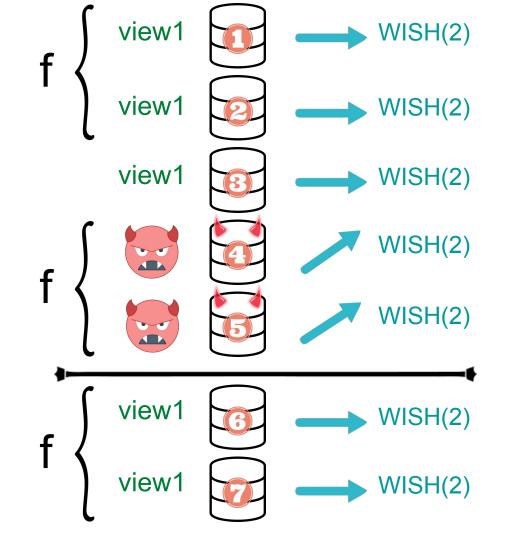








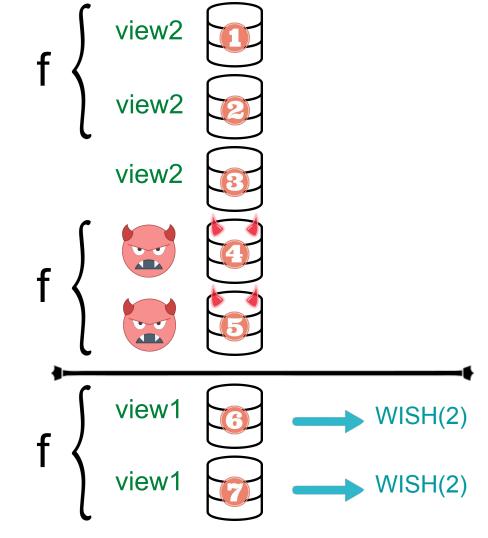




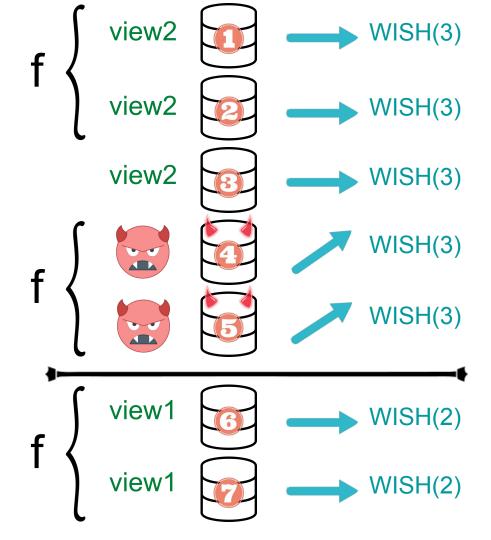
Faulty replicas send WISHes to 1~3, but not 6~7

Replicas 1~3
are partitioned from 6~7

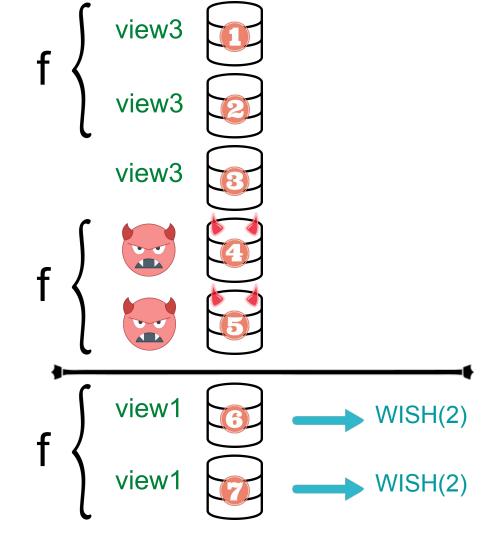




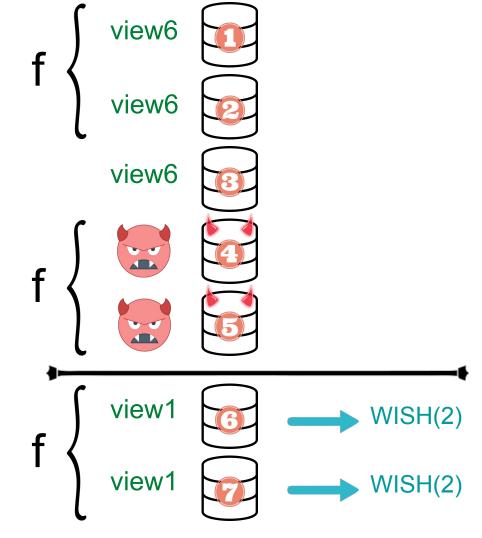




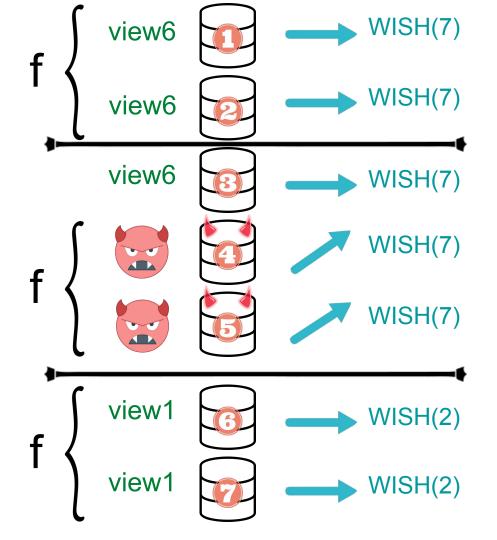






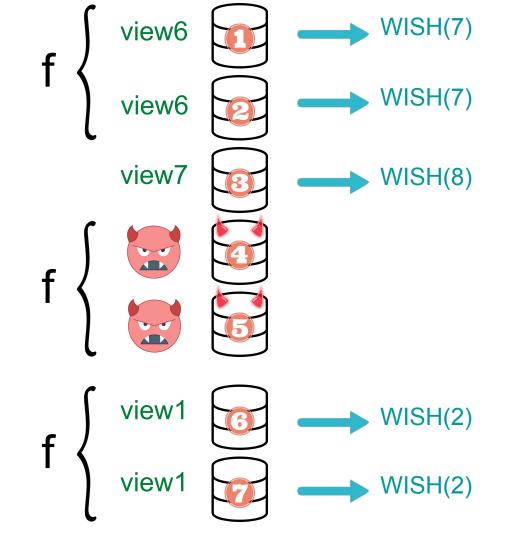






Faulty replicas send WISHes to replica 3 only



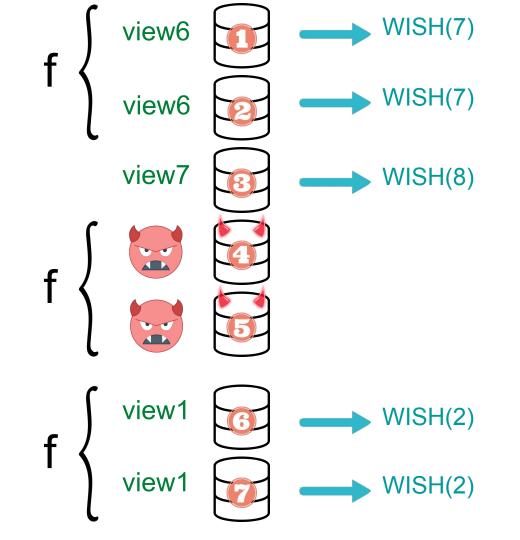


GST: WISHes get through

No matching f+1 WISHes

⇒ Bracha-like syncroniser
would get stuck

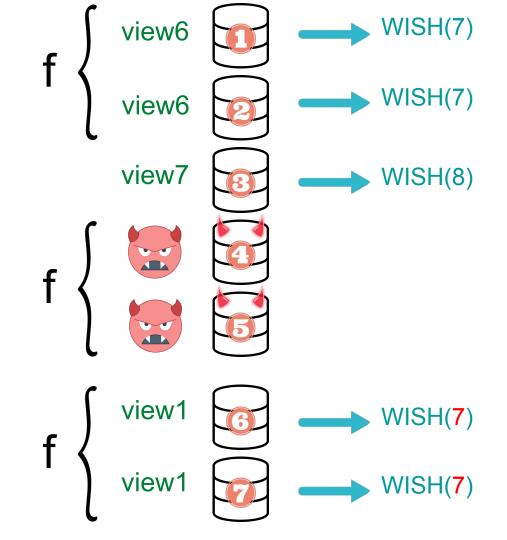




But can relay the minimal WISH of f+1

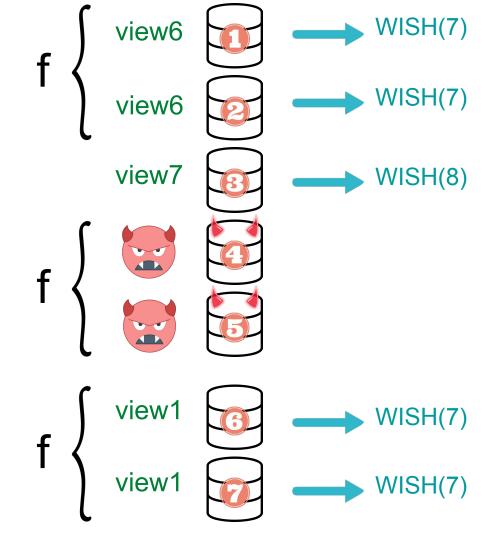
f+1







f+1

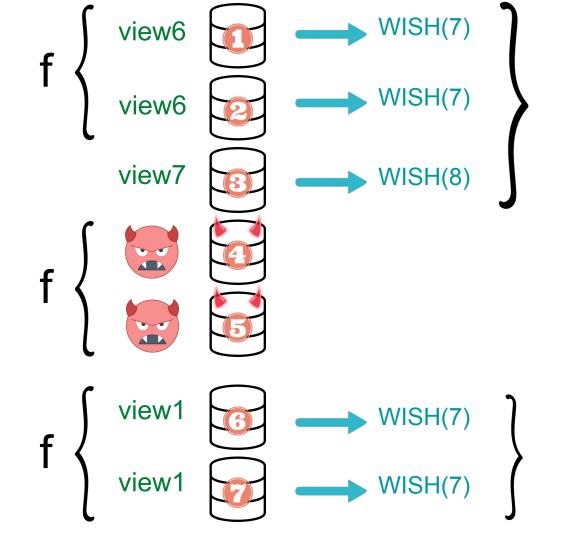


No matching 2f+1 WISHes

⇒ Bracha-like synchronizer

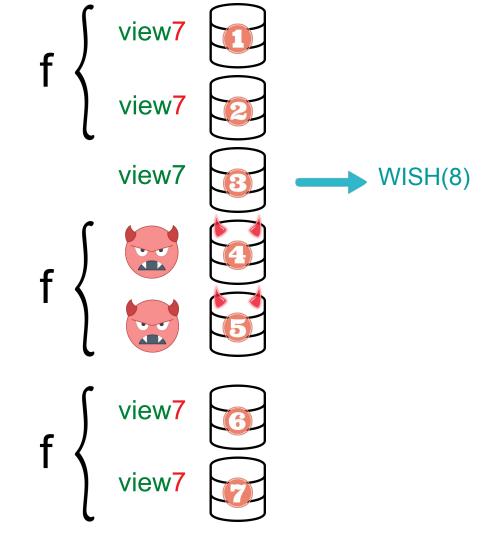
would get stuck again



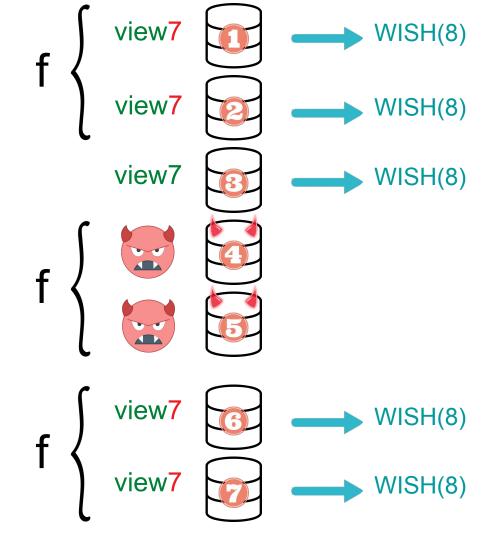


But can enter the view in the minimal WISH of 2f+1

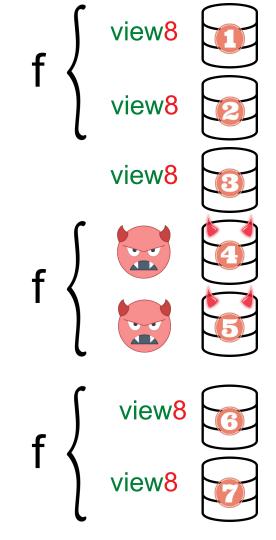












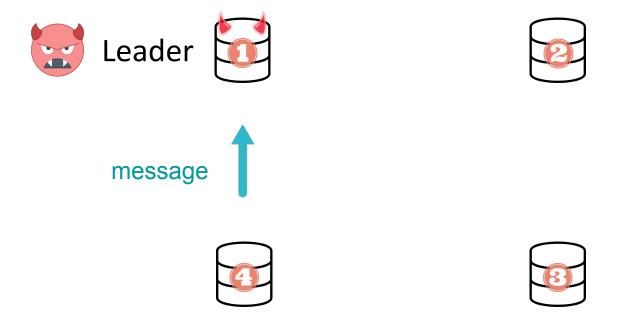
All correct replicas converge on the same view.

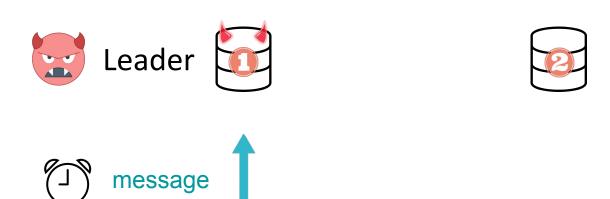


#### FastSync Synchronizer

- View synchronization mechanics hidden under the spec
- Can be proved liveness of different protocols
- FastSync + some of consensus protocols run in bounded space(PBFT, Hotstuff)
- Implementation has nice latency characteristics Allows establishing tight latency bounds for latency.

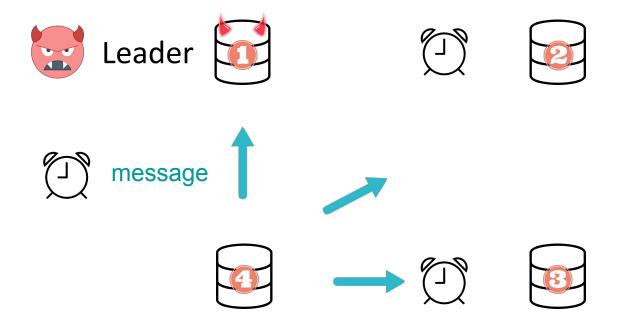












Leader













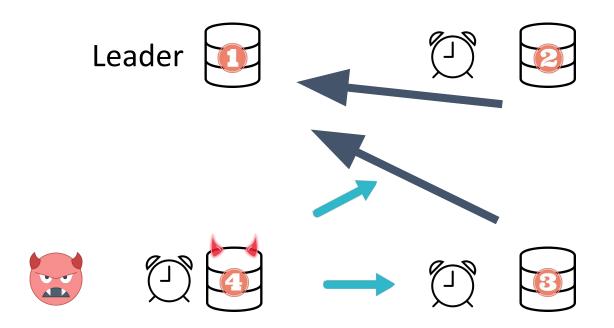




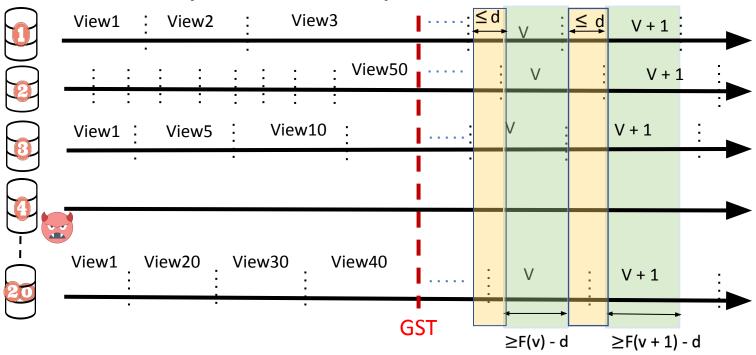








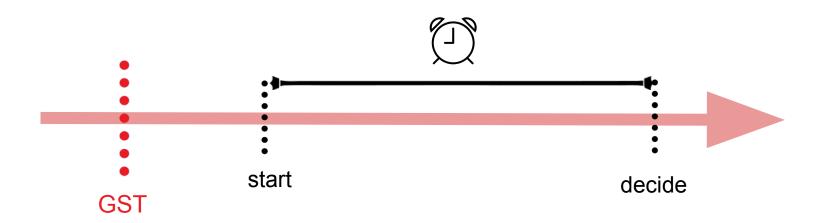
## Synchronizer Specification



• Non-faulty replicas enter each view within d (e.g.,  $d = 2\delta$ )

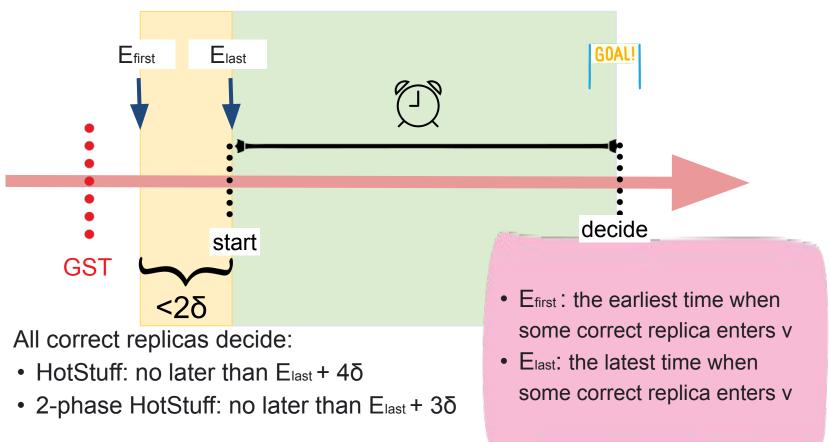


### **Latency Bounds for Consensus**

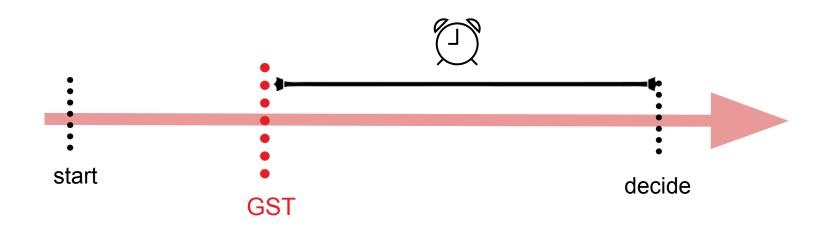


Under favorable conditions:
 if the protocol executes during a synchronous period, F(1) is
 big enough and leader(1) is correct

## Latency Bounds for Consensus



### **Latency Bounds for Consensus**



Under unfavorable conditions:
 If the protocol starts during an asynchronous period

# $v = GV(GST + \rho) + 1$ Leader view Efirst Elast start decide **GST**

Fp: When a replica enters a view where it is the leader, it sets a special timer for the duration determined by a

function Fp

All correct replicas decide:

HotStuff: no later than

$$GST + \rho + \sum_{k=v-1}^{v+f-1} (F(k) + \delta) + 7\delta$$

• 2-phase HotStuff: no later than  $GST + \rho + \sum_{k=0}^{\infty} (F(k) + \delta) + f_p(v + f) + 5\delta$ 

Latency Bounds for Consensus

# Thank you!

