




ByShard: Sharding in a Byzantine Environment

Hellings, J., & Sadoghi, M. (2021). ByShard. Proceedings of the VLDB Endowment, 14(11), 2230–2243.

October 30, 2023

Presented By: Amoolya, Manali, Mariana, Shravani

What is Sharding?



What is Sharding?

UC DAVIS STUDENT DATA



UC DAVIS STUDENT DATA SHARDED YEAR-WISE

1st YEAR

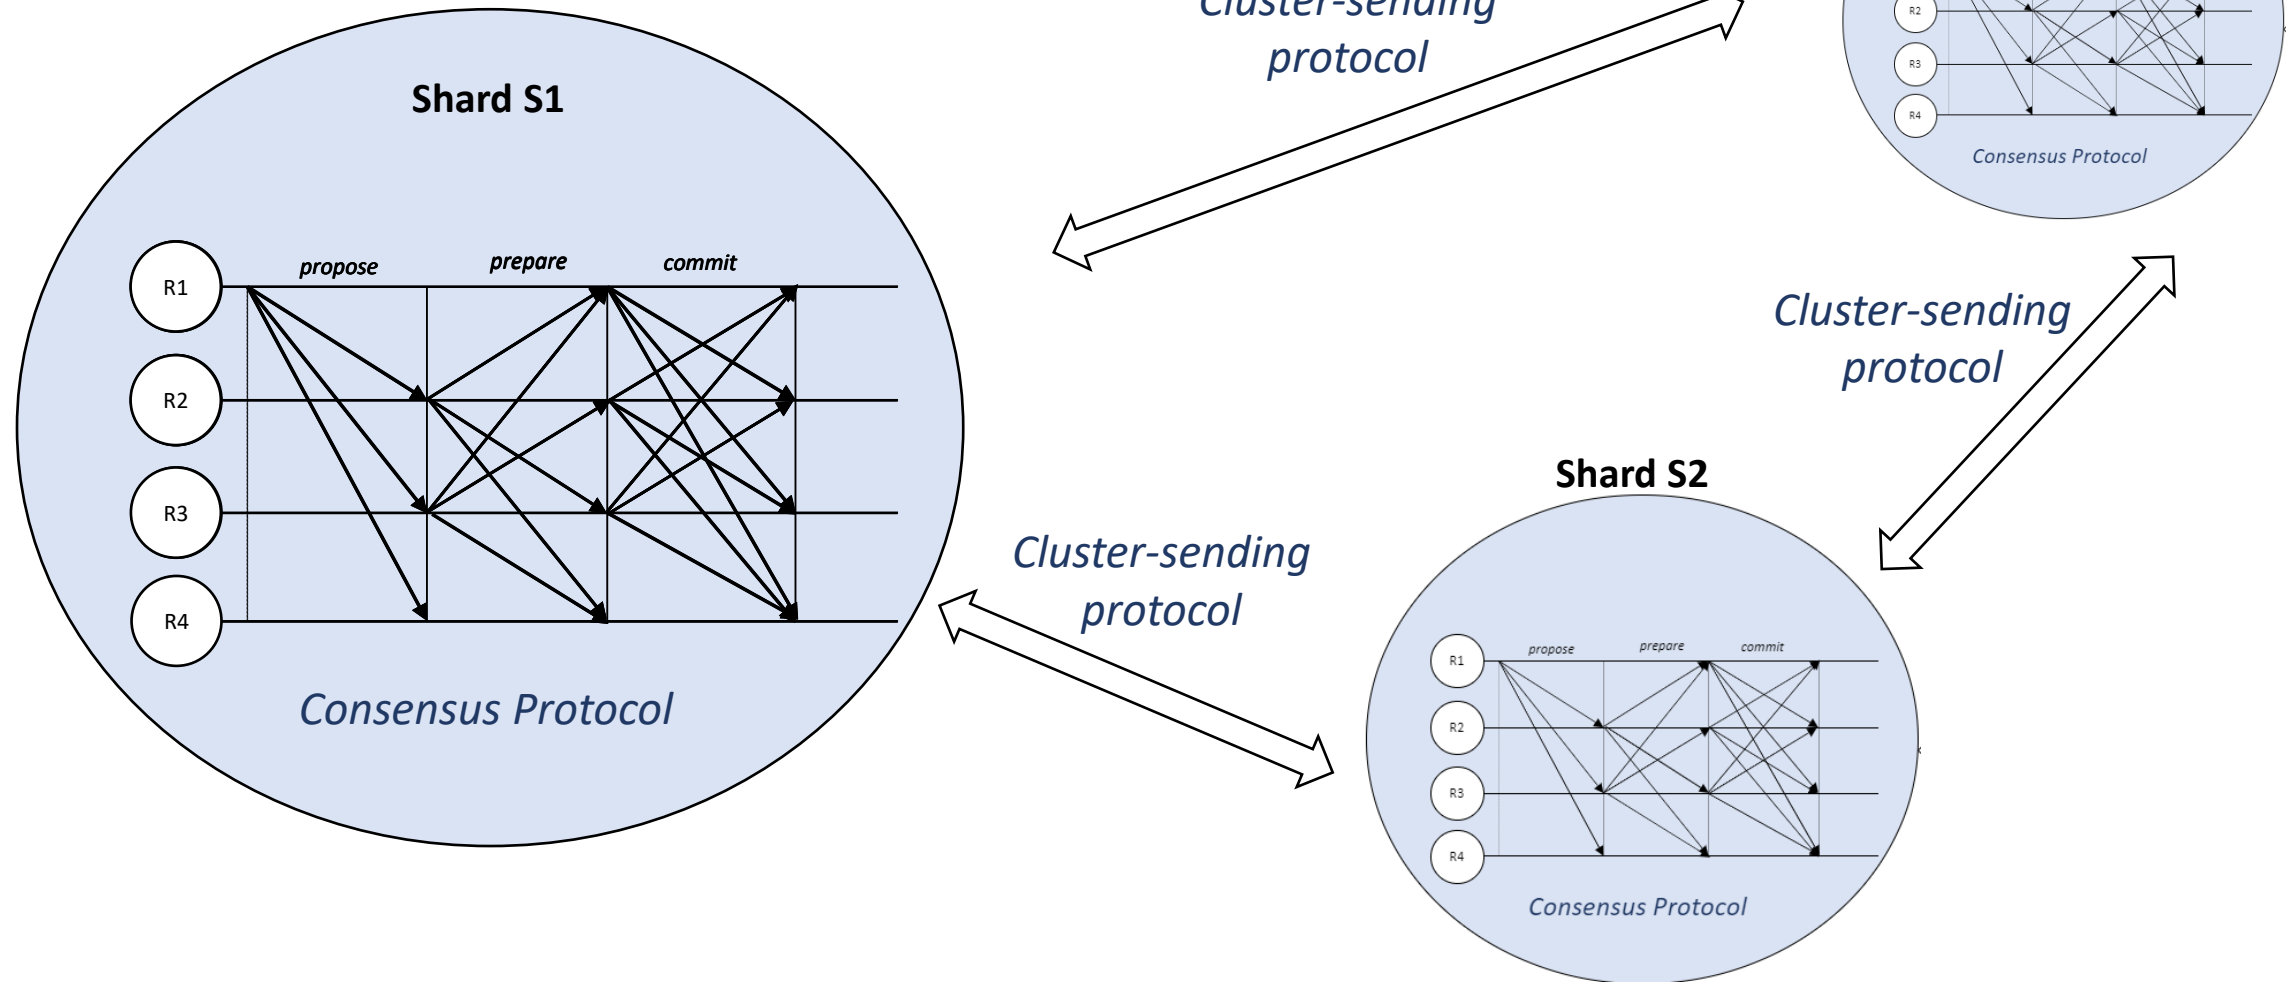
2nd YEAR

3rd YEAR

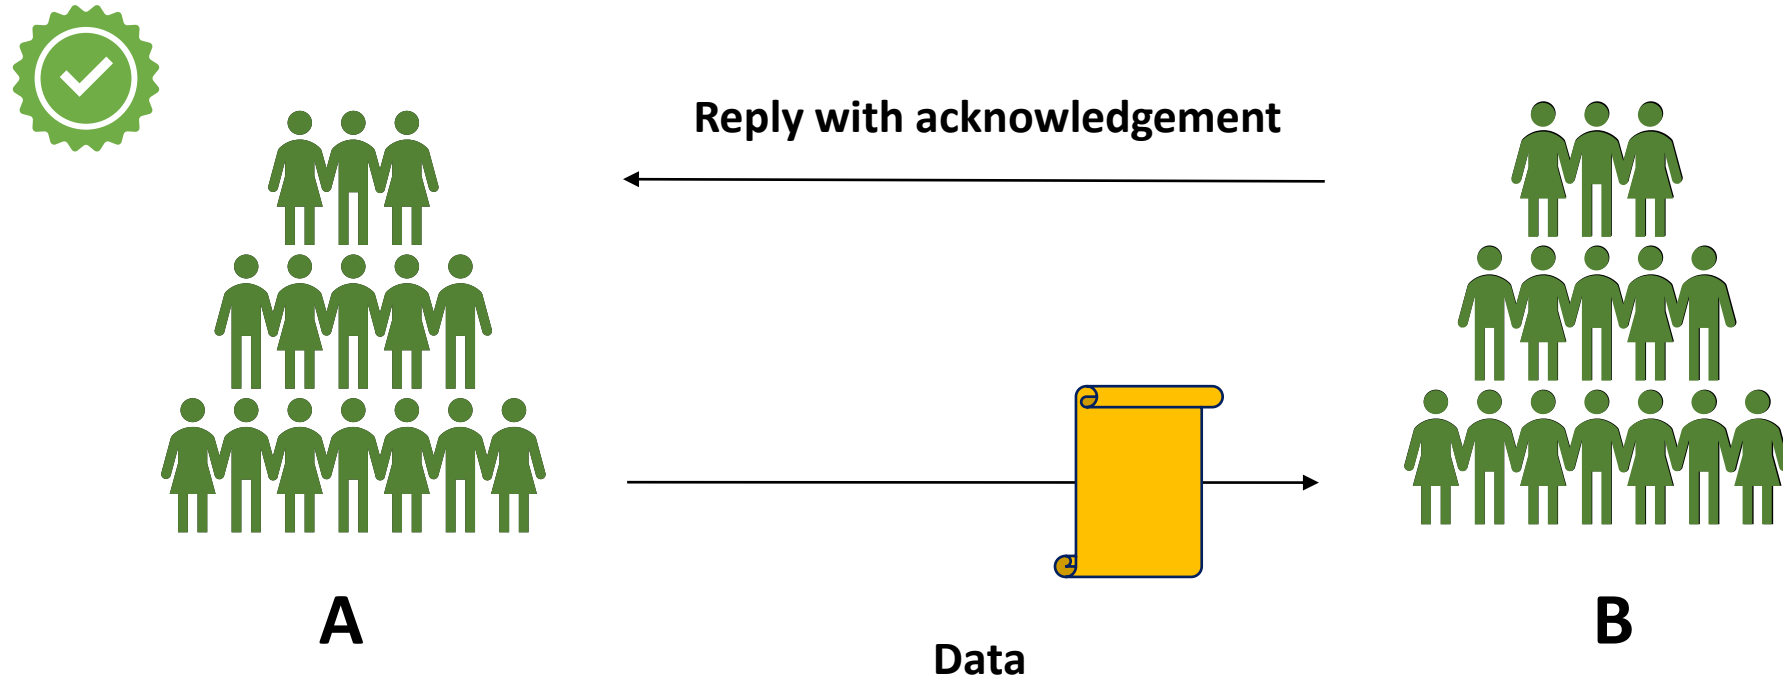
ByShard: A resilient sharding framework

- A base for general-purpose resilient sharded data management systems
 - Implemented principles of traditional distributed databases into sharded resilient systems
- Minimize the use of consensus

Overview of ByShard



Cluster Sending Protocol



Orchestrate-Execute Model



Orchestration

- Linear orchestration
- Centralized orchestration
- Distributed orchestration

Execution


- Isolated-free direct execution
- Lock-based execution

Orchestrate-Execute Model

This model reduces the number of consensus steps involved in each transaction.

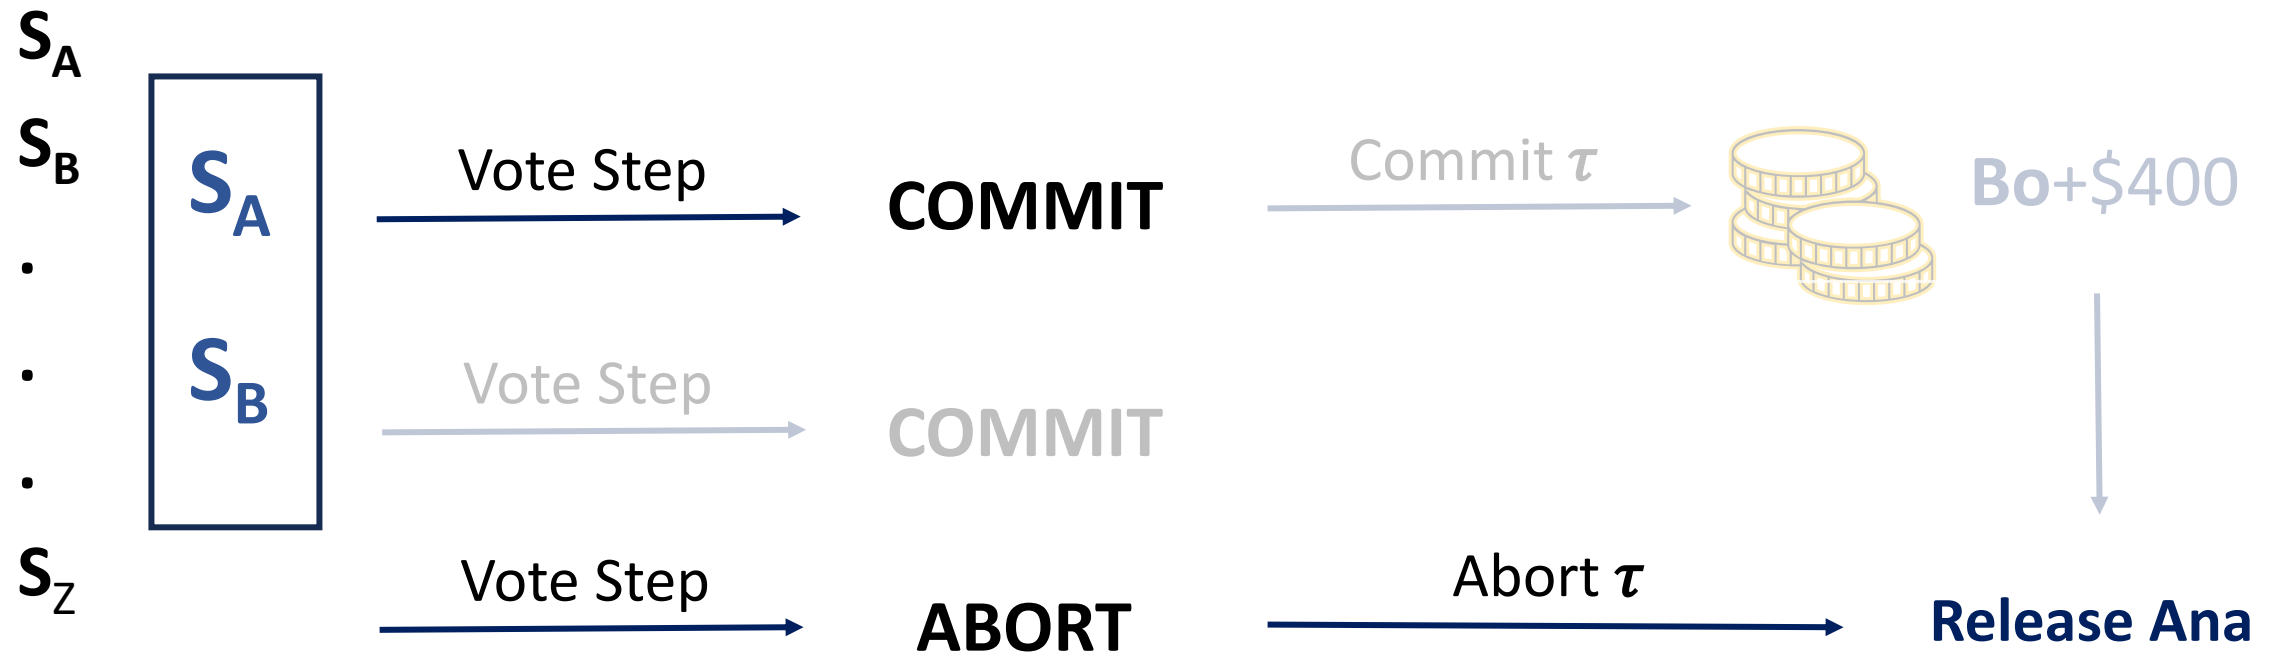
Ex: If Ana has at least \$500 and Bo at least \$200 then transfer \$200 from Ana and \$100 from Bo to Elisa.

- Vote Step
- Commit Step
- Abort Step



Sharded system - Example

τ = "if Ana has atleast \$500 and Bo has atleast \$200, then move \$400 from Ana to Bo"



Important Terminologies

τ - Transaction

nv - Vote Steps

na - Abort Steps


nc - Commit Steps

$\{S_1, \dots, S_n\}$ - Total no. Of Shards

$\text{shards}(\tau)$ - Shards involved in the transaction

S_r - Root Shard

Linear Orchestration



For transaction τ ,


Consensus Steps:

$$nv + nc$$

Cluster Sending:

$$nv + nc - 1$$

Centralized Orchestration



For transaction τ ,

- Consensus Steps:

$$nv + nc + 1$$

- Cluster Sending Steps:

$$2(nv - 1) + nc$$

Important Terminologies

τ - Transaction

nv - Vote Steps

na - Abort Steps


nc - Commit Steps

$\{S_1, \dots, S_n\}$ - Total no. of Shards

$\text{shards}(\tau)$ - Shards involved in the transaction

S_r - Root Shard

Distributed Orchestration



For transaction τ ,

Consensus Steps:

$$nv + nc$$

Cluster Sending:

$$nv (na + nc) + (nv - 1)$$

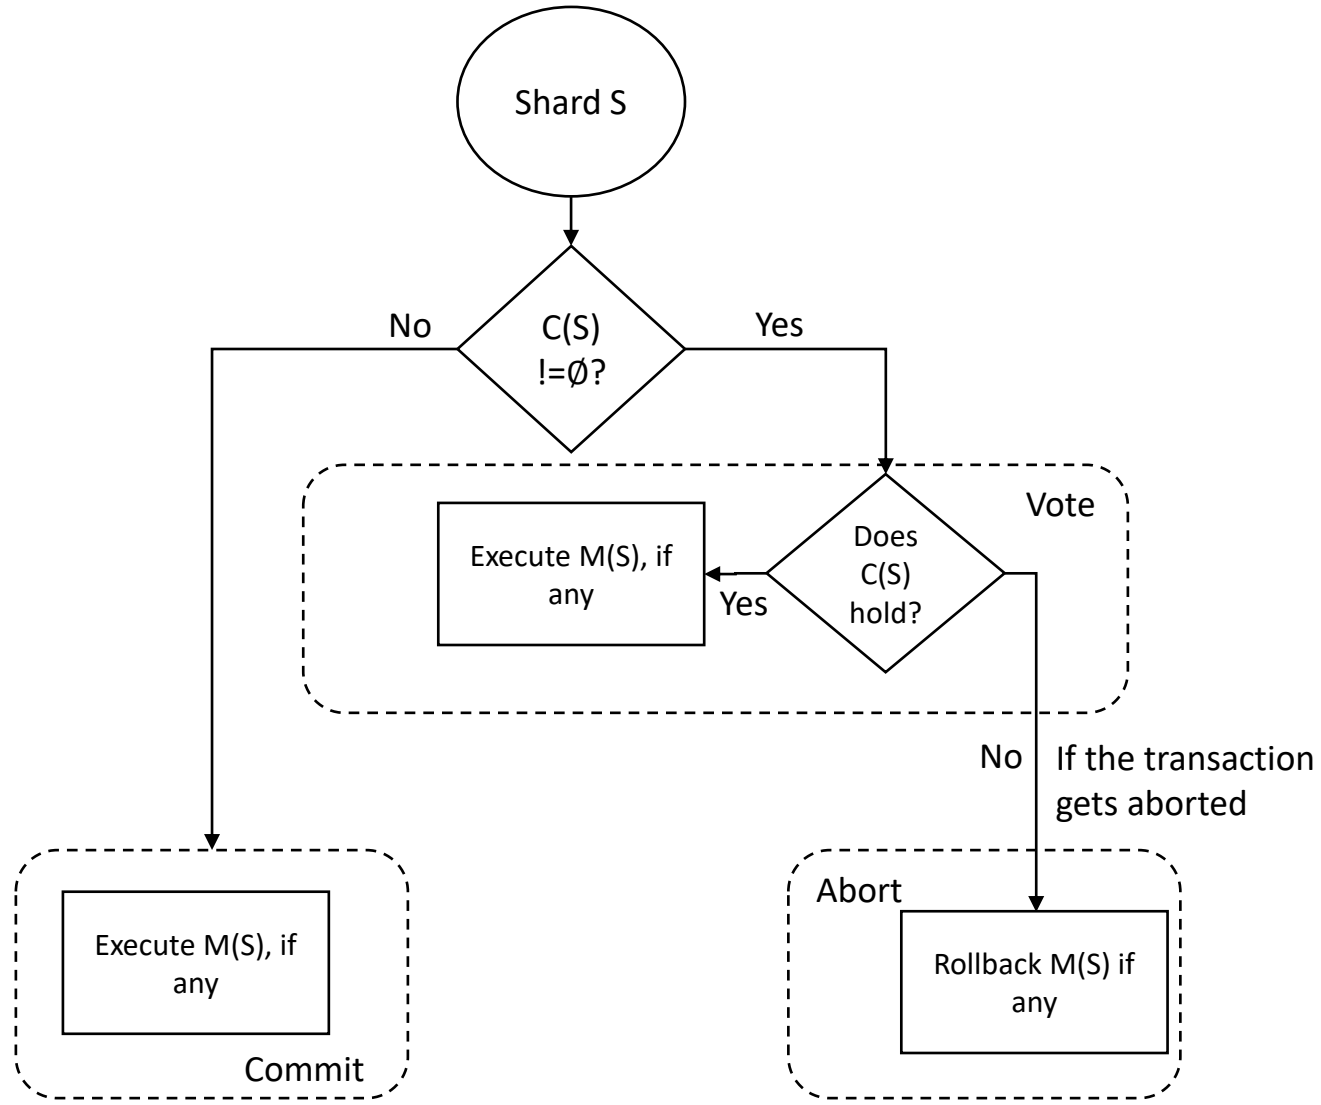
Orchestration Summary

Orchestration	Consecutive Consensus Steps	Consensus Steps ●	Cluster Sending Steps →
Linear	$nv+1$	$nv + nc$	$nv + nc - 1$
Centralized	4	$nv + nc + 1$	$2(nv - 1) + nc$
Distributed	3	$nv + nc$	$nv (na + nc) + (nv - 1)$

Execution in OEM

- Constraints
 $\text{CON}(X, y) = \text{“the balance of } X \text{ is at least } y\text{”}$
- Modifications
 $\text{MOD}(X, y) = \text{“add } y \text{ to the balance of } X\text{”}$

Isolated-free direct execution

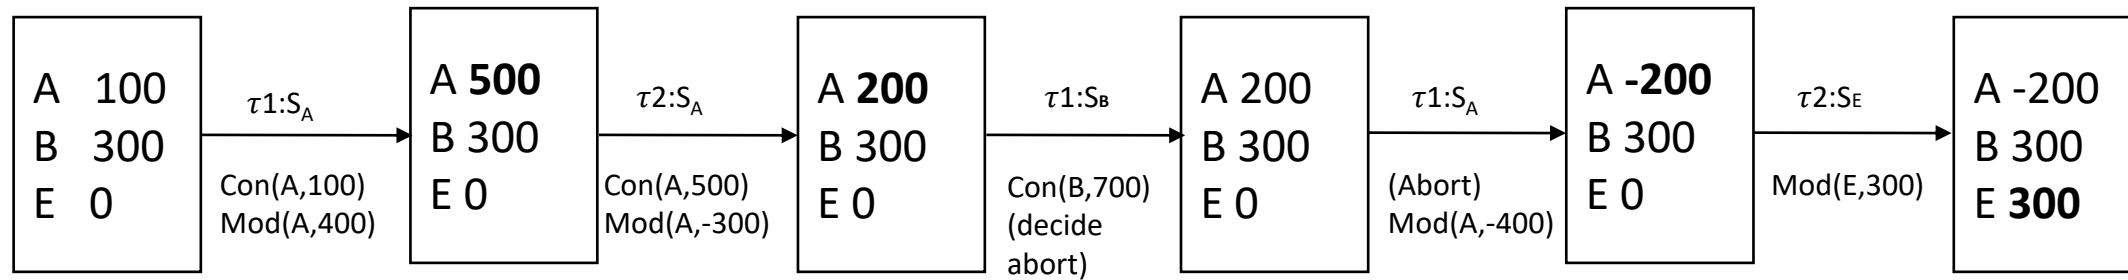


Isolated-free direct execution

Provides degree 0 isolation which can lead to violations of constraints.

$\tau_1 = \text{Con}(A, 100), \text{Con}(B, 700), \text{Mod}(A, 400), \text{Mod}(B, -400)$

$\tau_2 = \text{Con}(A, 500), \text{Mod}(A, -300), \text{Mod}(E, 300)$





Isolated-free direct execution

- Safe – Rolling back $M(X, y)$ with $y \leq 0$ (removal) \Rightarrow Increases the balance of X and X will never be negative \rightarrow Vote Step
- Unsafe – Rolling back $M(X, y)$ with $y \geq 0$ (addition) \Rightarrow Decreases the balance of X and X can be negative \rightarrow Commit Step
- Worst case: 2 shard steps when committing.

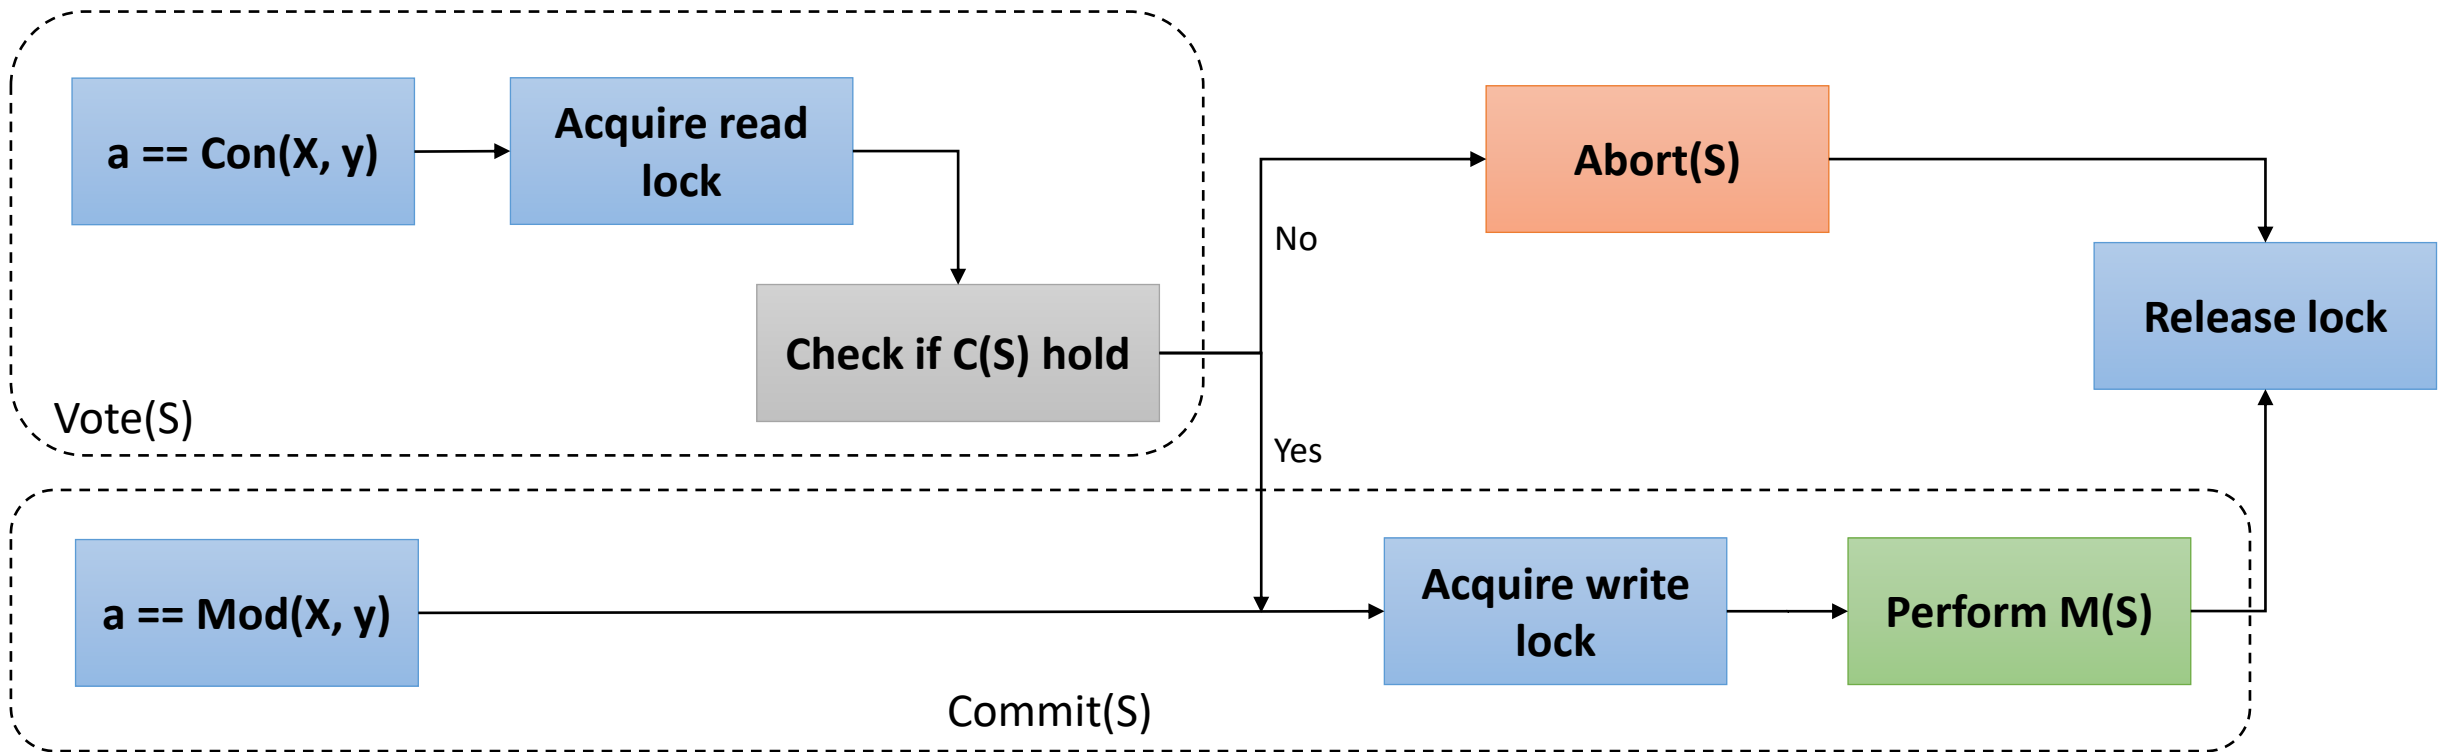
Lock-based execution

- Two-phase locking
- Possibility of higher isolation – degree 3




Lock-based execution

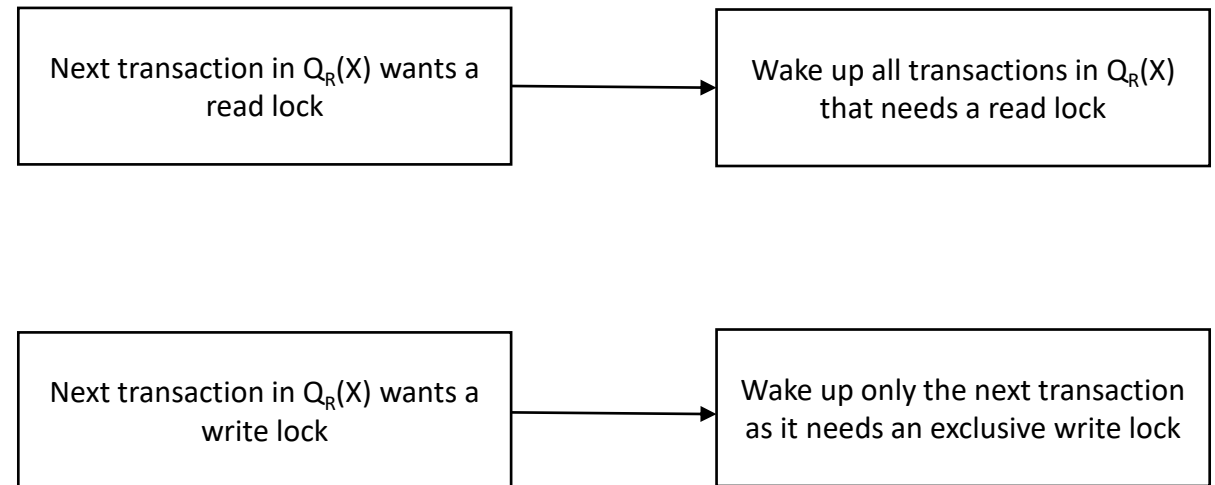
- Let $S \in \text{Shards}(\tau)$
- $\text{Accounts}(S) = \{a \mid \text{Con}(X, y) \in C(S) \vee \text{Mod}(X, y) \in M(S)\}$



Lock-based execution



After the existing lock is released



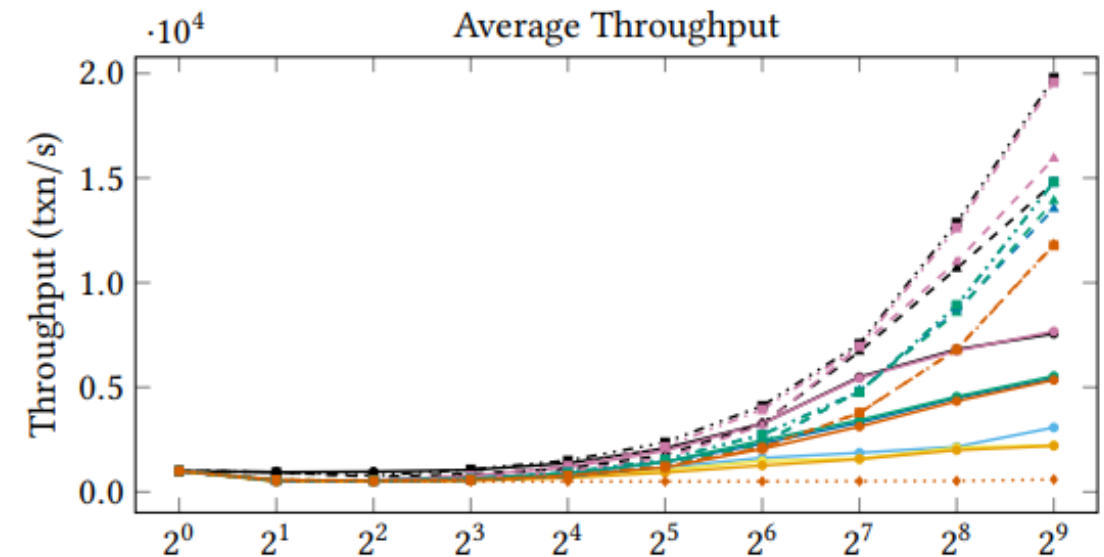
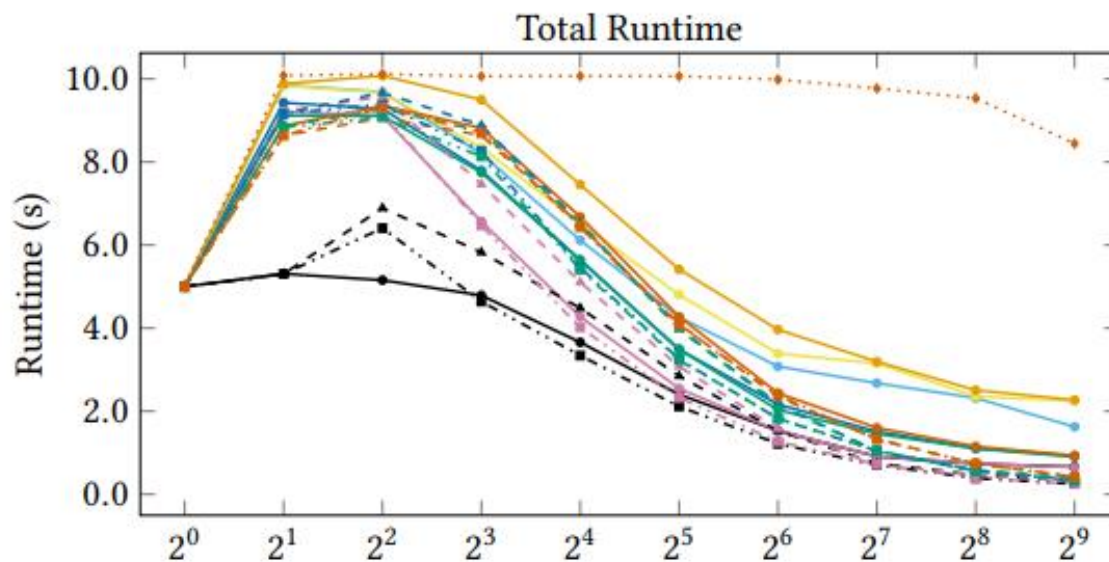
- Hence vote-step can obtain all its locks in only a single consensus step. Wake up step is deterministic. No extra consensus are necessary to determine which of the next transactions will be executed.

Lock-based execution

- Pro: Provides serializability and highest degree of isolation
- Con: Large transaction processing latency whenever contention is high
- Read uncommitted execution
 - Degree 1 isolation
 - No read locks on any data item
 - Very useful for read-heavy workloads
- Read committed execution
 - Degree 2 isolation
 - Read locks are released directly after reading an item
 - Minimizes the time read locks are held

Performance Evaluation

	Isolation-Free execution		Lock-based execution						AHL (reference committee)
	(write uncommitted)		Read Uncommitted		Read Committed		Serializable		
	<i>unsafe</i>	<i>safe</i>	<i>blocking</i>	<i>non-blocking</i>	<i>blocking</i>	<i>non-blocking</i>	<i>blocking</i>	<i>non-blocking</i>	
Linear	—●— LIFU	—●— LIFs	—●— LRUB	—●— LRUNB	—●— LRCB	—●— LRCNB	—●— LSB	—●— LSNB	
Centralized	-▲- CIFU	-▲- CIFs	-▲- CRUNB	-▲- CRUNB	-▲- CRCNB	-▲- CRCNB	-▲- CSNB	-▲- CSNB	
Distributed	-■- DIFU	-■- DIFs	-■- DRUNB	-■- DRUNB	-■- DRCNB	-■- DRCNB	-■- DSNB	-■- DSNB	



* Graphs taken from the paper figure 6.1

Conclusion

- ByShard - a general-purpose framework for sharded resilient data management systems
- Orchestrate-Execute Model (OEM)
- 18 multi-shard transaction processing protocols with tradeoffs between throughput, latency, isolation level and abort rate

Thank you