

# **RCC: Resilient Concurrent Consensus for High-Throughput Secure Transaction Processing**

**Authors: Suyash Gupta, Jelle Hellings, and Mohammad Sadoghi**

**Presenters:**

**Naitik Jain, Tanvi Mehta, Shivani Suryawanshi, Krithika Naidu**

# Contents

- Overview
- Traditional Consensus: PBFT
- Limitations of Traditional Consensus
- Promise of Concurrent Consensus
- Solution
- Resilient Concurrent Consensus (RCC)
- Design of RCC with Example
- Dealing with Detectable Failures
- Undetectable Failures
- Client Interactions with RCC
- RCC Improving Resilience of Consensus
- Evaluation of the Performance of RCC
- Conclusion
- References




# Overview

- Consensus-based database systems are emerged, offering resilience, strong data provenance, and federated data management.
- Rely on Primary-backup protocols.
- RCC enhances the throughput, resilience, and security by enabling concurrent consensus.


# Overview

At the core of any blockchain application or distributed database system, there is Byzantine Fault-Tolerant (BFT) consensus protocol.




# Overview

At the core of any blockchain application or distributed database system, there is Byzantine Fault-Tolerant (BFT) consensus protocol.




# Traditional Consensus: PBFT




Consensus among  $n$  replicas and atmost  $f$  byzantine  $\longrightarrow n \geq 3f+1$

# Traditional Consensus: PBFT



$n = 4$  replicas and  $f \leq 1$

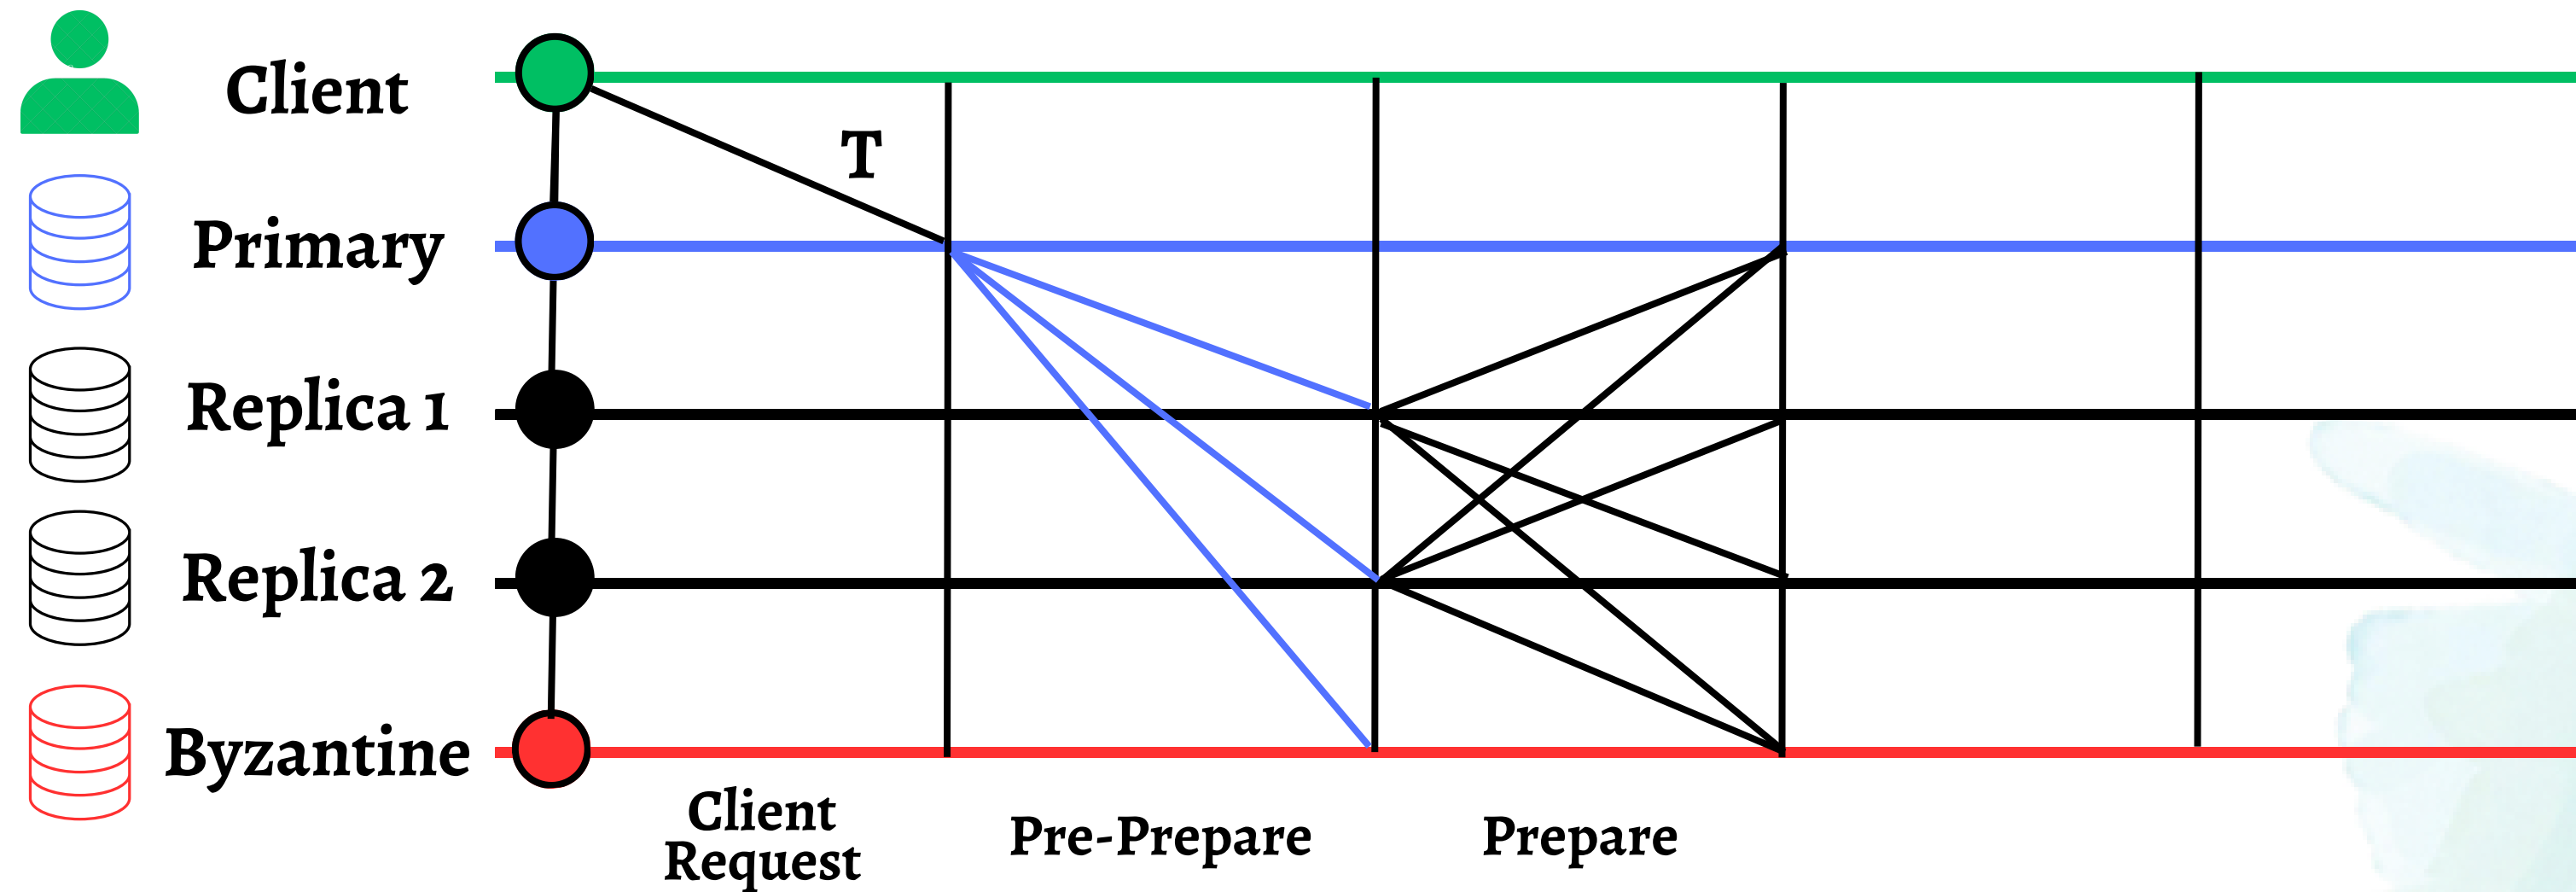
# Traditional Consensus: PBFT



$n = 4$  replicas and  $f \leq 1$

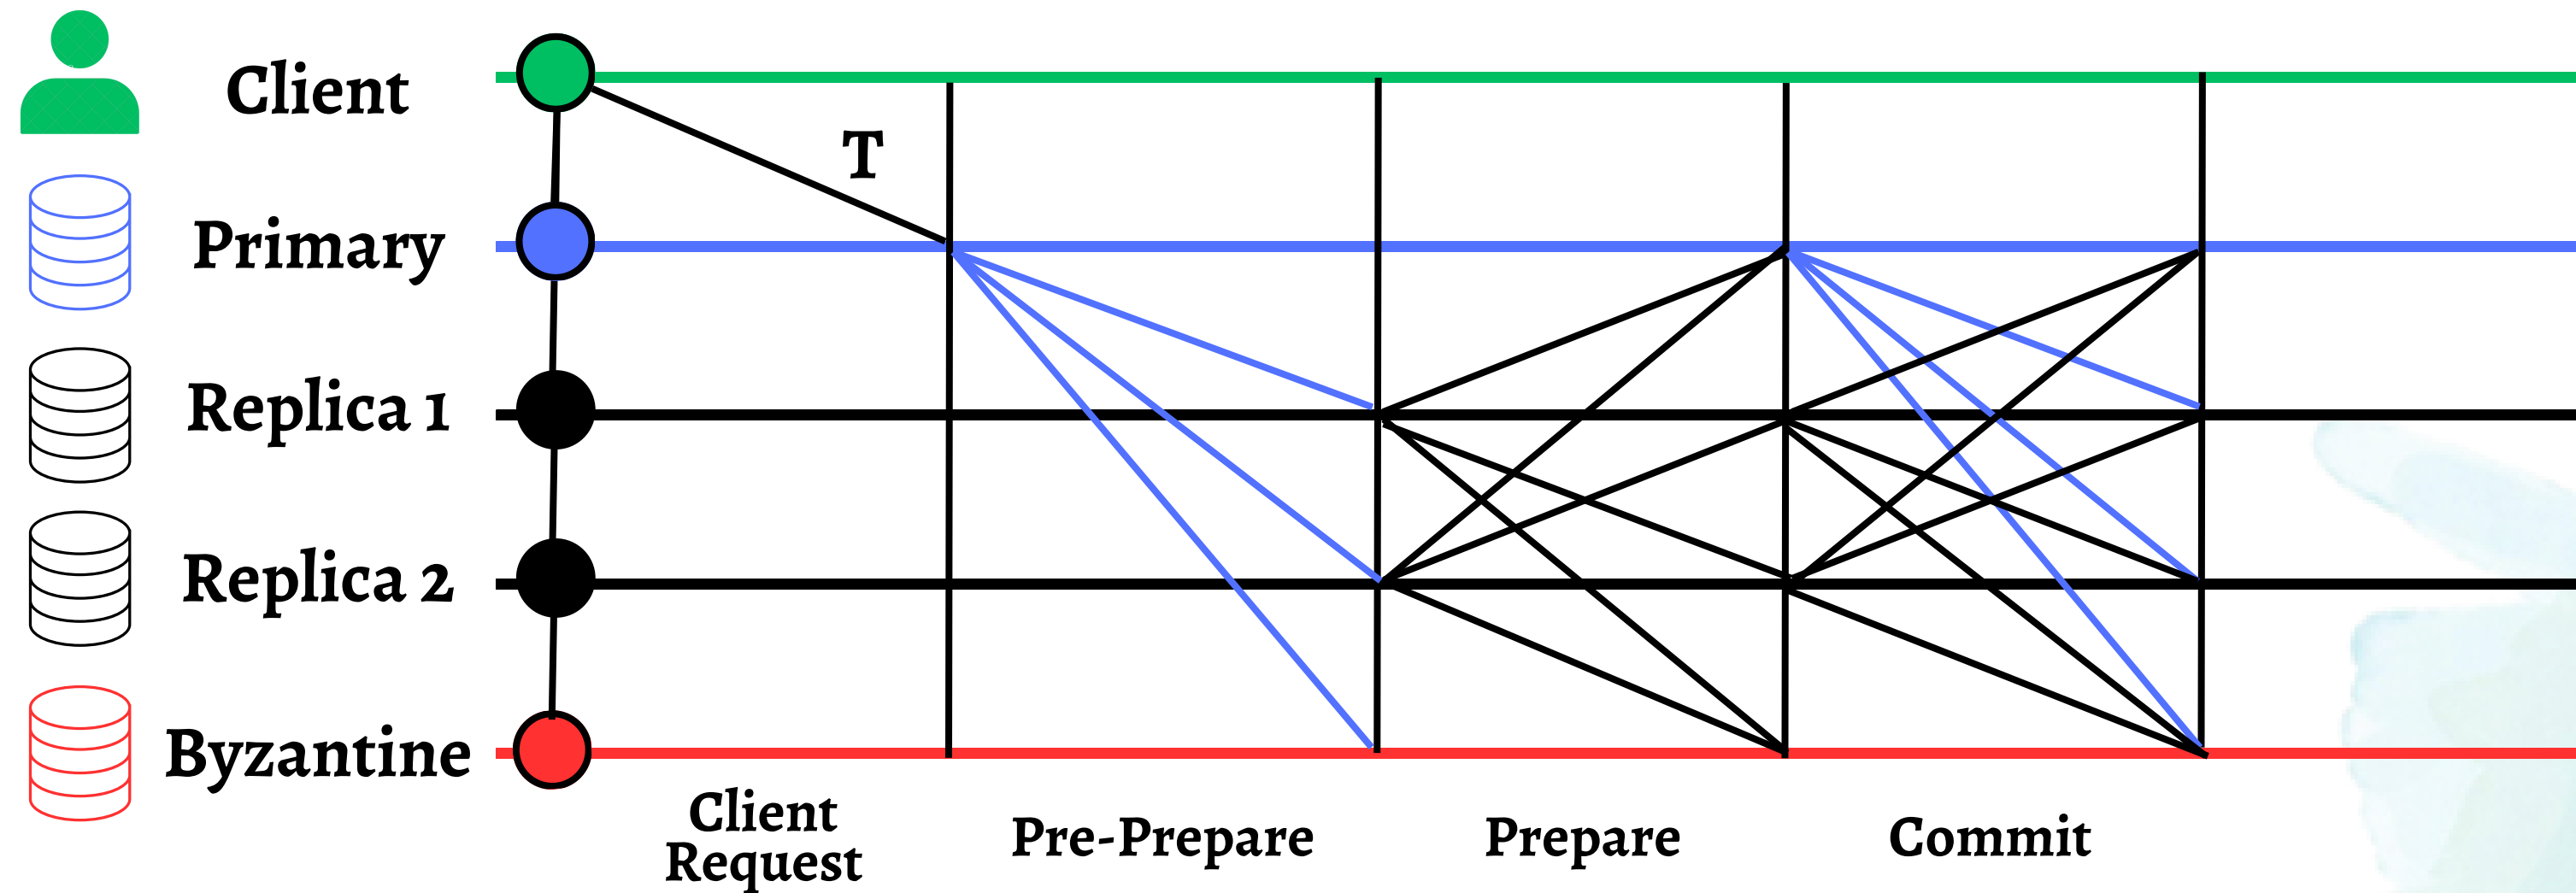


# Traditional Consensus: PBFT




$n = 4$  replicas and  $f \leq 1$

# Traditional Consensus: PBFT




$n = 4$  replicas and  $f \leq 1$

# Traditional Consensus: PBFT



# Traditional Consensus: PBFT



# Limitations

- Traditional primary-backup consensus protocols underutilize network resources and thus prevents maximization of transaction throughput, which is determined by the outgoing bandwidth of the primary.
- Considering a system,

$$T_{max} = \frac{B}{(n-1)st} \quad T_{PBFT} = \frac{B}{((n-1)(st + 3sm))}$$

$st$  is the size of each transaction  $sm$  is the size of each message

# Limitations

$$T_{PBFT} \approx T_{max} \text{ when } st \gg sm$$

- Replicas are underutilized: primaries must send  $(n-1)st$  while replicas only have to send and receive  $st$  bytes roughly, given that  $st \gg sm$ .
- Underutilization of non-primary replicas in comparison to primaries!



# Promise of Concurrent Consensus

- **Democracy** - Give all the replicas the power to be the primary.
- **Parallelism** - Run multiple parallel instances of a BFT protocol.
- **Decentralization** - Always there will be a set of ordered client requests.
- HotStuff balances load by consistently switching primaries, but does not address core issue of underutilization of resources.
- Concurrent consensus involves proposing at least  **$nf$  transactions** at a time.

$$T_{cmax} = nf \frac{B}{((n-1)st + (nf-1)st)}$$

# Solution

- Concurrent consensus can achieve higher levels of throughput by efficiently utilizing all available replicas.
- Resilient Concurrent Consensus (RCC) is a paradigm for transforming any primary-backup consensus protocol into a concurrent consensus protocol with increased throughput.



# Resilient Concurrent Consensus (RCC)

- To deal with underutilisation of resources and low throughput
- Turn any primary backup consensus protocol into a concurrent consensus protocol

## Design Goals:

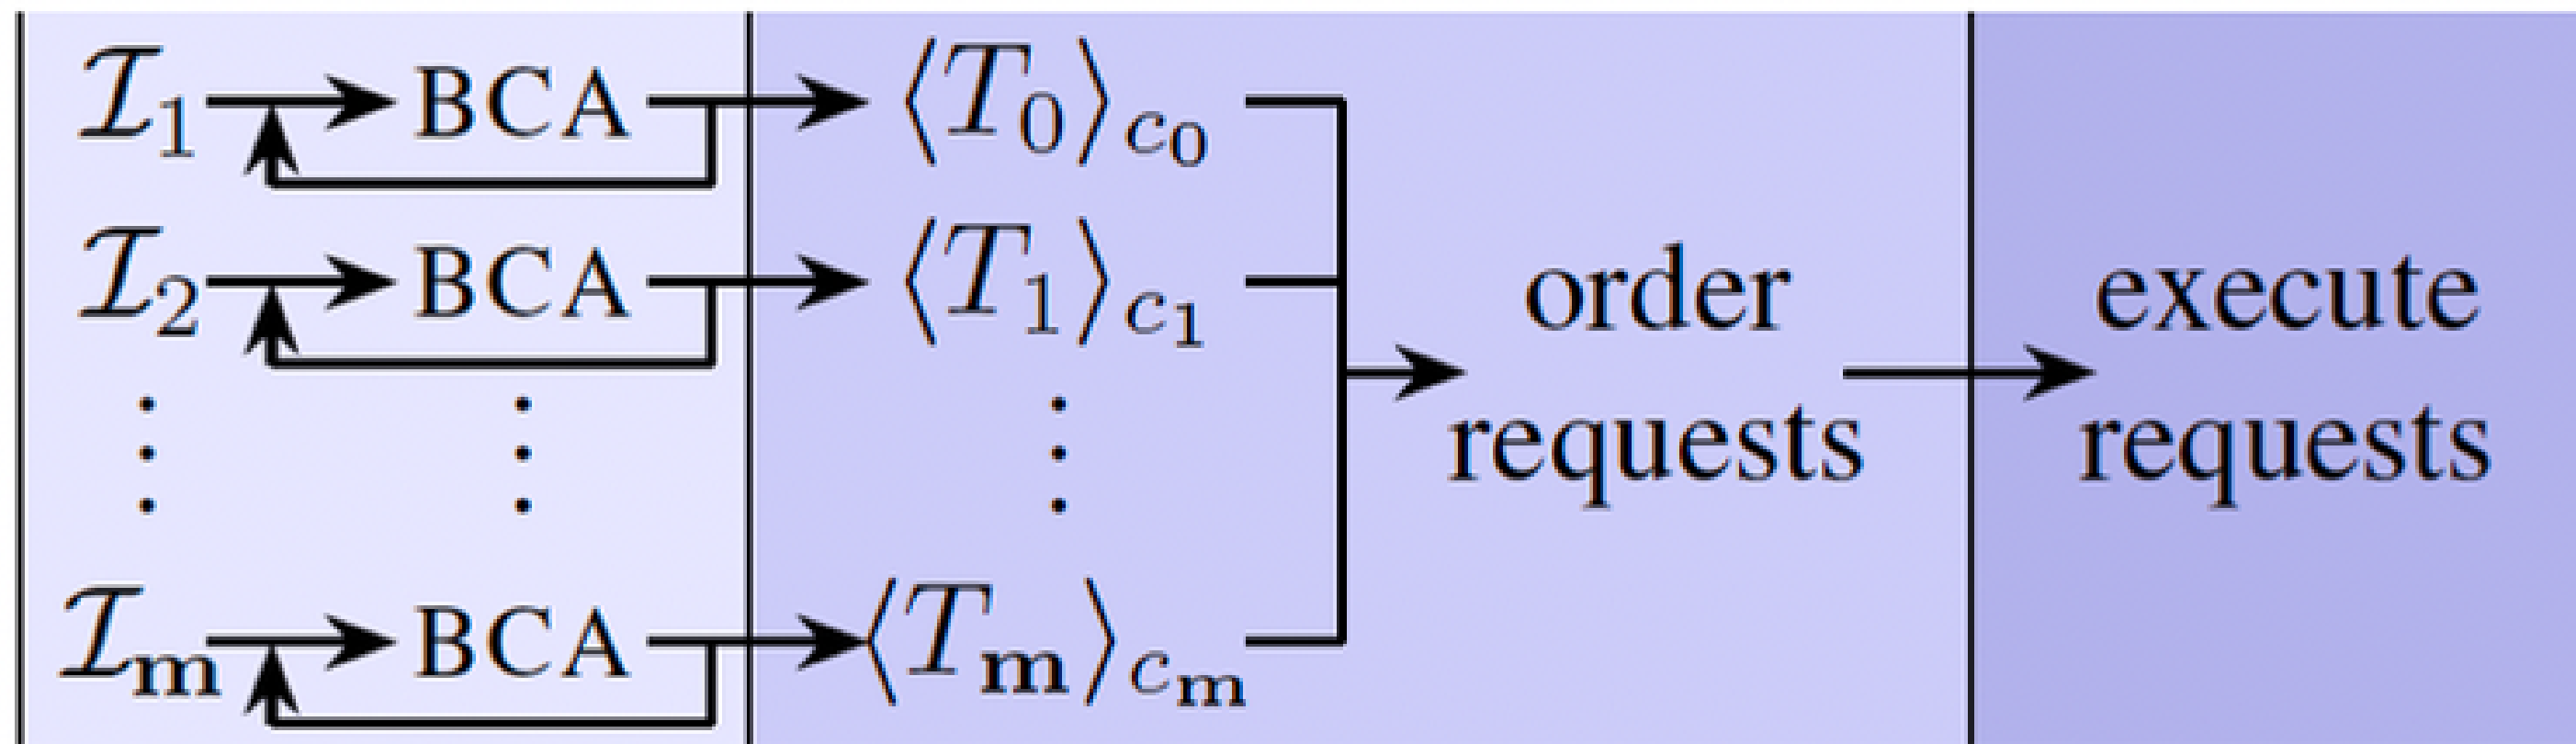
- Provide consensus among replicas on client transactions that are to be executed
- Applied to any other protocols
- Dealing with faulty primaries does not interfere with the operations of other consensus-instances

# Design of RCC

**Concurrent BCA**


**Ordering**

**Execution**

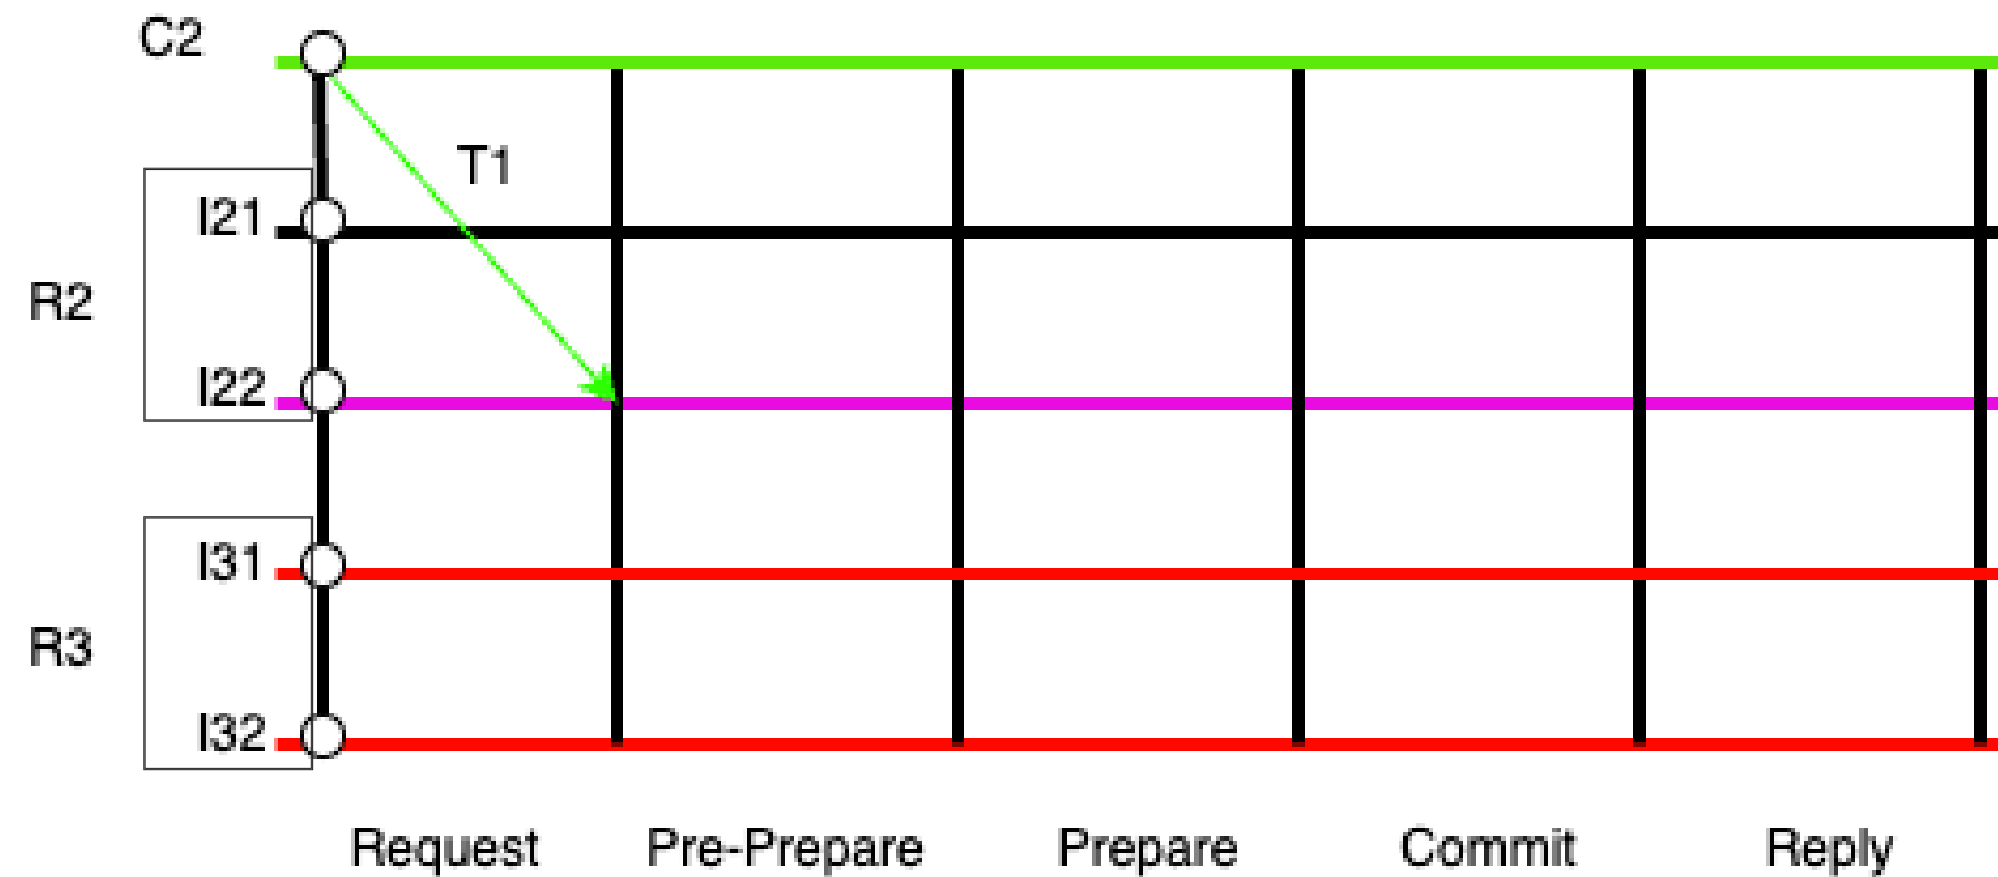


# Example


Primary Replica




Primary Replica



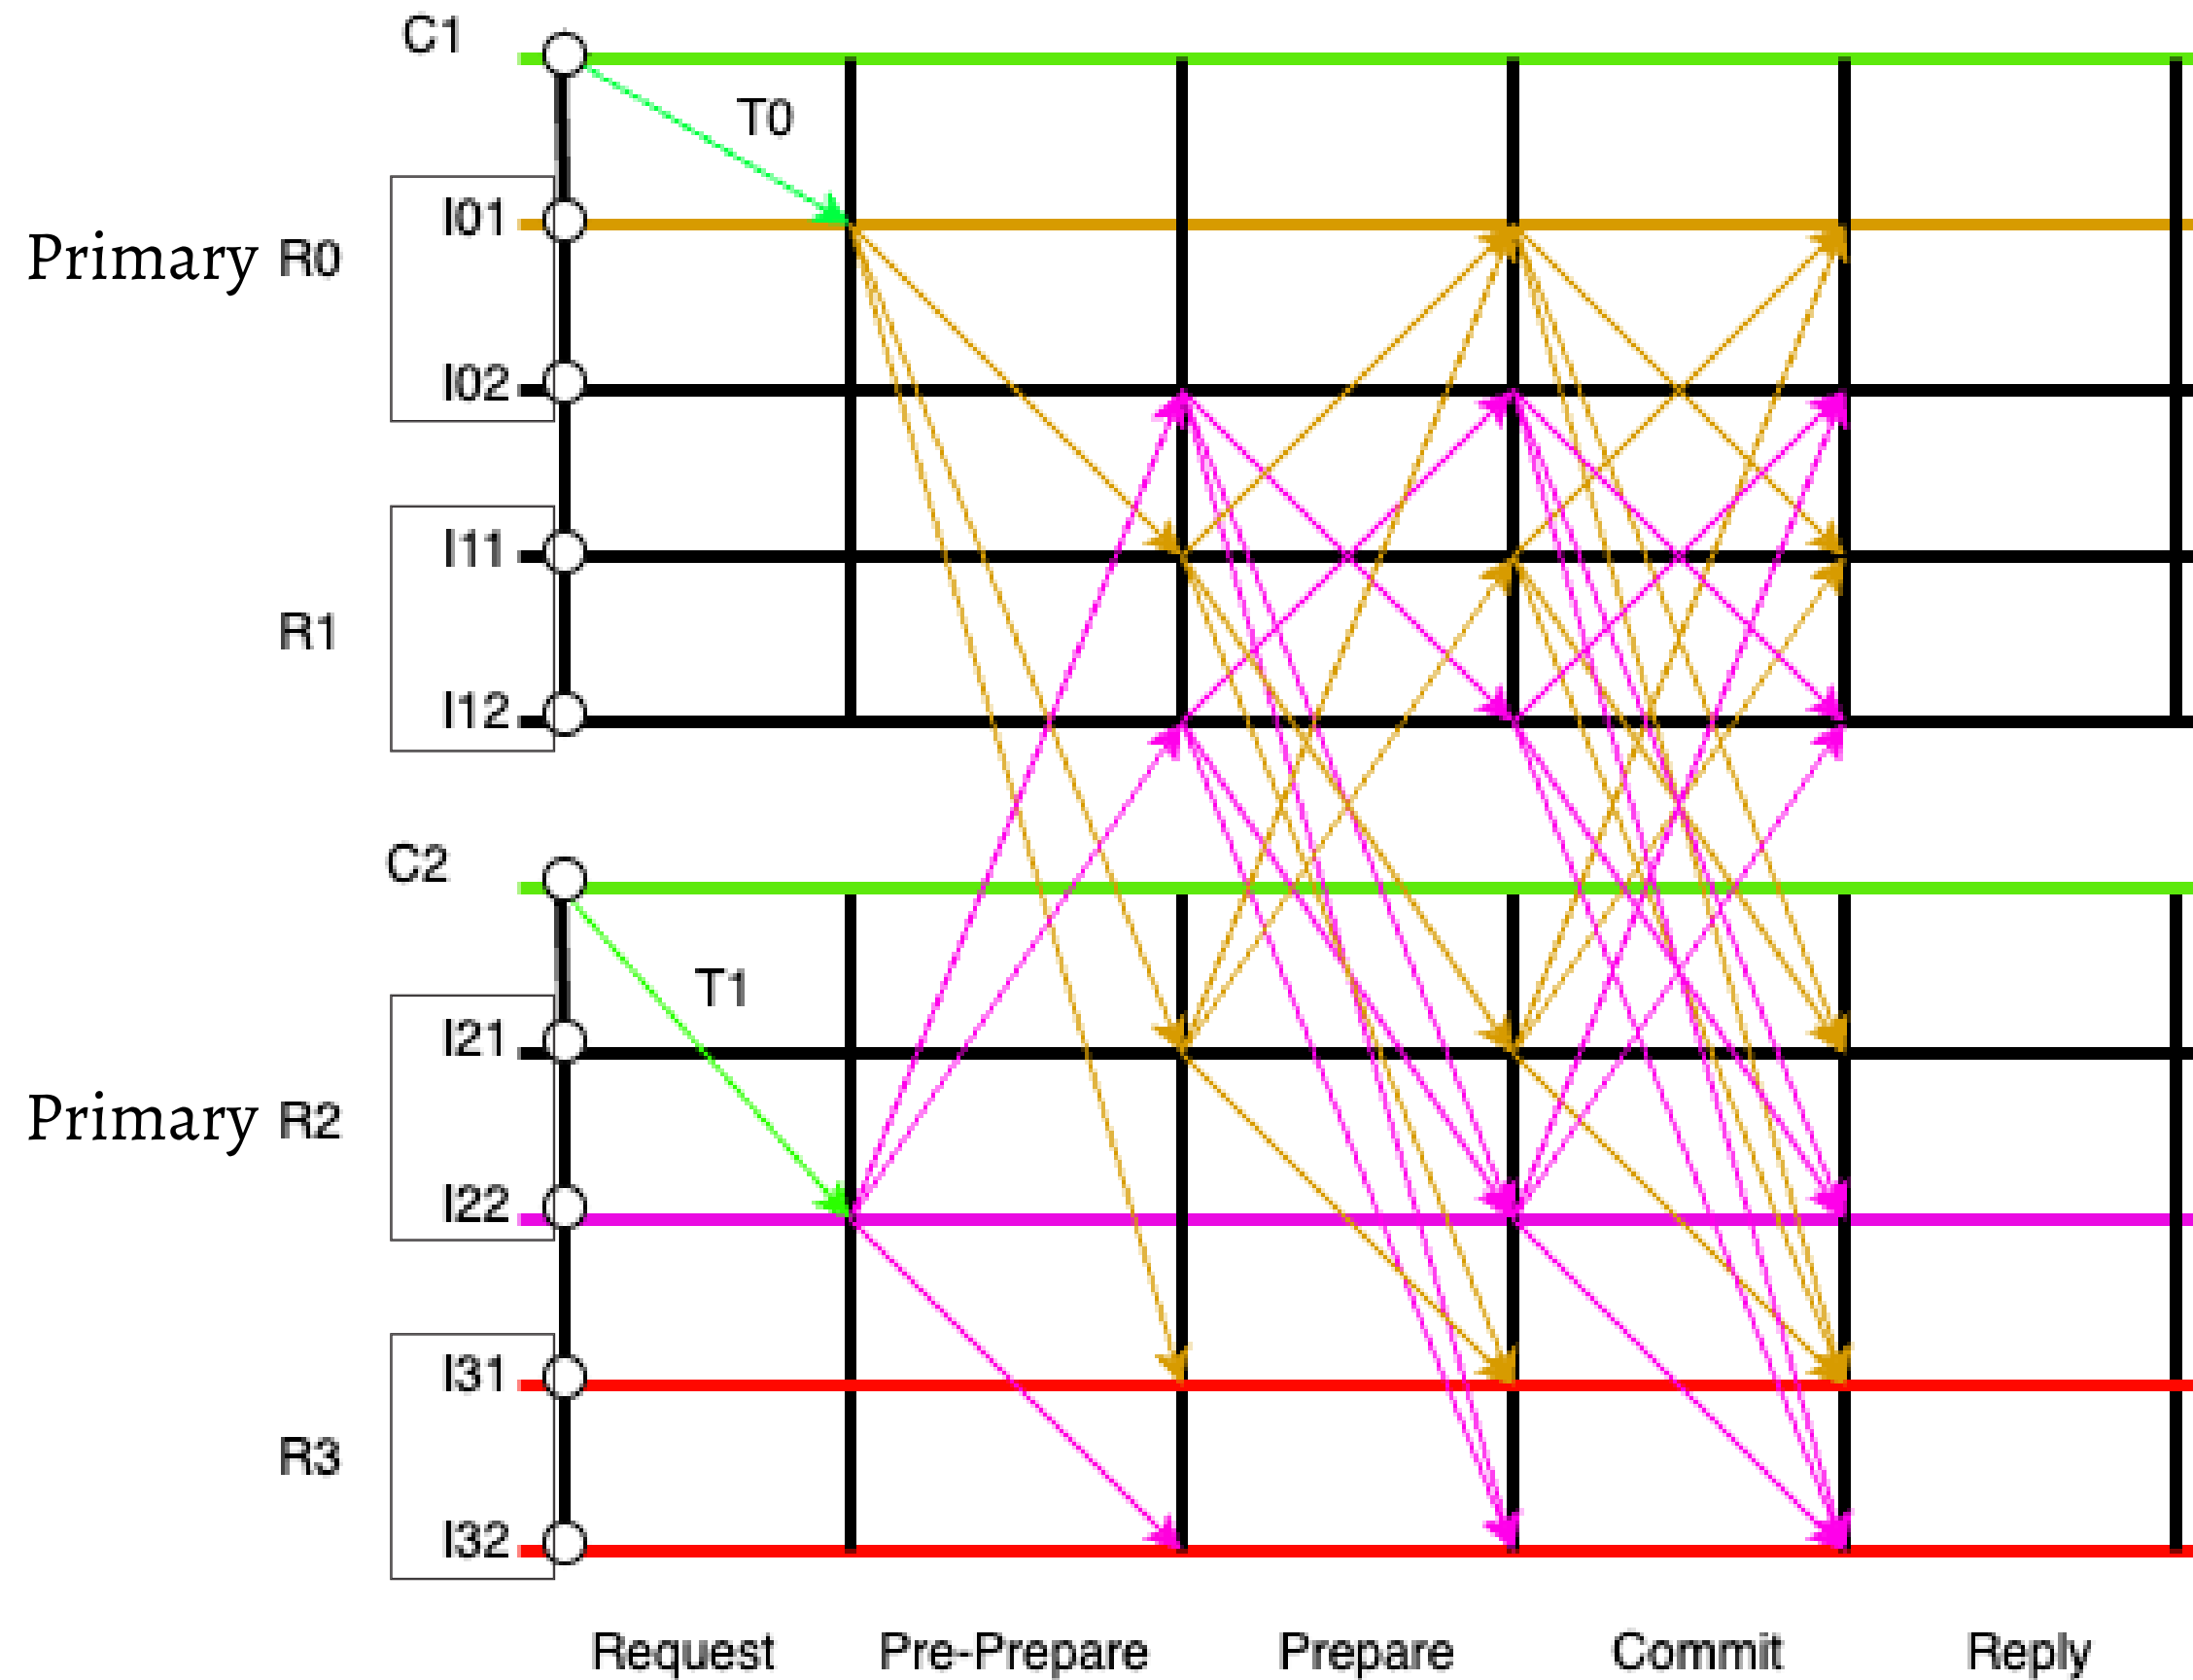
# Example




# Example



# Example



# Example




# Dealing with failures

- RCC follows asynchronous communication - Impossible to distinguish between malicious primary and primary whose messages get lost in the network.
- Reliable bounded delay communication - All messages sent by  $n$  replicas will arrive within some maximum delay.



# Dealing with Detectable Failures



$i$ : primary number  
 $\rho$ : round number  
 $P$ : state of  $R$ (that detects failure)  
 $E$ : set of  $nf$  failure messages

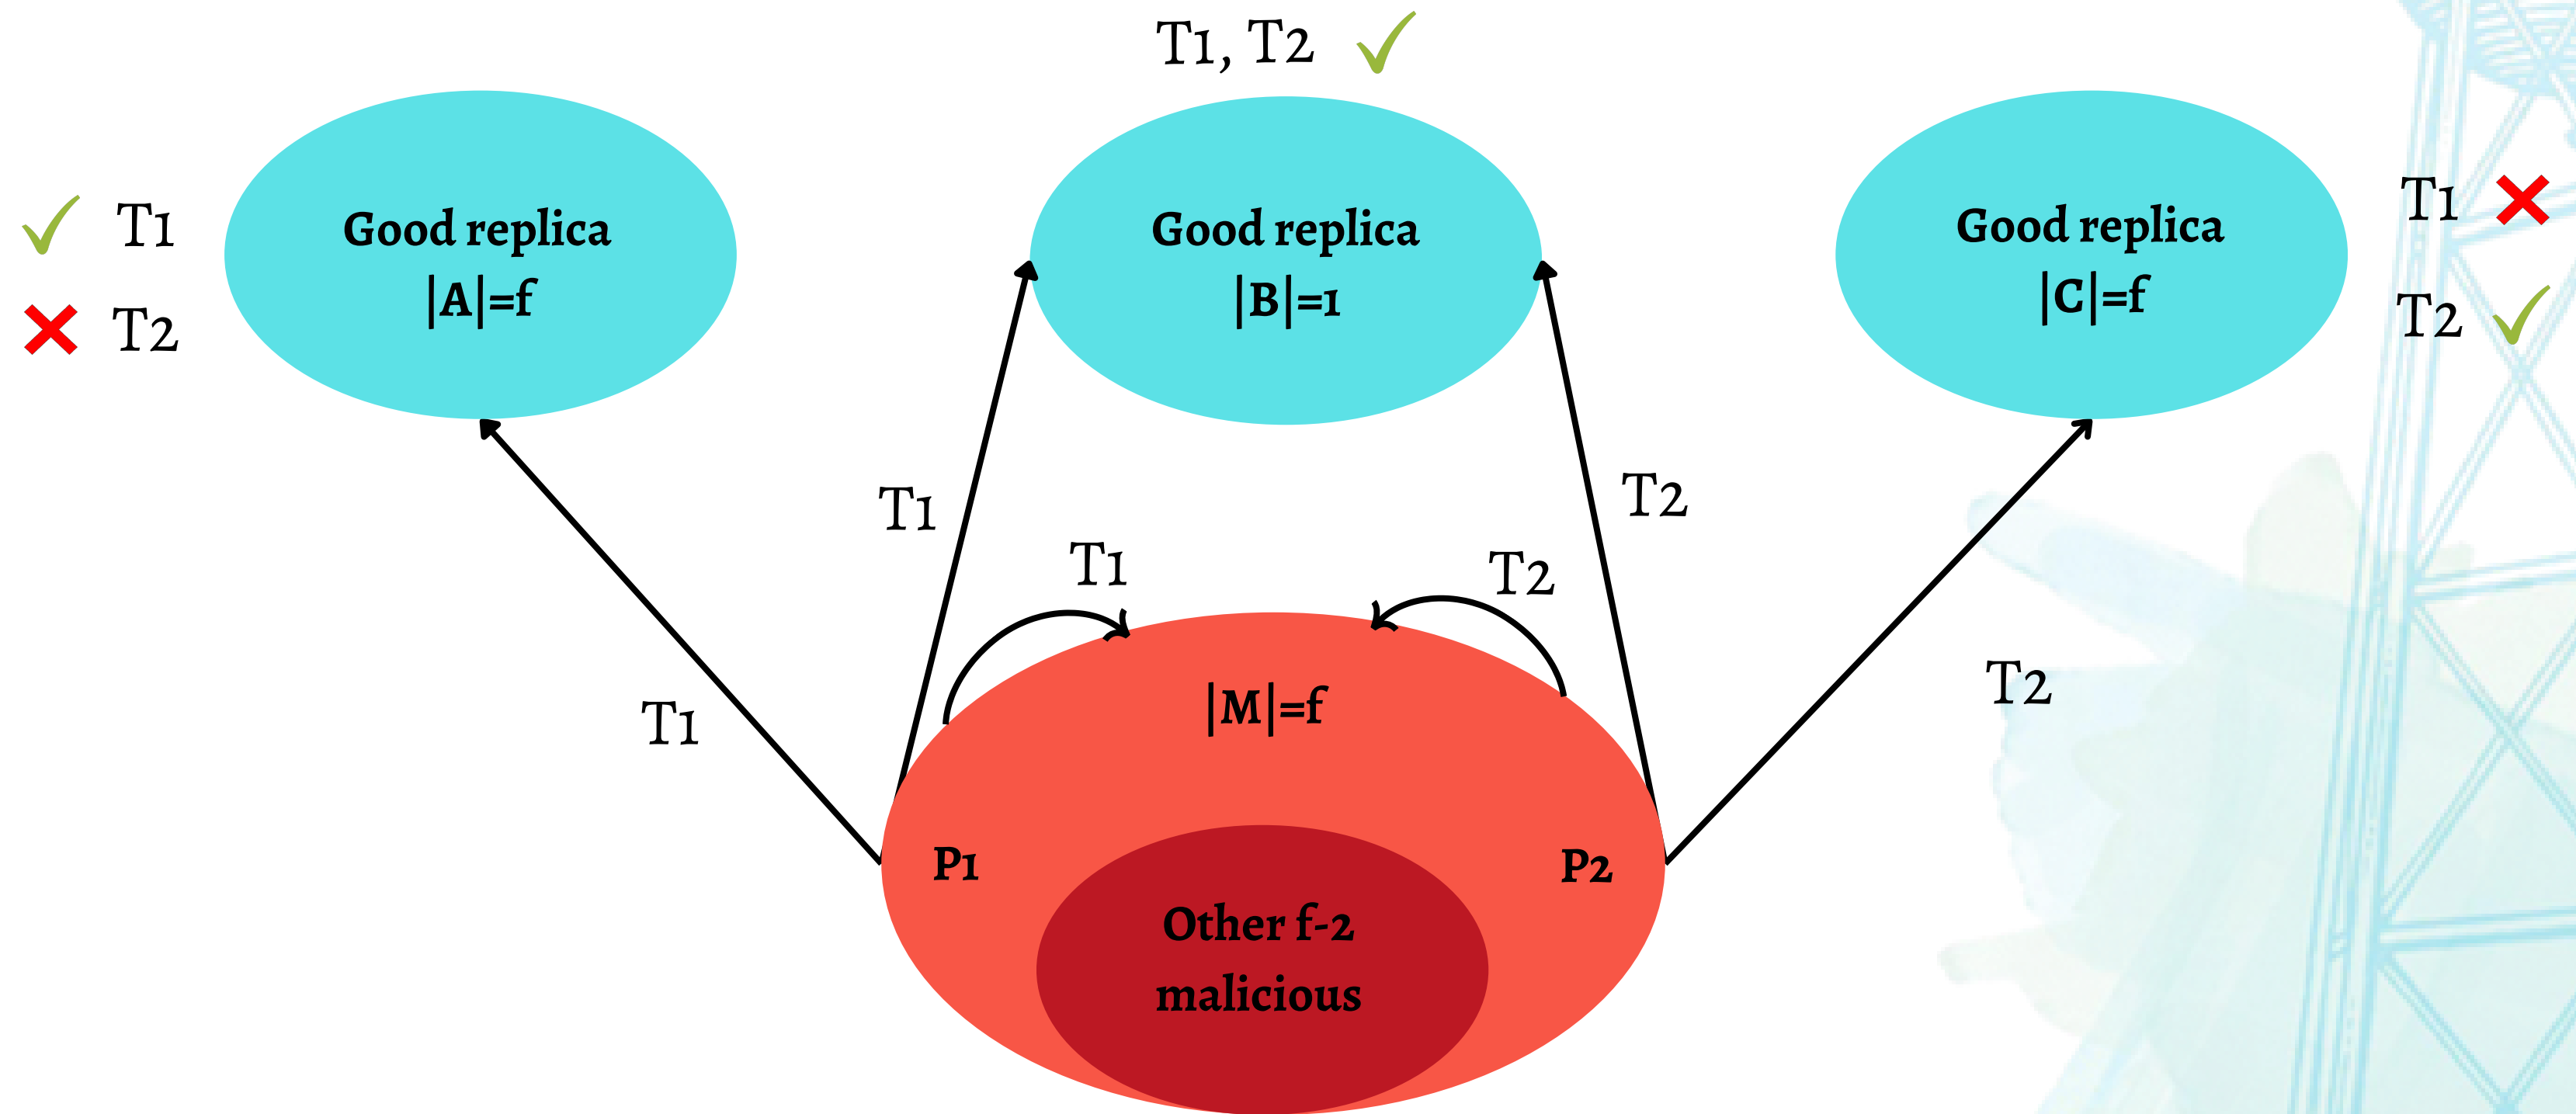
# Dealing with Detectable Failures

## Recovery process


All non faulty replicas need to:

- Detect failure of primary
- Reach agreement on state I (faulty instance)
- Determine the round in which primary can resume operation

# Dealing with Undetectable Failures



# Dealing with Undetectable Failures



# Client Interactions with RCC

- Importance of **distinct client transactions** for optimal throughput.
- It is designed to handle faulty clients without needing their cooperation.

## Design Optimization

- RCC optimized for scenarios with many more clients than replicas.
- Each client is assigned to a single primary ( $P_i$ ).
- Only ( $P_i$ ) can propose transactions for that client ( $c$ )

# Client Interactions with RCC

## Two Key Challenges

Primaries not  
receiving client  
requests



- **Problem:** Less concurrent clients than replicas.
- **Solution:** Propose small "no-op" requests in such situations.

Faulty Primaries



- **Problem:** Faulty primaries refusing to propose some clients' requests.
- **Solution:** Incentivize Malicious primary to not refuse service and detect primary failure.

# Client Interactions with RCC

- **Handling Unwilling or Incapable Primaries**

**Situation:** Primary ( $P_i$ ) is unwilling or crashes.

**Solution:** Client ( $c$ ) requests reassignment to another instance ( $I_j$ ) by broadcasting a request ( $m := \text{SWITCHINSTANCE}(c, j)$ ) to all replicas.



# Improving Resilience of Consensus

## What are the challenges in Traditional Consensus Protocols?

- Heavy Reliance on the primary
- Designed to handle primaries that completely fail to propose client transactions.
- Not effectively designed to handle various forms of malicious behavior..



# Improving Resilience of Consensus

## Ordering Attack

**T<sub>1</sub>:** Transfer (Alice, Bob, 3)

**T<sub>2</sub>:** Transfer (Bob, Eve, 2)

**Alice**



8

**Bob**



0

**Eve**



0

```
Transfer (A,B,m) -> If    bal(A) > m  
                      then  
                      bal(A) = bal(A) - m  
                      bal(B) = bal(B) + m
```


**Case 1: T<sub>1</sub> then T<sub>2</sub>**

# Improving Resilience of Consensus

## Ordering Attack

**T<sub>1</sub>:** Transfer (Alice, Bob, 3)

**T<sub>2</sub>:** Transfer (Bob, Eve, 2)



Transfer (A,B,m) -> **If**  $\text{bal}(A) > m$

**then**

$\text{bal}(A) = \text{bal}(A) - m$

$\text{bal}(B) = \text{bal}(B) + m$

**Case 1: T<sub>1</sub> then T<sub>2</sub>**

# Improving Resilience of Consensus

## Ordering Attack

**T<sub>1</sub>:** Transfer (Alice, Bob, 3)

**T<sub>2</sub>:** Transfer (Bob, Eve, 2)

**Alice**



8

**Bob**



0

**Eve**



0

Transfer (A,B,m) -> **If**  $\text{bal}(A) > m$

**then**

$\text{bal}(A) = \text{bal}(A) - m$

$\text{bal}(B) = \text{bal}(B) + m$


**Case 1: T<sub>2</sub> then T<sub>1</sub>**

# Improving Resilience of Consensus

## Ordering Attack

**T<sub>1</sub>:** Transfer (Alice, Bob, 3)

**T<sub>2</sub>:** Transfer (Bob, Eve, 2)



Transfer (A,B,m) -> **If**  $\text{bal}(A) > m$   
**then**  
 $\text{bal}(A) = \text{bal}(A) - m$   
 $\text{bal}(B) = \text{bal}(B) + m$

**Case 1: T<sub>2</sub> then T<sub>1</sub>**

# Improving Resilience of Consensus

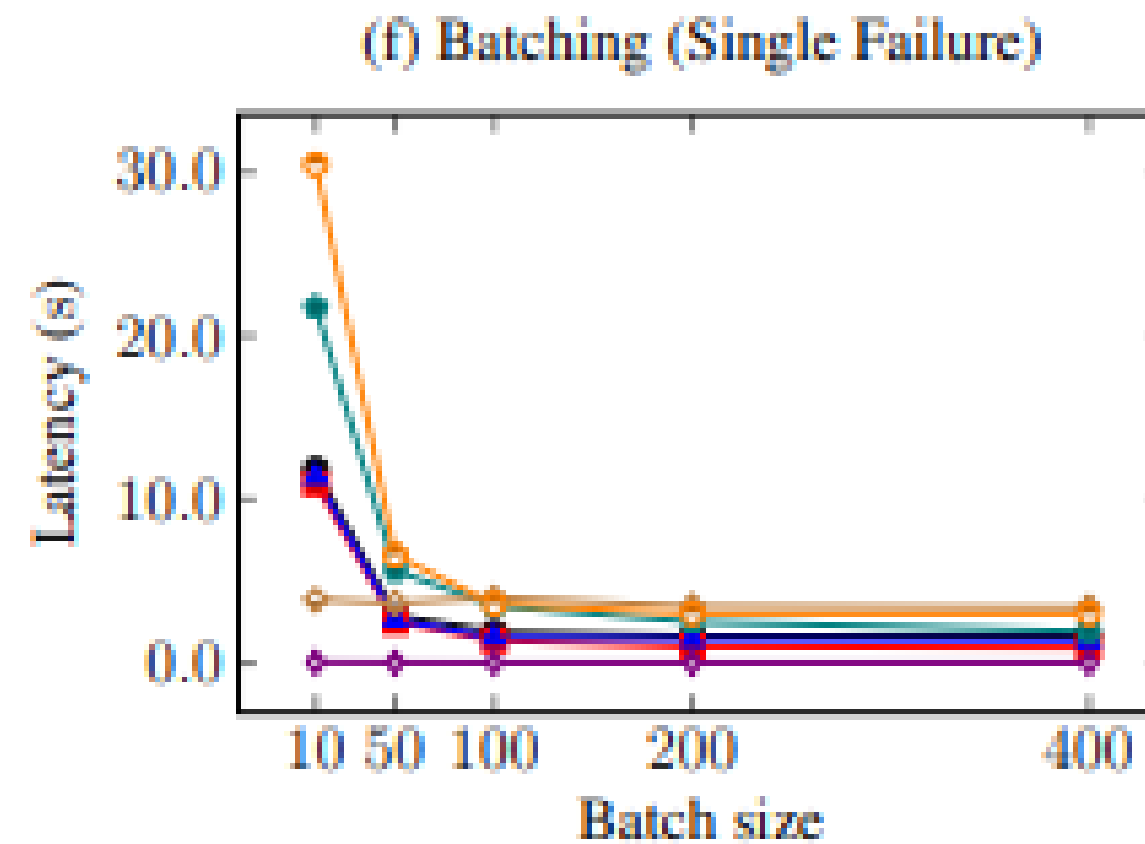
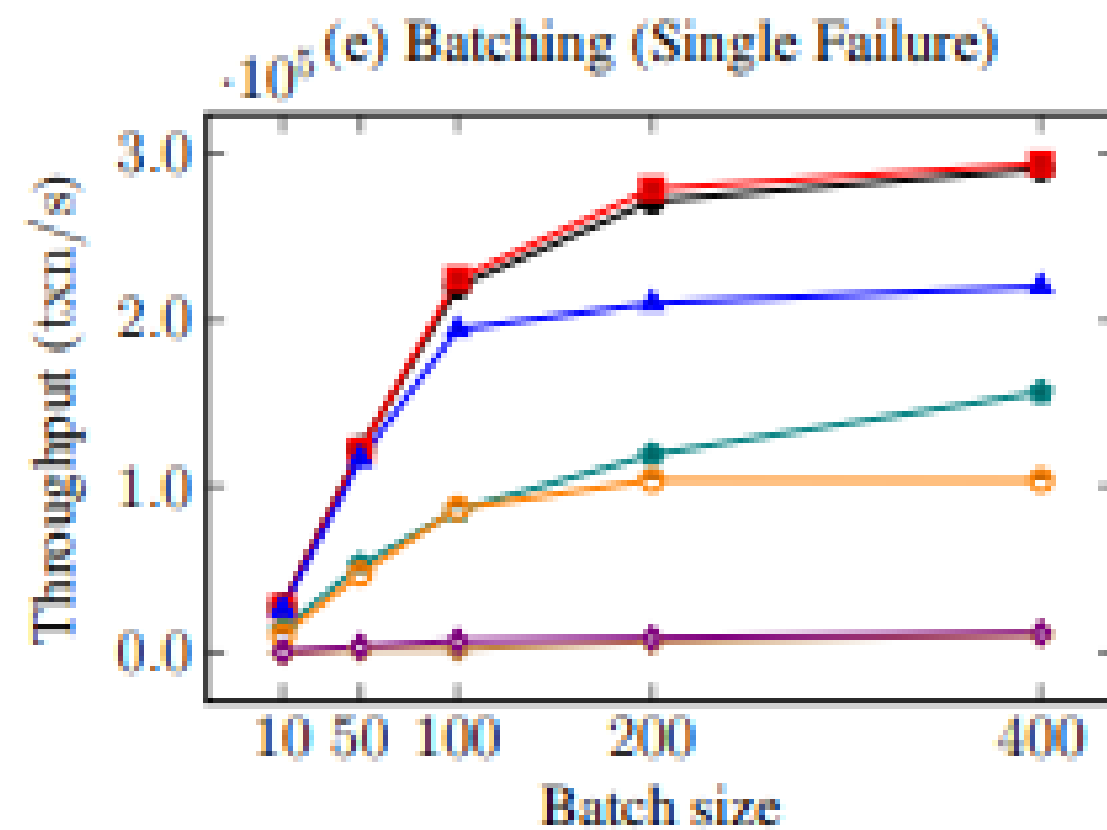
## How does RCC mitigate the ordering attack

- A method to deterministically select different permutation of the order of execution in every round.
- The order is practically impossible to predict or influence by faulty replicas.

# Evaluation

## Varying batch size

- Performance increases when batch size increases.

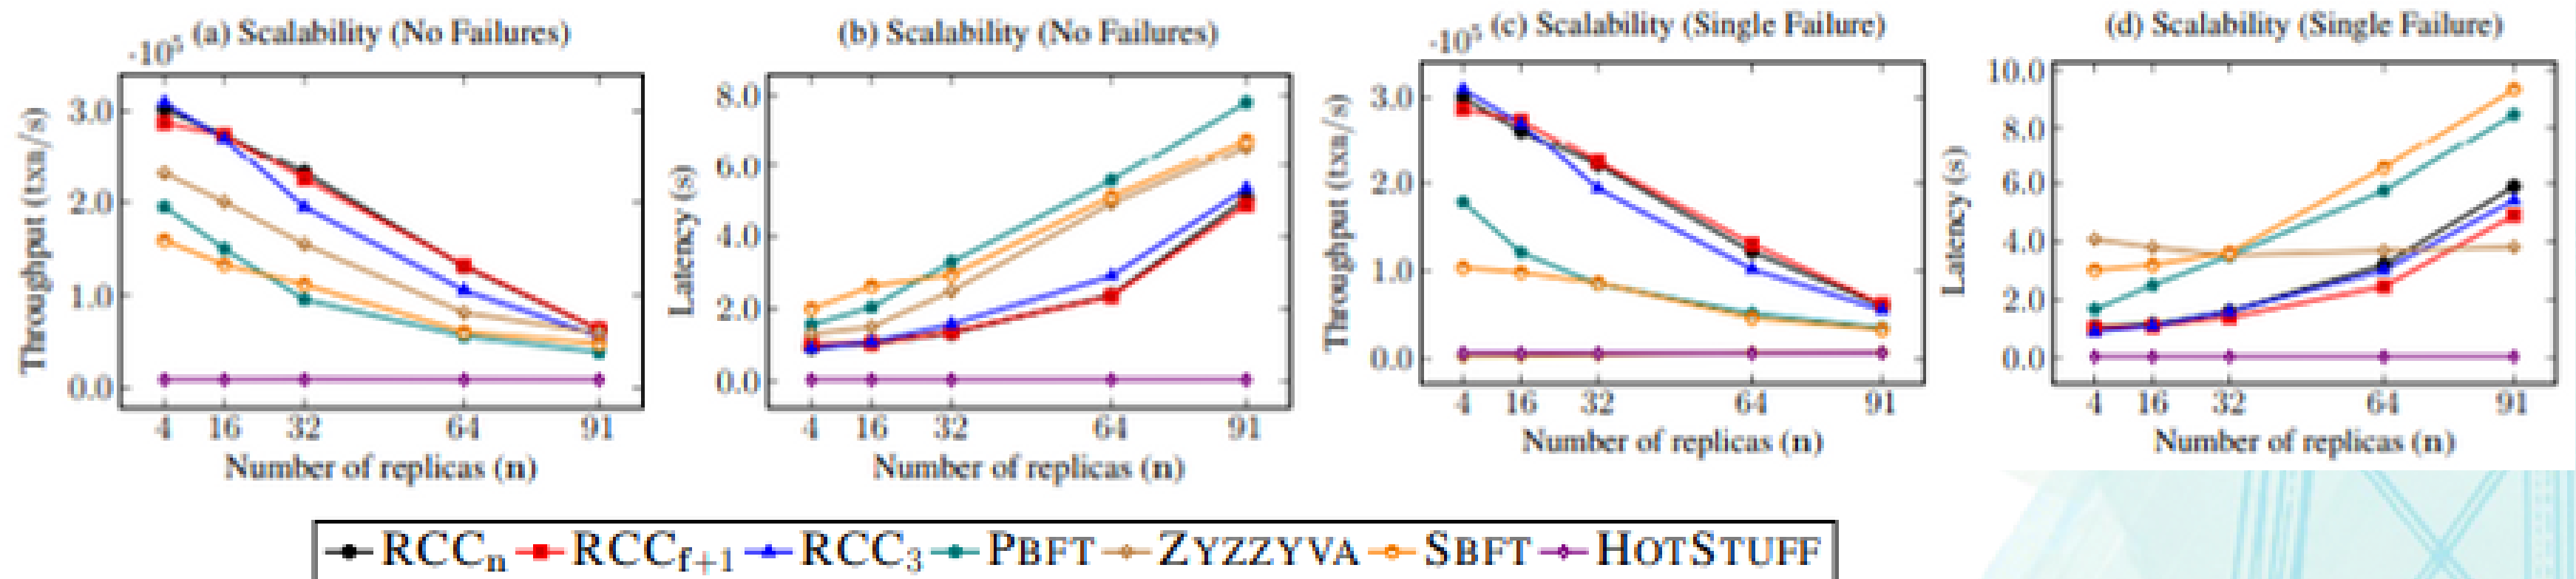


—●— RCC<sub>n</sub> —■— RCC<sub>f+1</sub> —▲— RCC<sub>3</sub> —◆— PBFT —◇— ZYZZYVA —○— SBFT —◇— HOTSTUFF

# Evaluation

## Varying Replicas

- Even in the best-case situation of no failures, RCC easily outperforms ZYZZYVA.
- RCC performs better than all other protocols in three versions.
- Performance is enhanced through improving concurrency by adding more instances.





# Conclusion

In RCC,

- Every replica is allowed to be a primary (parallel consensus).
- Enhanced security features are provided.
- A set of ordered client requests are guaranteed.
- More throughput than primary-backup consensus is provided.



# References

- S. Gupta, J. Hellings and M. Sadoghi; "RCC: Resilient Concurrent Consensus for High-Throughput Secure Transaction Processing"; IEEE 37th International Conference on Data Engineering (ICDE); Chania; Greece; 2021

**Thank You !**

