

All about Eve: Execute-Verify Replication for Multi-Core Servers

Authors: Manos Kapritsos, Yang Wang, Vivien Quema, Allen
Clement, Lorenzo Alvisi, Mike Dahlin

Presenter: Xuerui Li

Table of Contents

- 1 System Model
 - Background
 - Characteristics
 - Challenges
 - Improvement Based on Previous Techniques
- 2 Execution Stage
 - Basic Steps
 - Mixer Design
 - State Management
- 3 Verification Stage
 - Introduction
 - Basic Operations
 - Asynchronous Byzantine Fault Tolerant
 - Synchronous Primary-backup
 - Tolerating Concurrency Bugs
- 4 Results
- 5 Comparison with other Works
- 6 Conclusion

Comparison with Zyzzyva

- In Zyzzyva the agreement on inputs guarantees agreement on outputs: hence, a replica need only send (a hash of) the sequence of requests it has executed to convey its state to a client.
- In contrast, in Eve there is no guarantee that correct replicas will be in the same state, as the mixer may have incorrectly placed conflicting requests in the same parallelBatch.
- Verification stage is moved to the clients to reduce cost
- Shortcomings: It may introduce corner cases

Execution Stage

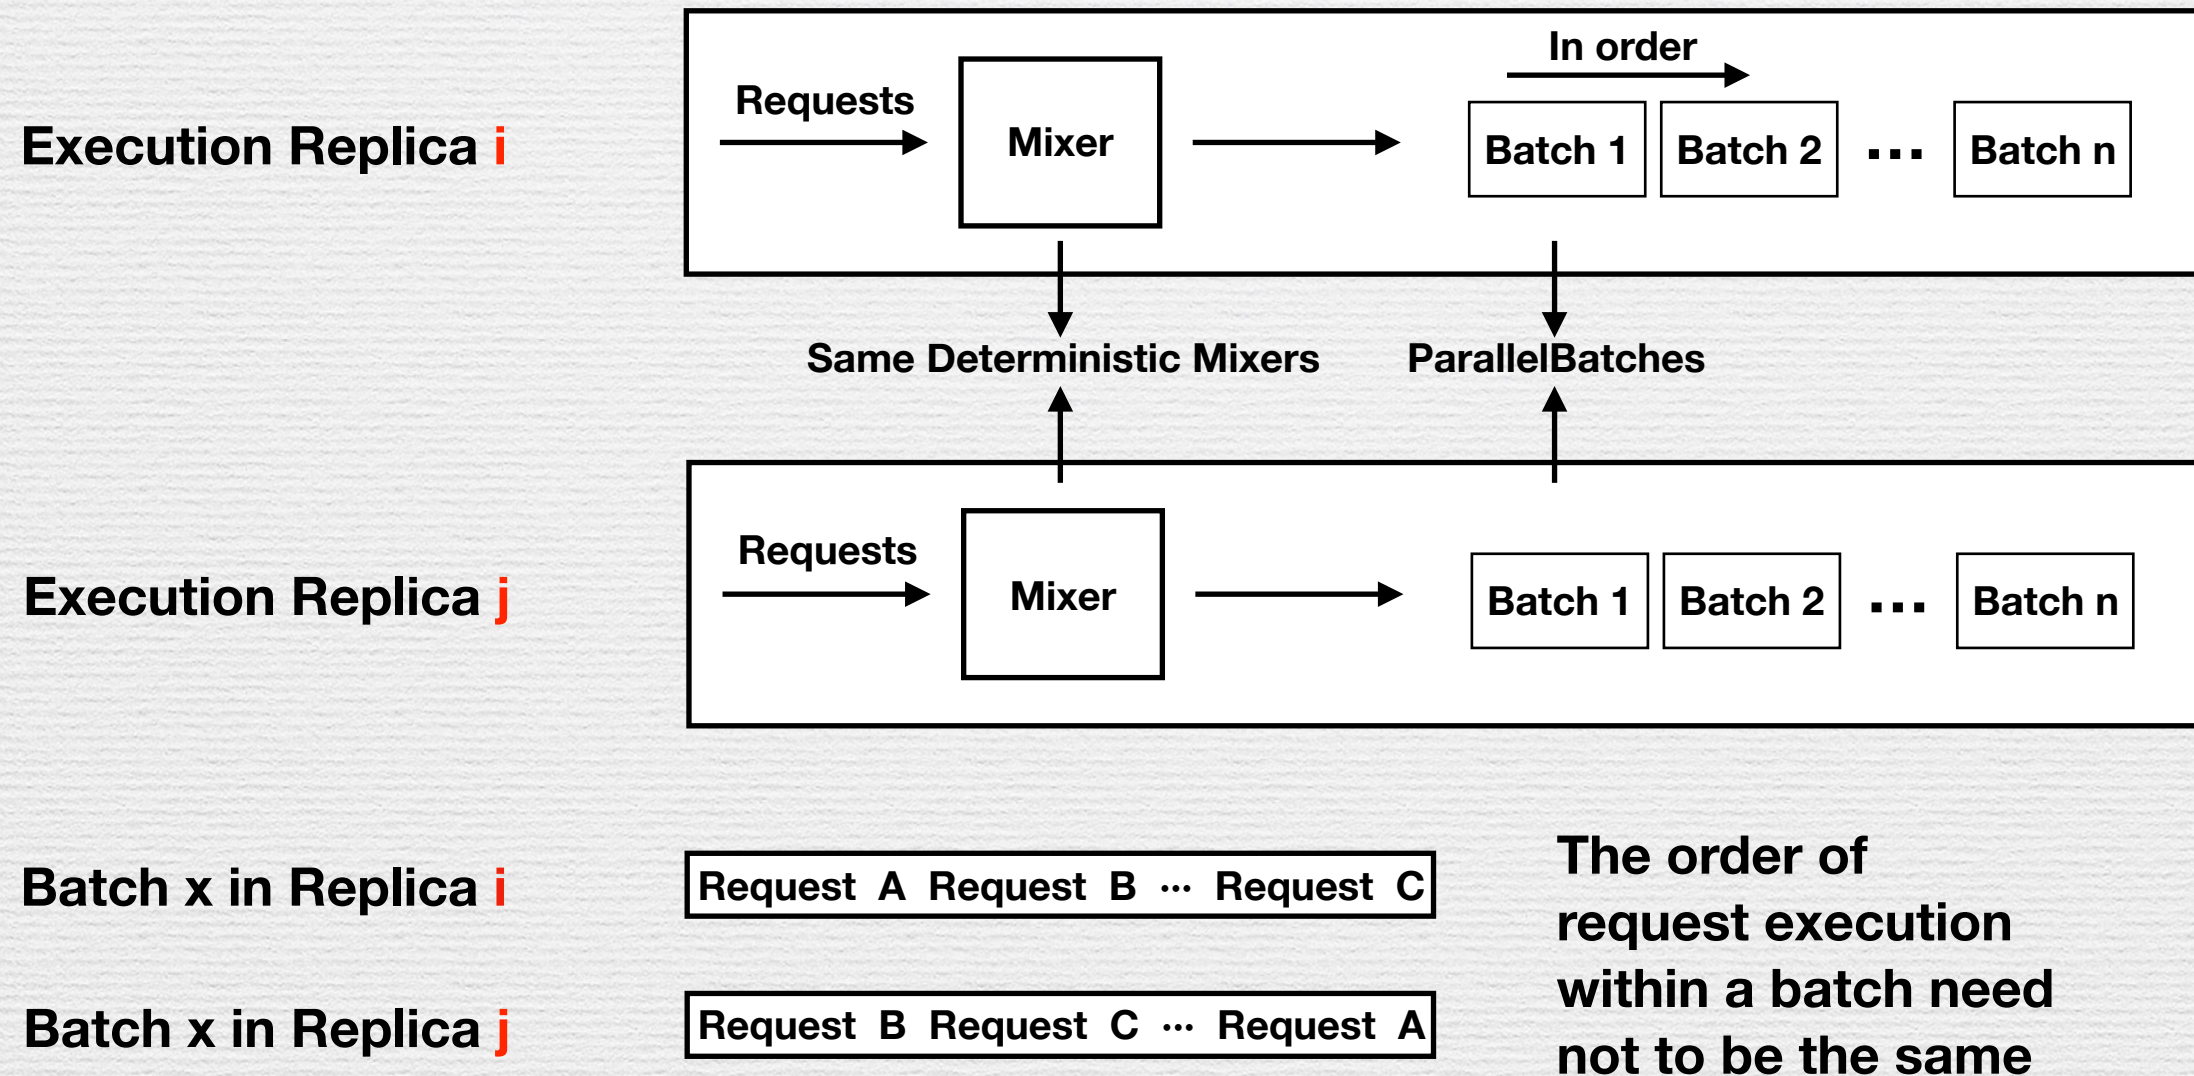


Fig. 2 ParallelBatches

Execution Stage

- Divergence: When one object is written by a request, at the exactly same time it is also used by another request.

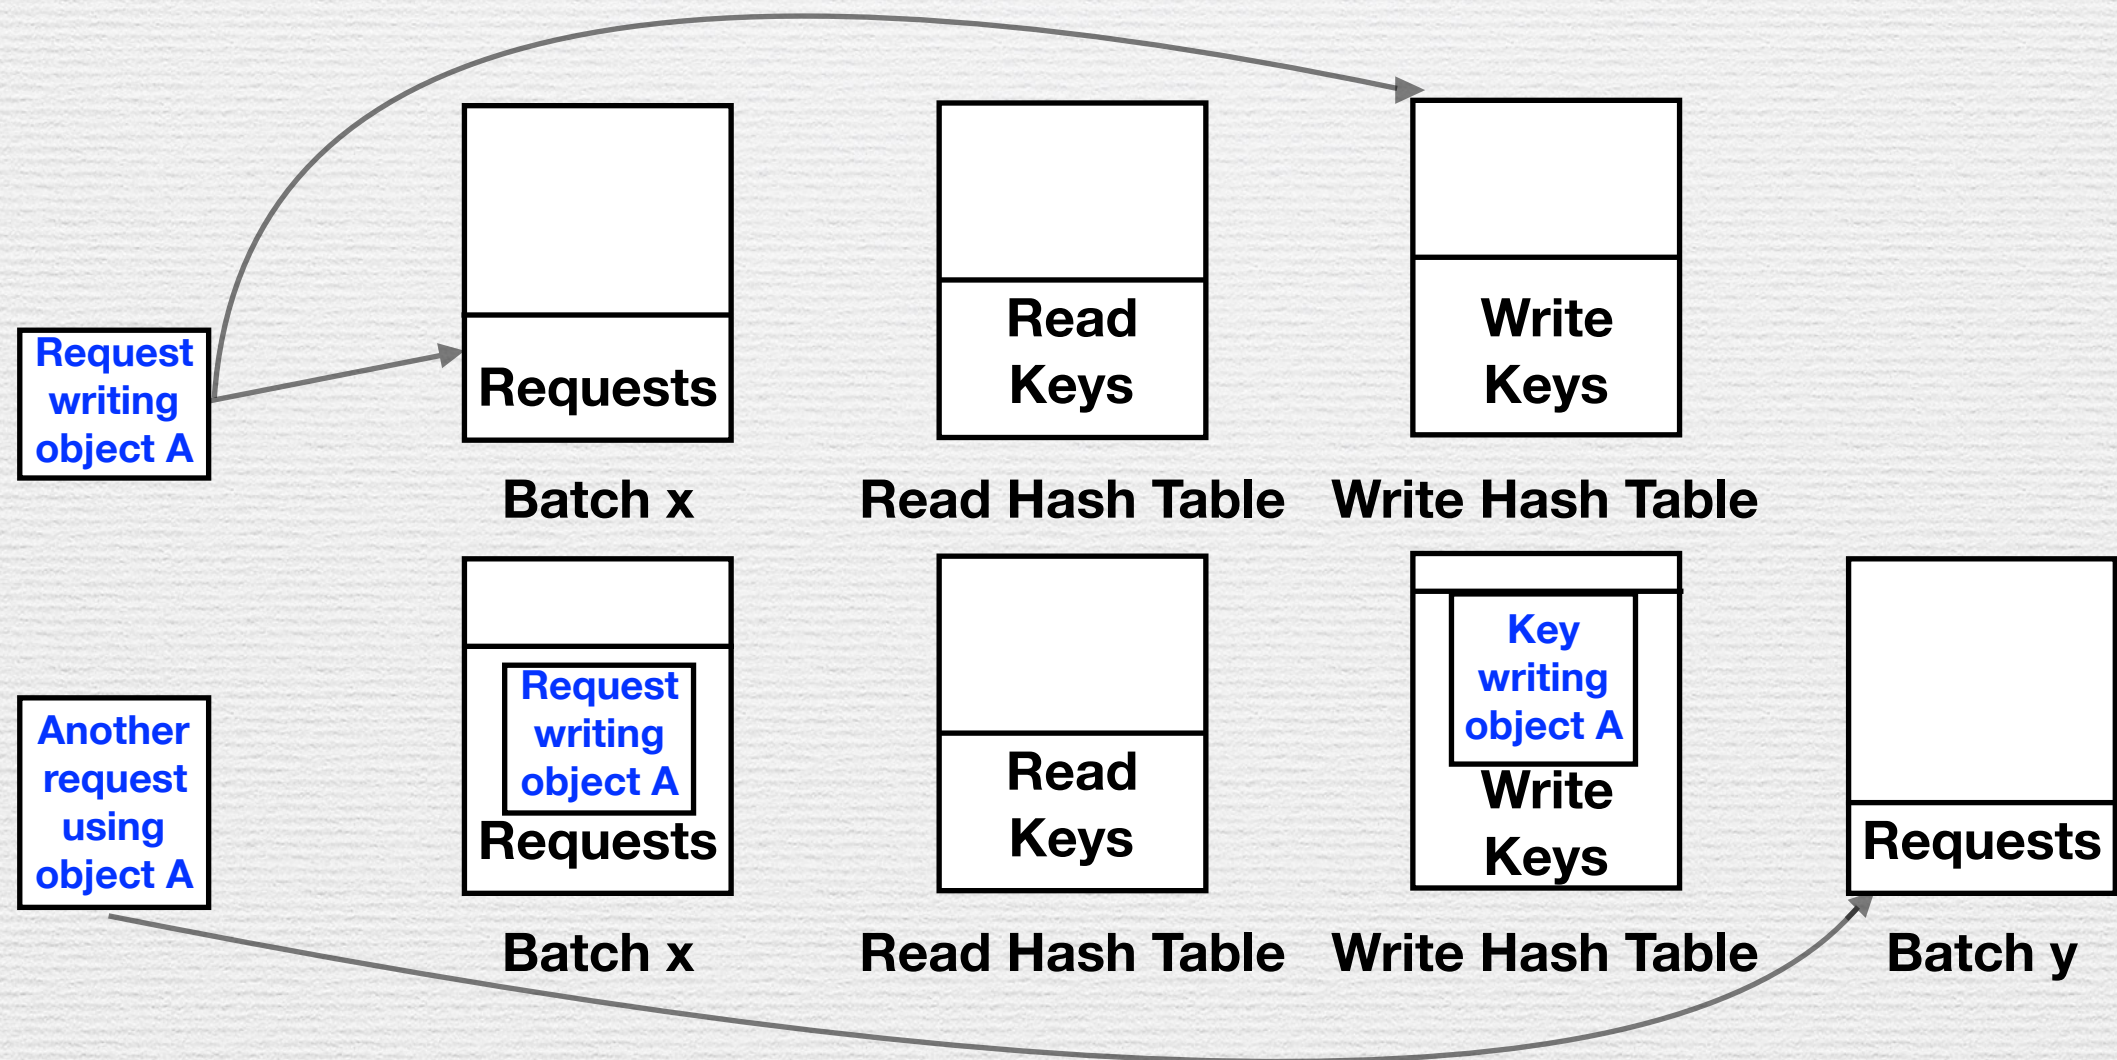


Fig. 3 Mixer

Execution Stage

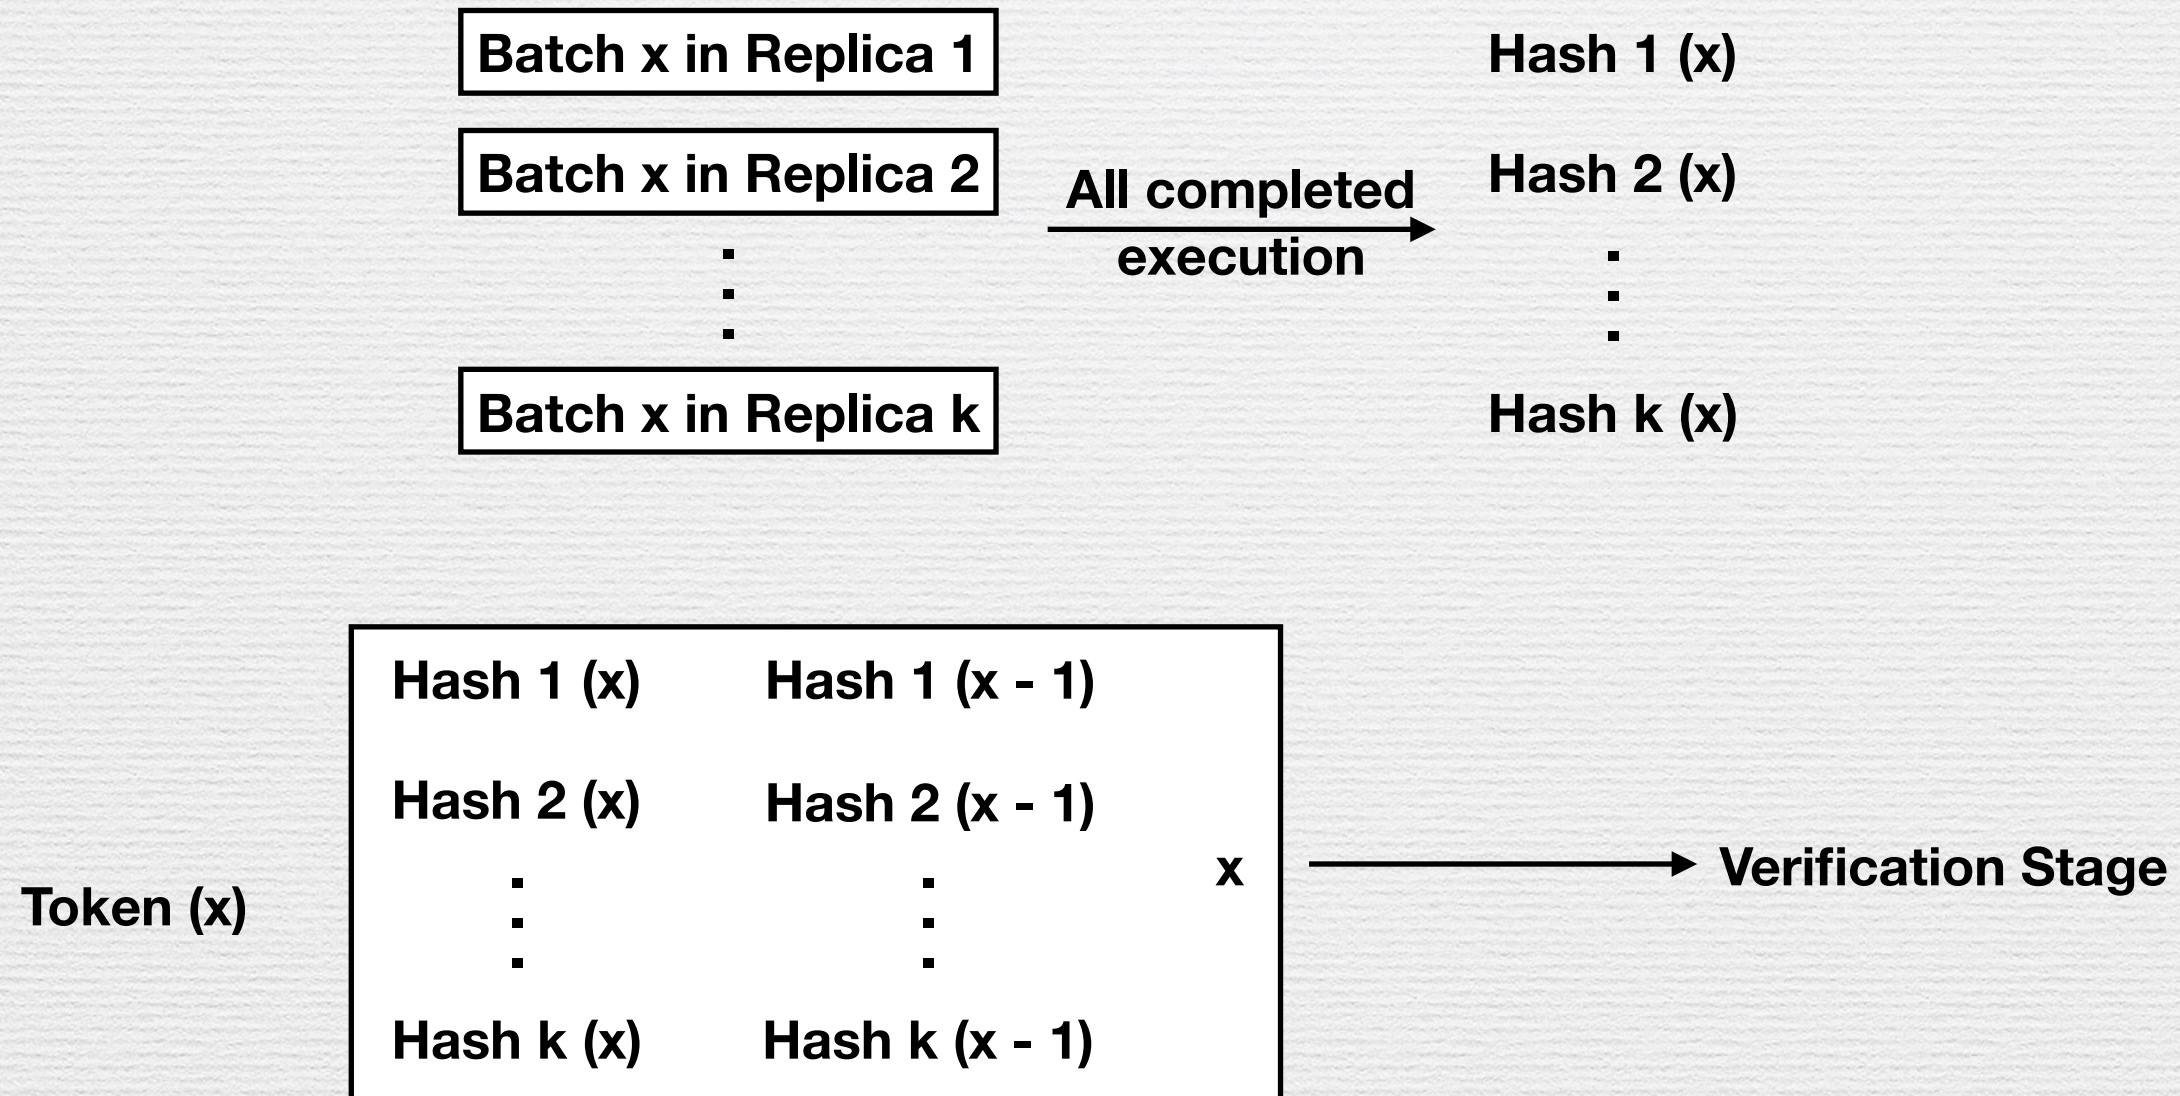


Fig. 4 Token for Verification

Verification Stage

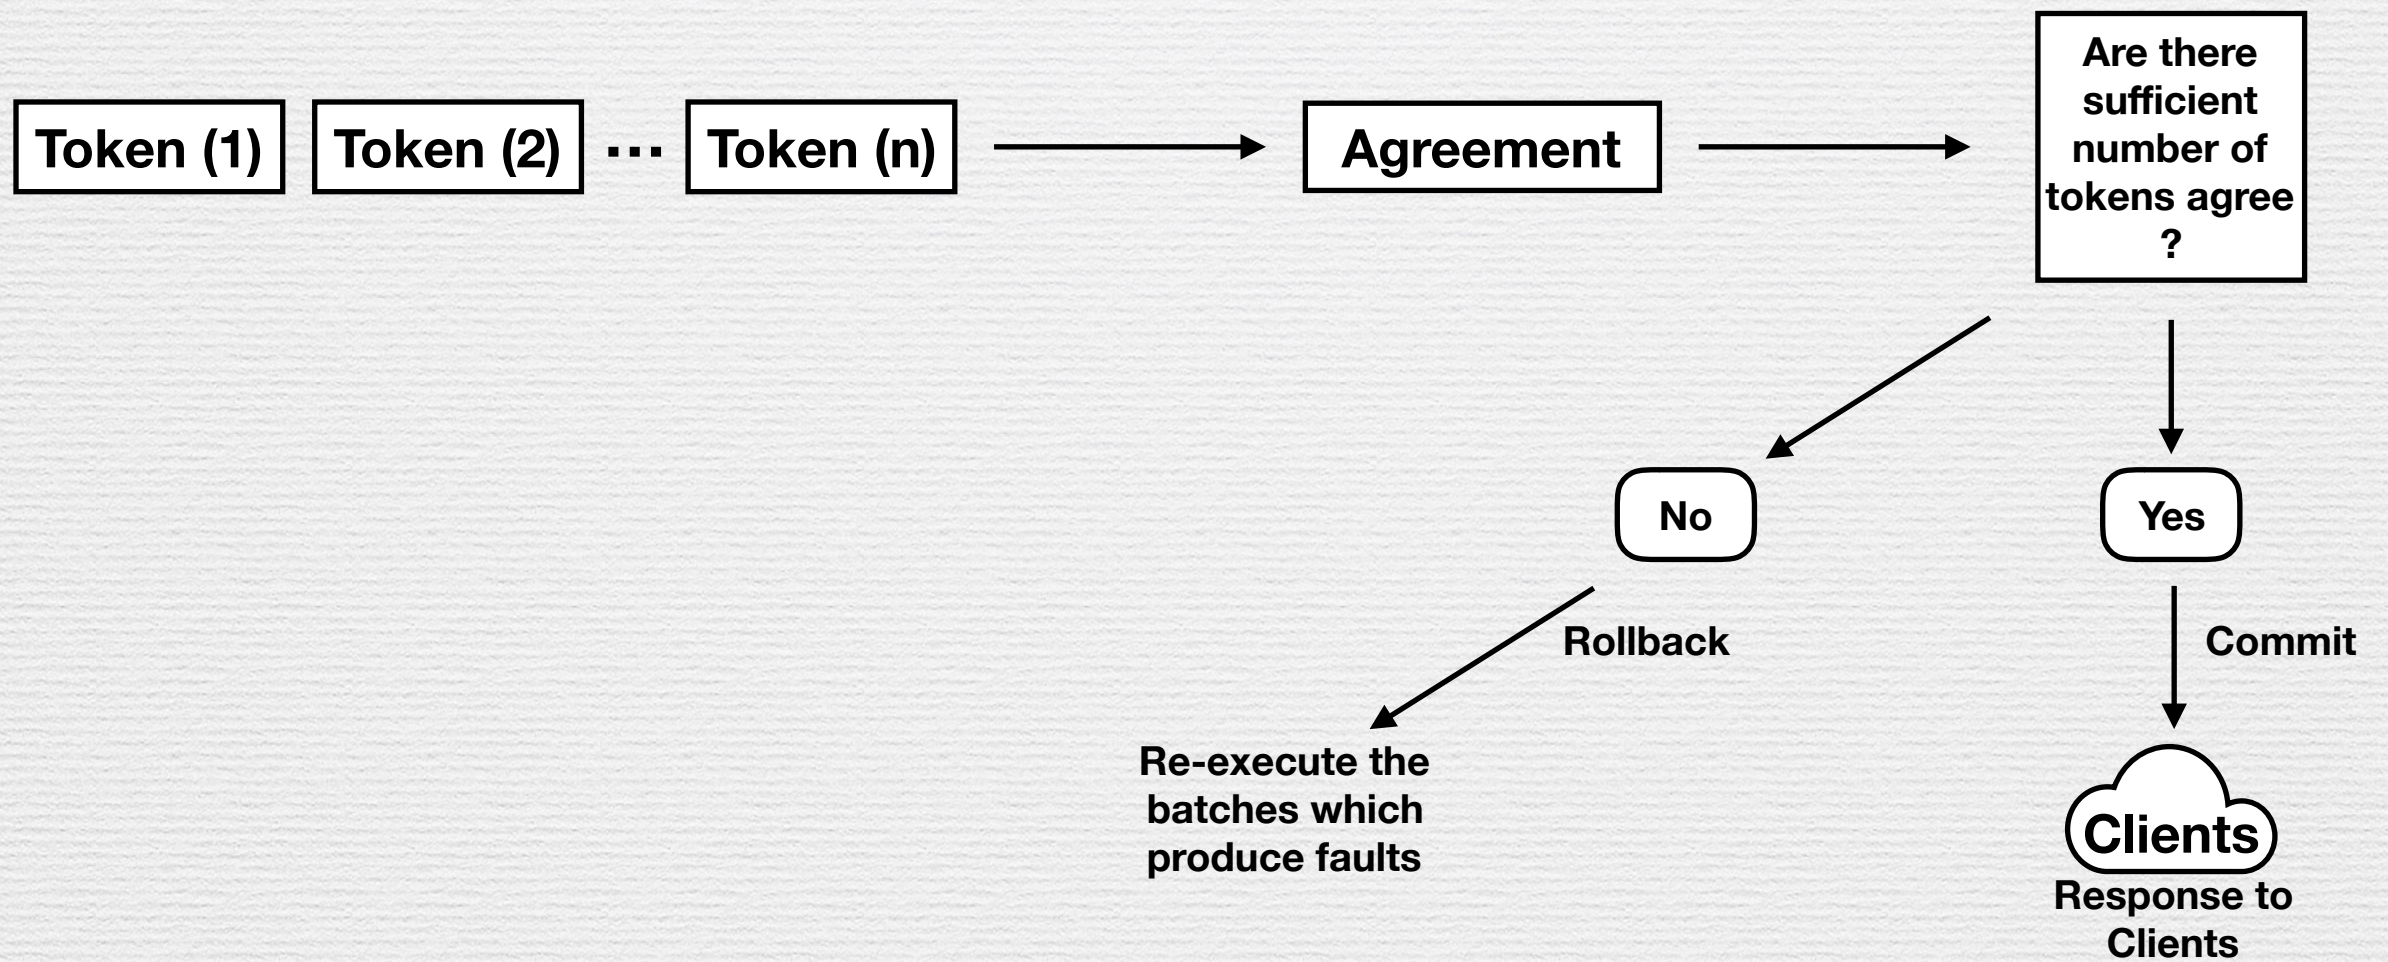
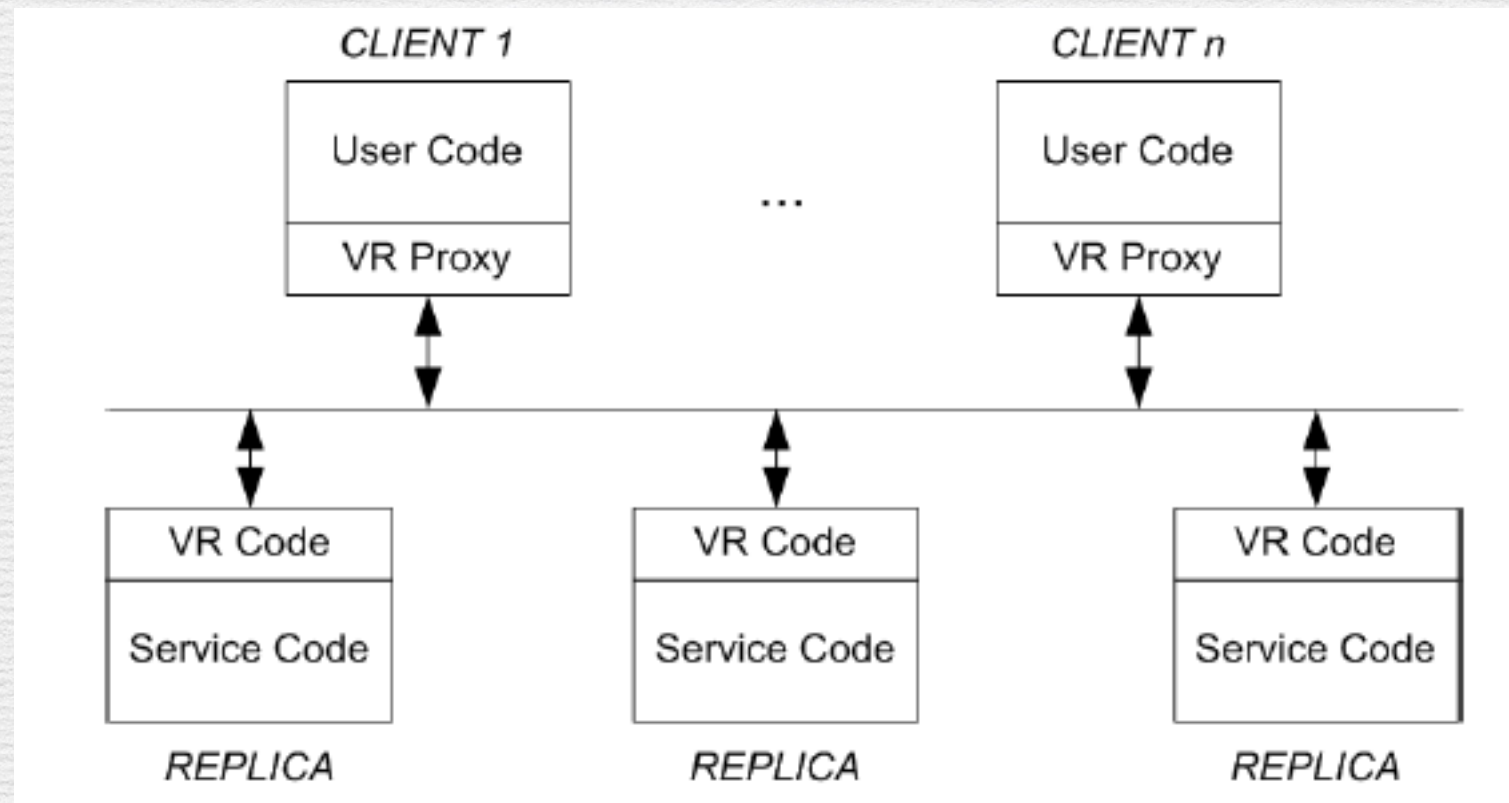


Fig. 5 Verification Stage

Background

- State machine
 - An abstract machine concerning states
- State machine replication
 - Replication: keeping the same state among distinct nodes through placing **copies** of the State Machine on **distinct servers**
 - Provides multiple equivalent execution for the same program
 - Enables implementation of analysis tools on some of the replica nodes while not disturbing the work of other replica nodes instead of slow them down.
 - Steps:
 - Receive client requests as Inputs
 - Choose an ordering for the Inputs and **execute them in this order on each server**
 - Respond to clients with the Output from the State Machine

Background



Overall Model

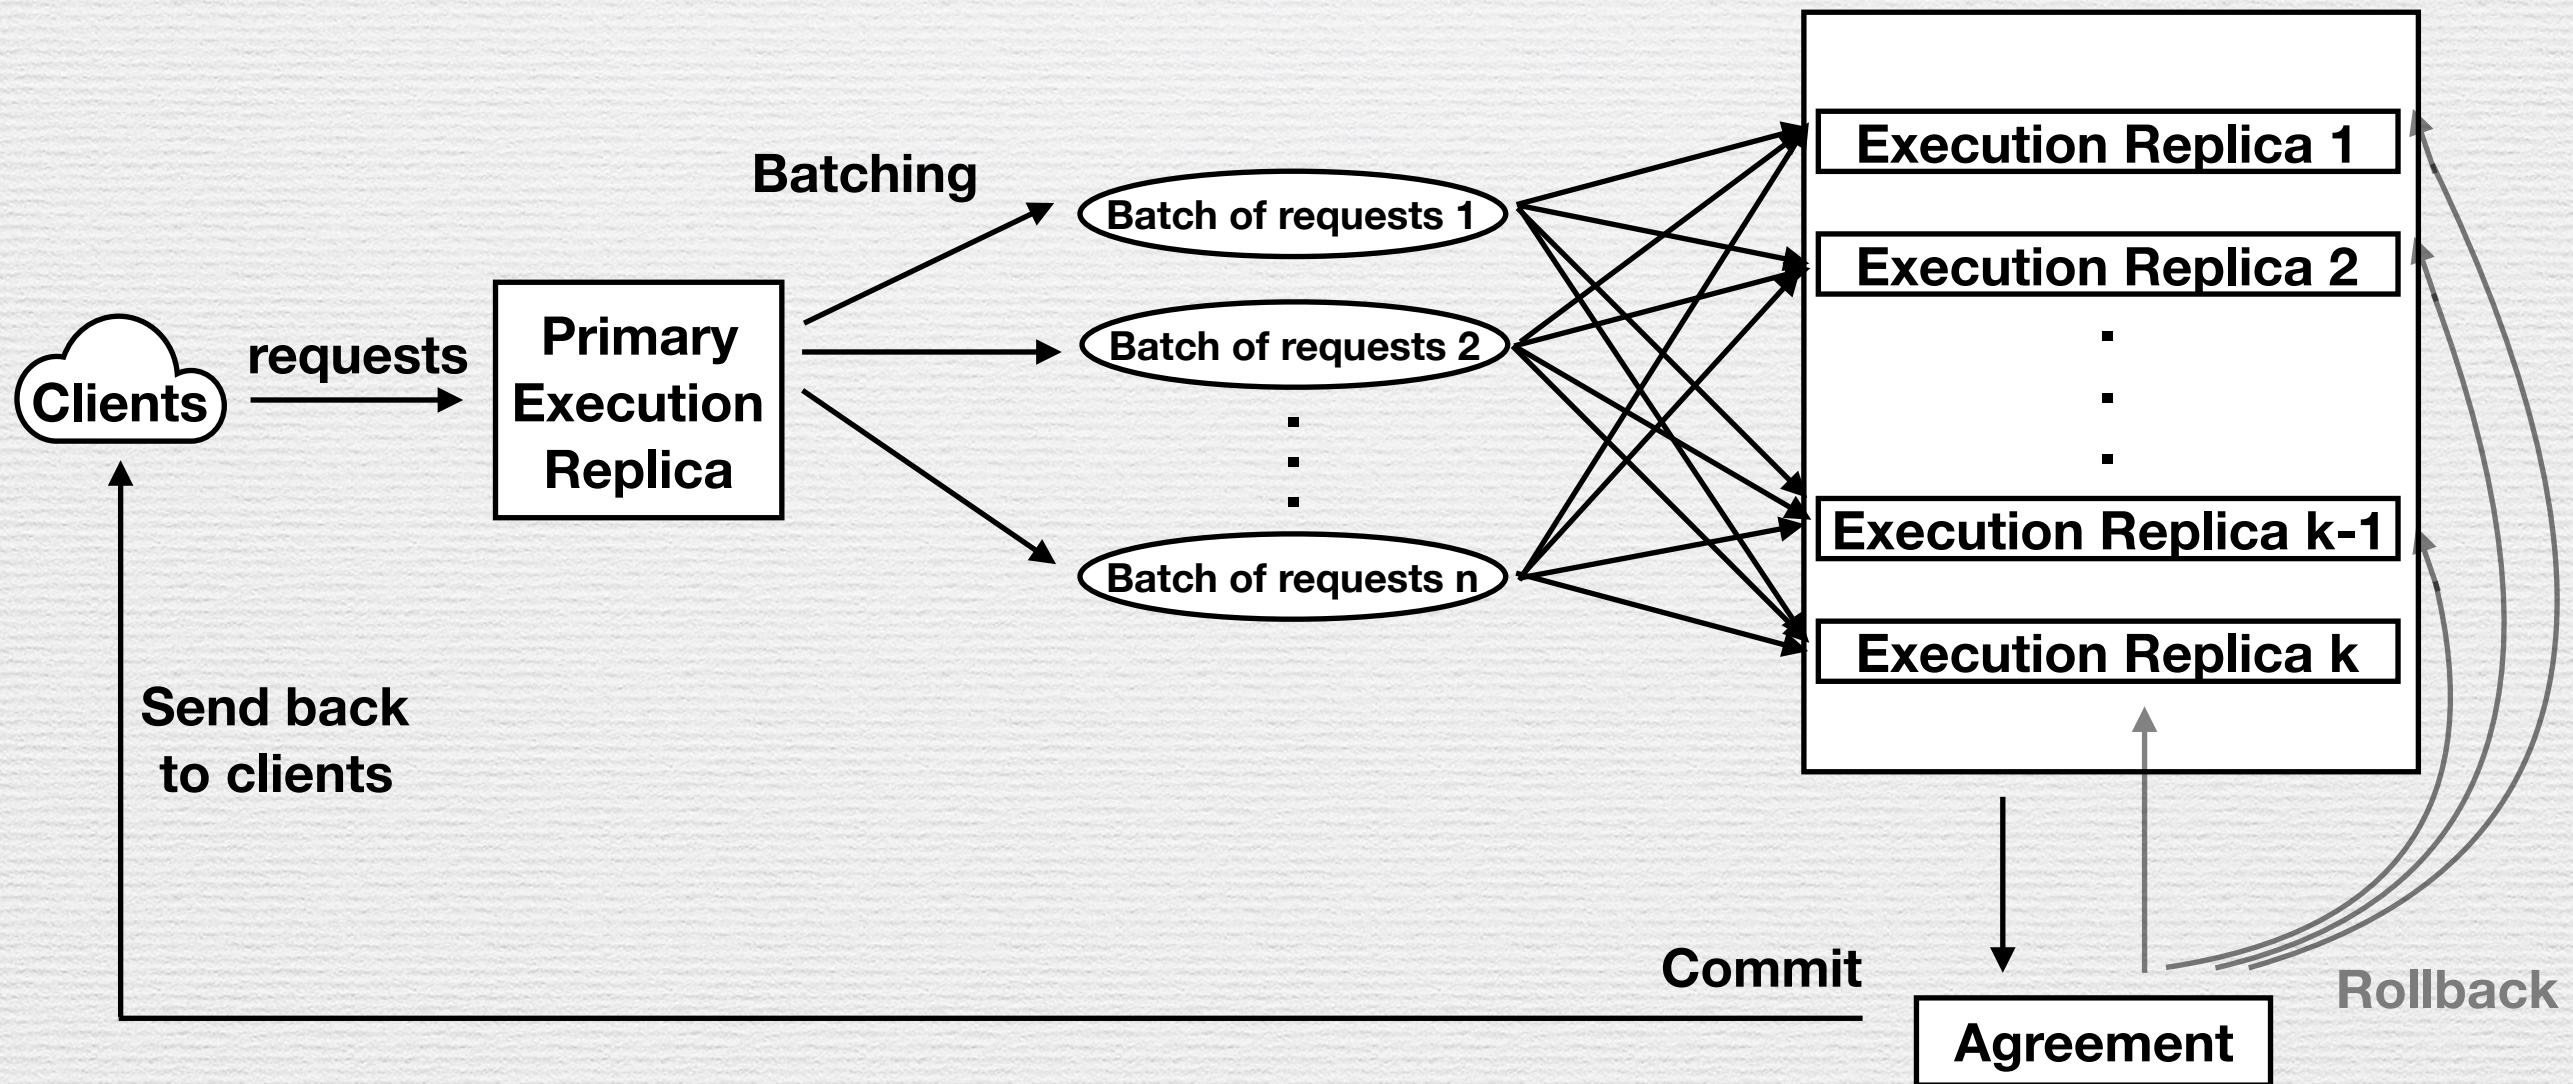


Fig. 1 Overall Model

Characteristics

- A new **Execute-Verify** architecture that allows **state machine replication** to scale to **multi-core servers**
- Can be applied to **both synchronous and asynchronous systems**
- Guarantees of robustness:
 - Attempting to **run only unlikely-to-interfere requests** in parallel
 - Ability to **detect and recover when** concurrency causes executions to **diverge**

Characteristics

- Nondeterministic interleaving of requests ensures high-performance replication for multi-core servers.
 - It avoids the overhead of enforcing determinism
- Independence enables **tolerance of a wide range of faults**
 - It helps tolerating crash, omission, or Byzantine faults

Characteristics

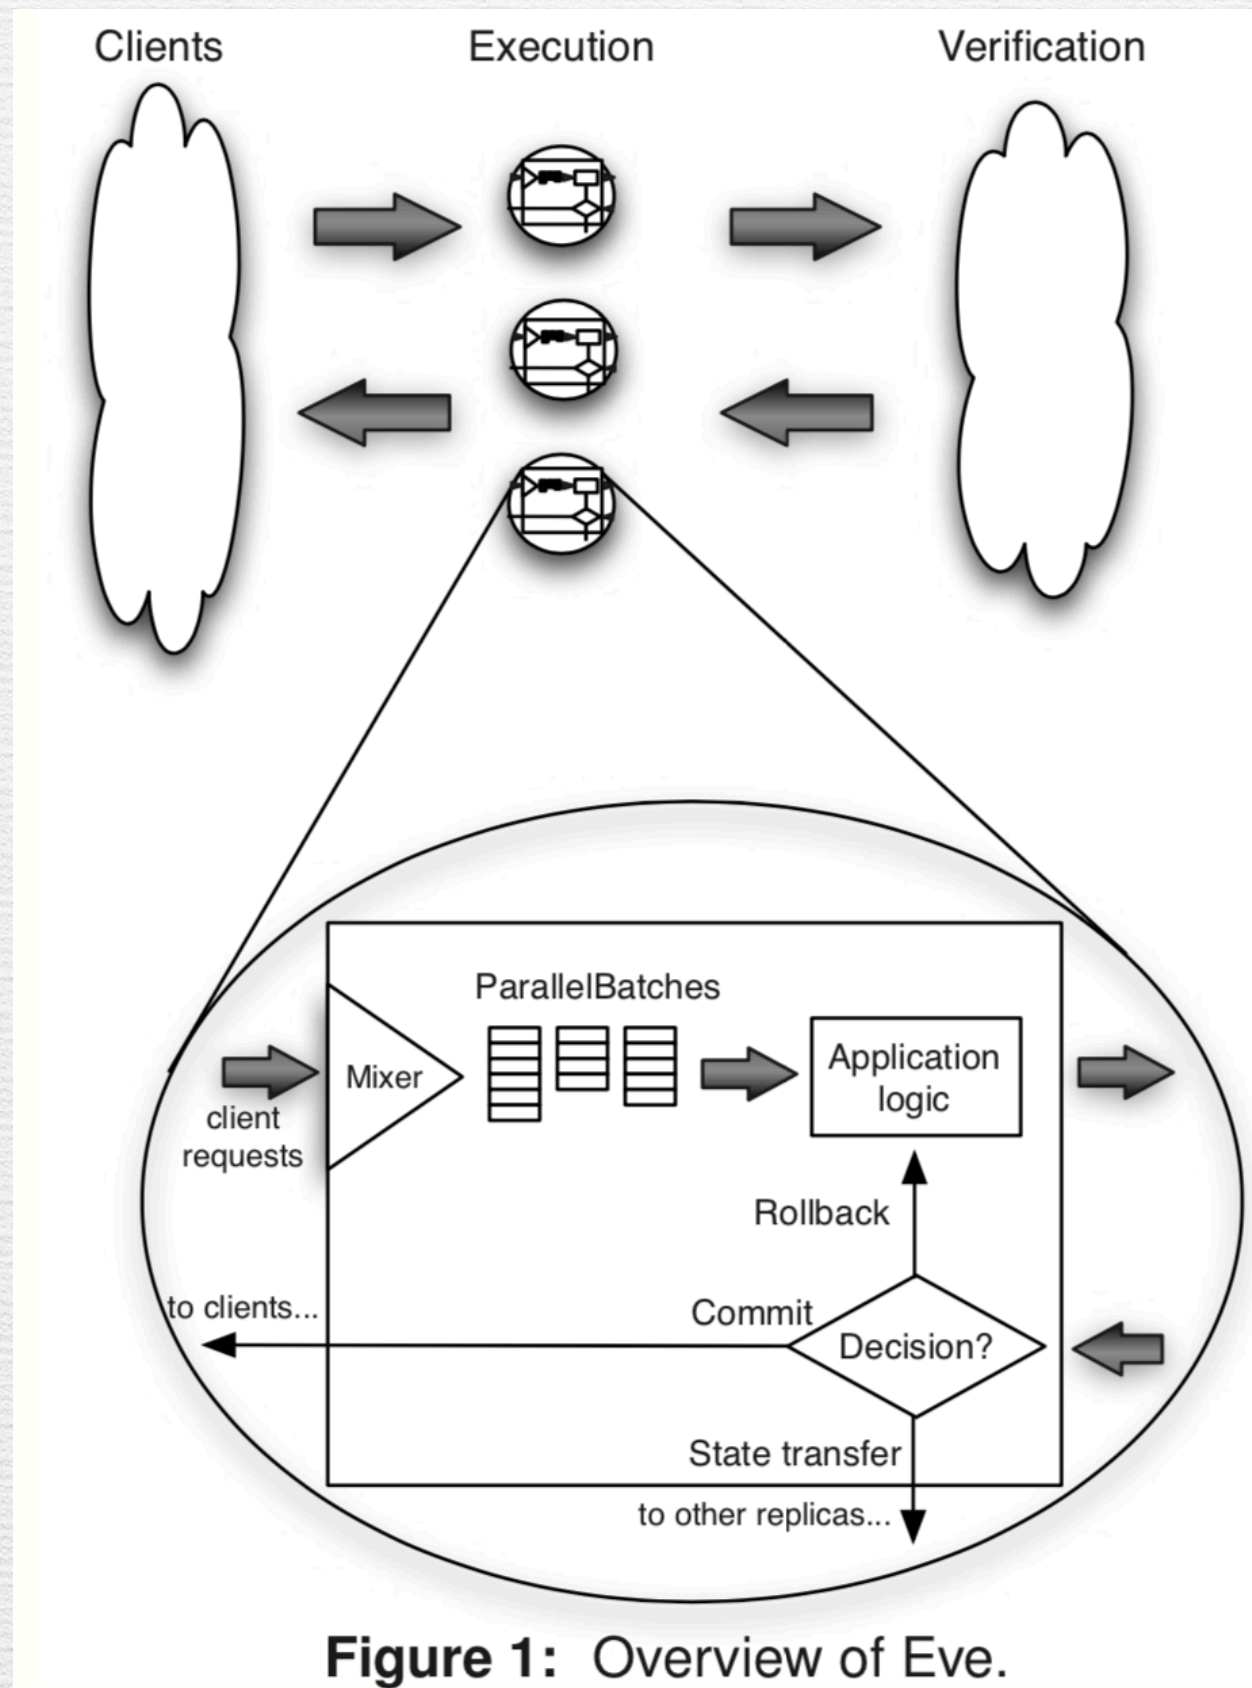


Figure 1: Overview of Eve.

Challenges

- Traditional idea for replicas: **hard to implement** on **multi-core servers**
 - Processing the **same sequence** of requests
- Modern servers: **parallel** execution can cause **divergence** of correct states and outputs
 - Most SMR systems require servers to process requests sequentially: a replica finishes executing one request before beginning to execute the next
- Traditional agree-execute method can cause violation of the safety requirements (correct replicas agree on the same state and output)
 - First agree on the order in which requests are to be executed and then execute them

Challenges

- Recent efforts to enforce deterministic parallel execution do not work well
 - Reason 1: the practical limitations of current implementations (e.g. high overhead)
 - Reason 2: many modern replication algorithms do not actually execute operations in the same order at every replica to achieve better performance
- Problems of deterministic execution:
 - Current techniques for deterministic multithreading either **require hardware support or are too slow**
 - **Semantic gap** exists between modern SMR protocols and the techniques used to achieve deterministic multithreading.

Challenges

- Current methods to discover the semantics of the requests do not work well
 - Current methods:
 - Aims to achieve replica coordination without forcing all replicas to process identical sequences of inputs
 - Read-only optimization: many modern SMR systems no longer insist that read requests be performed in the same order at all replicas
 - Preferred quorum optimization: allows read requests to be executed only at a preferred quorum of replicas, rather than at all replicas
 - Problems:
 - Deterministic multithreading techniques **know nothing of the semantics** of the operations they perform
 - Read-only optimizations and preferred quorum operations **violate the assumption that all replicas receive identical sequences of inputs** and hence lead correct replicas to diverge
 - Read-only requests advance a replica's instruction counter and may cause the replica to acquire additional read locks

Improvement Based on Previous Techniques

- Replicas do not have to execute requests in the same order
- Requests are partitioned in **batches**
- Allowing **different replicas** to execute requests within **their own batches in parallel**
 - Ensuring that the system's important state and output at each replica will match across enough replicas.
- Execute requests concurrently before verify that they have **agreed on the state and the output produced by a correct replica**
 - If too many replicas diverge so that a correct state/output cannot be identified, it guarantees safety and liveness by rolling back and sequentially and deterministically re-executing the requests

Improvement Based on Previous Techniques

- Minimizes divergence through a **mixer** stage
 - Applying application-specific criteria to produce groups of requests that are unlikely to interfere
 - Making repair efficient through incremental state transfer and fine-grained rollbacks
- Taking methods to guarantee that replication remains safe and live
 - The agreement is delayed until after execution, even if the program is correct under un-replicated parallel execution
 - When necessary, falling back to sequential re-execution even if the mixer allows interfering requests in the same group.

Basic Steps of Execution Stage

- Eve's execution stage tries to **make divergence unlikely**
- Batching
 - Clients send their requests to the current **primary execution replica**
 - The primary execution replica divide clients' requests into batches, **assigns each batch a sequence number**, and sends them back to **all other execution replicas**.
 - The primary execution replica sends other data needed to process the requests along with the sequence number
 - Ensuring that distinct batches are processed in order.

Basic Steps of Execution Stage

- Mixing
 - **ParallelBatches**: groups of requests that can be executed in parallel
 - Each replica runs the same deterministic mixer to partition each batch into these **ParallelBatches**
 - The mixer ensures that different interleavings will not produce diverging results at distinct replicas.

Basic Steps of Execution Stage

- Parallel Execution
 - Each replica executes the parallelBatches in the order specified by the deterministic mixer.
 - Then a replica computes a hash of its application state and of the outputs corresponding to requests in that batch.
 - This hash, along with its sequenceNumber and the hash for previous batch, constitute a token that is sent to the verification stage in order to discern whether the replicas have diverged

Mixer Design

- The mission of the mixer: to **identify requests that may be executed in parallel** with low false negative and false positive rates
 - False negatives will cause conflicting requests to be executed in parallel, creating the potential for divergence and rollback
 - False positives will cause requests that could have been successfully executed in parallel to be serialized, reducing the parallelism of the execution
- The mixer parses each request and tries to predict which state it will access
 - It may cause two requests to conflict when they access the same object in a read/write or write/write manner.

Mixer Design

- Methods to avoid putting together conflicting requests:
 - The mixer starts with an empty parallelBatch and two (initially empty) hash tables, one for objects being read, the other for objects being written
 - The mixer then scans in turn each request, **mapping** the objects accessed in the request **to a read or write key**
 - Before adding a request to a parallelBatch, the mixer checks **whether that request's keys have read/write or write/write conflicts** with the keys already present in the two hash tables
 - If not, the mixer adds the request to the parallelBatch and adds its keys to the appropriate hash table
 - When a conflict occurs, the mixer tries to add the request to a different parallelBatch — or creates a new parallelBatch, if the request conflicts with all existing parallelBatches
- Experiment Results:
 - This simple mixer achieves good parallelism (acceptably few false positives), and has not presented any rollbacks (few or no false negatives)

Mixer Design

- Using the names of the tables accessed in read or write mode as read and write keys for each transaction

Transaction	Read and write keys
getBestSellers	read: item, author, order_line
getRelated	read: item
getMostRecentOrder	read: customer, cc_xacts, address, country, order_line
doCart	read: item write: shopping_cart_line, shopping_cart
doBuyConfirm	read: customer, address write: order_line, item, cc_xacts, shopping_cart_line

Figure 2: The keys used for the 5 most frequent transactions of the TPC-W workload.

State Management

- Problems
 - Moving from an agree-execute to an execute-verify architecture puts pressure on the implementation of state checkpointing, comparison, rollback, and transfer.
- Eve stores the state using a copy-on-write **Merkle tree**, whose root is a concise representation of the entire state
 - To achieve efficient state comparison and fine-grained checkpointing and rollback
 - First, it includes **only the subset of state that determines the operation of the state machine**, omitting other state that has no semantic effect on the state and output produced by the application
 - Second, it provides an abstraction wrapper on some objects to **mask variations across different replicas**
 - Eve manually annotates the application code to denote the objects that should be added to the Merkle tree and to mark them as **tainted** when they get modified.

State Management

- Problems of determinism Merkle trees:
 - It is challenging to maintain a deterministic Merkle tree structure under parallel multithread execution and parallel hash generation
 - First current solution: making memory allocation synchronized and deterministic
 - Ignores efforts paid in concurrent memory allocation
 - Second current solution: generating an ID based on object content and to use it to determine an object's location in the tree
 - Fails to work because many objects have the same content
- Solution by Eve
 - **Postponing adding newly created objects** to the Merkle tree until the end of the batch.
 - Eve scans existing modified objects, and if one contains a reference to an object not yet in the tree, Eve adds that object into the tree's next empty slot and iteratively repeats the process for all newly added objects
 - Reason
 - First, existing objects are already put at deterministic locations in the tree
 - Second, for a single object, Eve can iterate all its references in a deterministic order

State Management

- Challenges when implementing on Java:
 - Objects to which the Merkle tree holds a reference to are **not eligible for Java's automatic garbage collection**
 - Several standard set-like data structures in Java are **not oblivious to the order in which they are populated**
 - Two set-like data structures at different replicas may contain the same elements but may generate different checksums when added to a Merkle tree, which causes divergence
- Solution by Eve
 - Creating wrappers that **abstract away semantically irrelevant differences** between instances of set-like classes kept at different replicas
 - For each set-like data structure, the wrappers generate a deterministic list of all the elements it contains and a corresponding iterator
 - Two wrappers are needed for two Java interfaces (Set and Map) respectively

Verification Stage

- Problems of Execution Stage
 - There are still some divergencies that cannot vanish after execution stage
- Mission of Verification Stage
 - Determining whether **enough execution replicas have same state and responses** after executing a batch of requests
 - Ensuring that **divergences** such as conflicting requests in the same parallelBatch **cannot affect safety**
 - Ensuring that all correct replicas that have executed the i th batch of requests are guaranteed to have reached the **same final state** and produced the **same outputs**

Verification Stage

- Concurrency Bugs: deviations from an application's intended behavior that is caused by particular thread interleaving
- Reasons for Eve's ability to mask replica divergences caused by concurrency bugs
 - The mixer makes concurrency bugs less likely to occur by trying to avoid parallelizing requests that interfere
 - Concurrency bugs may manifest differently on different replicas

Basic Operations

- Agreement
 - The verification stage runs an agreement protocol to determine the final decision, which is either **commit** (if enough tokens match) or **rollback** (if too many tokens differ)
- Verification on whether replicas have diverged
 - If all tokens agree, the replicas' common final state and outputs are committed
 - If there is divergence, the agreement protocol tries to find a pair of final state and outputs that leads to a correct replica, and ensures that all correct replicas commit to that state and outputs.
 - Otherwise, it rolls back.

Basic Operations

- Commit
 - The execution replicas mark the corresponding sequence number as committed and **send the responses to corresponding clients**
- Rollback
 - The execution replicas **roll back** their state **to the latest committed sequence number** and **re-execute the batch** sequentially to guarantee progress

Asynchronous Byzantine Fault Tolerant

- Model
 - $u + \max(u, r) + 1$ execution replicas and $2u + r + 1$ verification replicas, which allows the system to remain live despite u failures (whether of omission or commission), and safe despite r commission failures and any number of omission failures
- Similarities between Eve and PBFT
 - Both protocols attempt to perform agreement among $2u + r + 1$ replicas ($3f + 1$ in PBFT terminology)
- Differences between Eve and PBFT
 - In PBFT the replicas try to agree on the output of a single node—the **primary**, while in Eve the object of agreement is the behavior of **a collection of execution replicas**
 - In PBFT the replicas try to agree on the **inputs** to the state machine, while in Eve replicas try to agree on the **outputs** of the state machine

Asynchronous Byzantine Fault Tolerant

- Steps:
 - 1 When an **execution replica** executes a batch of requests, it sends a $\langle \text{VERIFY}, v, n, T, e \rangle$ message to all **verification replicas**, where v is the current view number, n is the batch sequence number, T is the computed token for that batch, and e is the sending execution replica.
 - 2 When a **verification replica** receives sufficient **VERIFY** messages with matching tokens, it marks this sequence number as **preprepared** and sends a $\langle \text{PREPARE}, v, n, T, v \rangle$ message to all **other verification replicas**.
 - 3 When a **verification replica** receives sufficient matching **PREPARE** messages, it marks this sequence number as **prepared** and sends a $\langle \text{COMMIT}, v, n, T, v \rangle$ to all **other verification replicas**.

Asynchronous Byzantine Fault Tolerant

- Steps:

- 4 When a **verification replica** receives sufficient matching COMMIT messages, it marks this sequence number as **committed** and sends a $\langle \text{VERIFY-RESPONSE}, v, n, T, v \rangle$ message to **all execution replicas**.
- 5 If agreement can not be reached, it sends a $\langle \text{VERIFY-RESPONSE}, v, n, T, v, f \rangle$ message to **all execution replicas**, where **f** is a flag that indicates that the next batch should be executed sequentially to ensure progress
- 6 Upon receipt of $r + 1$ matching VERIFY-RESPONSE messages, an **execution replica** **e** judges whether the view number has increased and whether the agreed-upon token matches the one **e** previously sent, before choosing to **commit, transfer its state or rollback**.

Asynchronous Byzantine Fault Tolerant

- Commit
 - View number has **not increased**, and the token **matches** the execution replica previously sent
 - Releasing the responses to the corresponding clients.
- State Transfer
 - View number has **not increased**, but the token **does not match** the execution replica previously sent
 - This replica has diverged from the agreed-upon state and thus issues a state transfer request to other replicas.
- Rollback
 - View number has **increased**
 - Agreement could not be reached and the primary is changed

Synchronous Primary-backup

- The primary receives client requests and groups them into batches. When a batch **B** is formed, it sends a $\langle \text{EXECUTE-BATCH}, n, \mathbf{B}, \mathbf{ND} \rangle$ message to the backup, where n is the batch sequence number and **ND** is the data needed for consistent execution of nondeterministic calls.
- The backup sends its token to the primary, which compares it to its own token.
- If the tokens match, the primary marks this sequence number as stable and releases the responses to the clients.
- If the tokens differ, the primary rolls back its state to the latest stable sequence number and notifies the backup to do the same.

Protection Against Concurrency Bugs

- Concurrency bugs
 - It can lead to both omission faults and commission faults
 - It's **easy to repair** compared with other Byzantine faults
- Asynchronous case
 - When configured with $n = 2u + 1$ and $r = 0$, Eve guarantees the safety and liveness corresponding to the requirements of state machine replication, and correctness of the state machine itself despite up to u concurrency or omission faults.
 - It ensures that the committed state and outputs at correct replicas match and that **requests eventually commit**

Protection Against Concurrency Bugs

- Synchronous case
 - When configured with just $u + 1$ execution replicas, Eve can continue to operate with one replica if u replicas fail by omission.
 - Eve does not have spare redundancy and can not mask concurrency faults at the one remaining replica.
 - In both the synchronous and asynchronous case, when configured for $r = 0$, Eve enters extra protection mode after k consecutive batches for which all execution replicas provided matching, timely responses.
- Extra protection during good intervals
 - Good intervals: when there are no other replica faults or time-outs except than those caused by concurrency bugs
 - Eve uses **spare redundancy** to boost its best-effort protection against concurrency bugs in both the synchronous and asynchronous cases.

Results

- Eve provides **gain on execution speed** compared to traditional sequential execution approaches
- The mixer never parallelized requests it should have serialized
- The mixer could completely mask concurrency bugs
 - When the bug manifests only in one replica, Eve detects that the replicas have diverged and repairs the damage by rolling back and re-executing sequentially.
- Eve has better performance under heavier workloads
- The scalability decreases for larger objects, which is an artifact of the hashing library.

Results

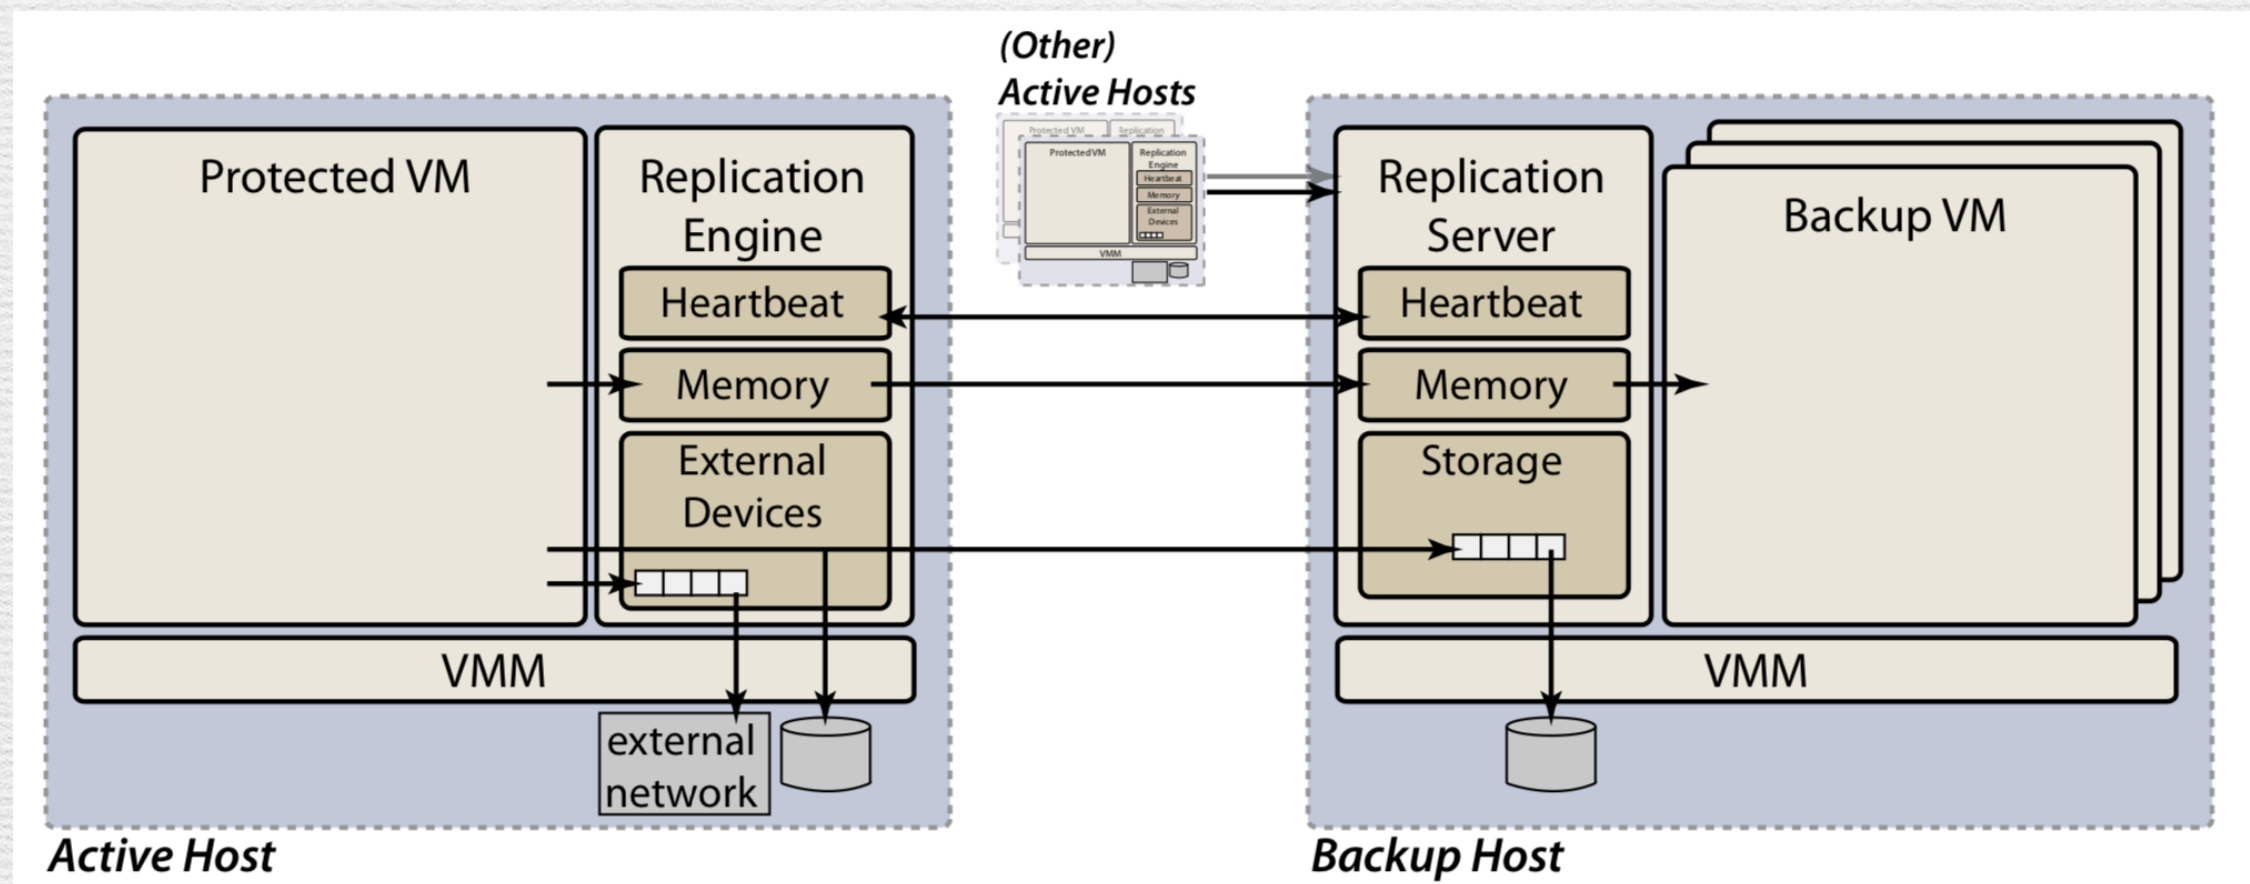
- The effect of false negatives on throughput
 - Even for 0% false negatives, the throughput drops as the pairwise conflict chance increases due to the decrease of available parallelism.
 - With a 1% conflict rate, the batch will be divided into only a few parallelBatches, so there is a good chance that conflicting requests will land in the same parallelBatch.
 - Increased false positive ratios can lead to lower throughput.
- Eve uses **significantly less bandwidth**, achieves **higher throughput**, and provides **stronger guarantees** compared to passive replication approaches.
- Eve uses a **dynamic batching scheme**:
 - The batch size decreases when the demand is low (providing good latency), and increases when the system starts becoming saturated, in order to leverage as much parallelism as possible.
 - It could keep its latency low while maintaining a high peak throughput

Comparison with other Works

- Semi-active Replication
 - It **weakens** state machine replication both in **determinism** and in **execution independence**
 - The primary executes nondeterministically and logs all the nondeterministic actions it performs
 - All other replicas then execute by deterministically reproducing the primary's choices
 - Relevant Works
 - Vandiver et al. : a Byzantine-tolerant semi-active replication scheme for transaction processing systems
 - Kim et al. : applying it to a transactional operating system.
- Shortcomings
 - Relaxing the requirement of independent execution makes these systems **vulnerable to commission failures**
 - That the same input is given to all replicas is violated in modern replication systems

Comparison with other Works

- Remus primary-backup system
 - The backup does not execute requests, but instead passively absorbs state updates from the primary
 - Shortcomings: it cannot tolerate any commission failures



Conclusion

- Goal of Eve:
 - Solving the divergence that may occur when executing State Machine Replication in multi-core servers
- Core Method:
 - Mixer that produce groups of requests that are unlikely to interfere
 - Fine-grained rollbacks in verification stage

dank u
ju faleminderit
Tack
Asante 谢谢 Tak mulțumesc
kiitos
Salamat! Gracias
Terima kasih Aliquam
Merci
ありがとう Dankie Obrigado
köszönöm grazie
Aliquam Go raibh maith agat
děkuji
Thank you